



DEPLOYMENT AND PERFORMANCE TESTING

NodeJS for Enterprise : Day 2/2

DAY 2 - OUTLINE

- Deployment
 - NodeJS as Service
 - Nginx as Reverse Proxy
 - Nginx Upstreams
- Basic Performance Testing with JMeter
- Performance Monitoring Tools with NewRelic and Keymetrics
- Real-time Application Analytics Sample Case (Hummingbird)
- Security Discussion
- Q&A



DEPLOYMENT

SIMPLE DEPLOYMENT

Start small with built-in static server, node.js app & database all on the same server

Sample of Tools:

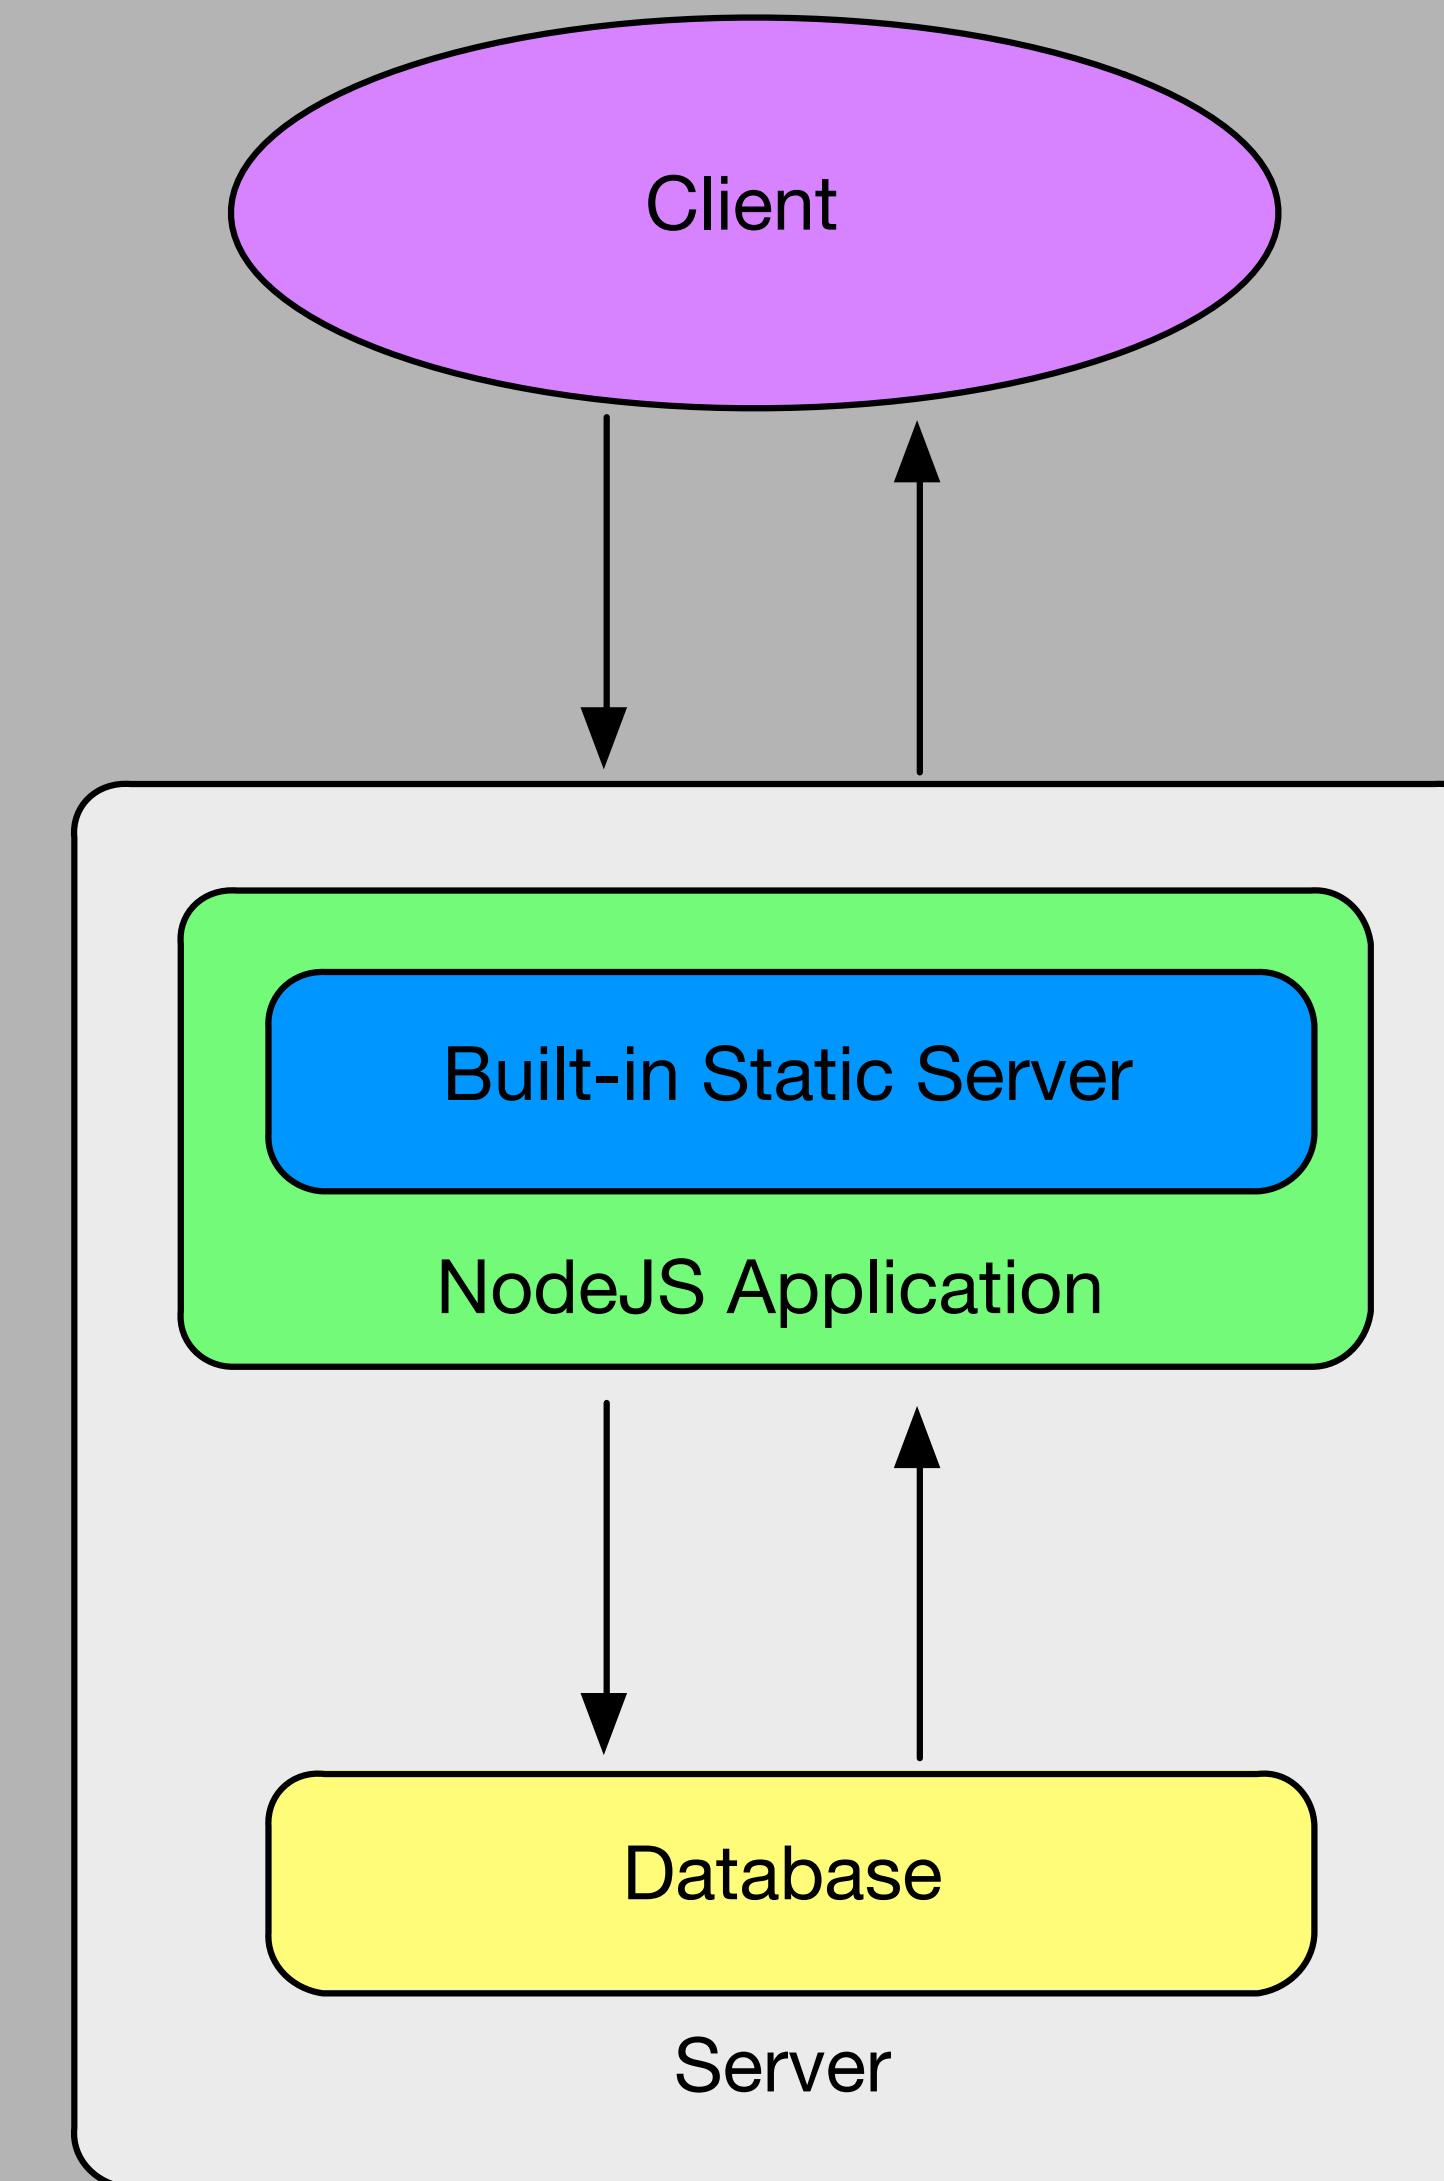
built-in static server

OR

node-windows

OR

pm2





LAB 2-I

node server.js

BUILT-IN STATIC SERVER

```
$ node server.js
```



LAB 2-2

node-windows, node-mac, node-linux

INSTALLATION

```
#node-windows, node-mac, node-linux
```

```
$ npm install -g node-windows
```

in your project root

```
$ npm link node-windows
```

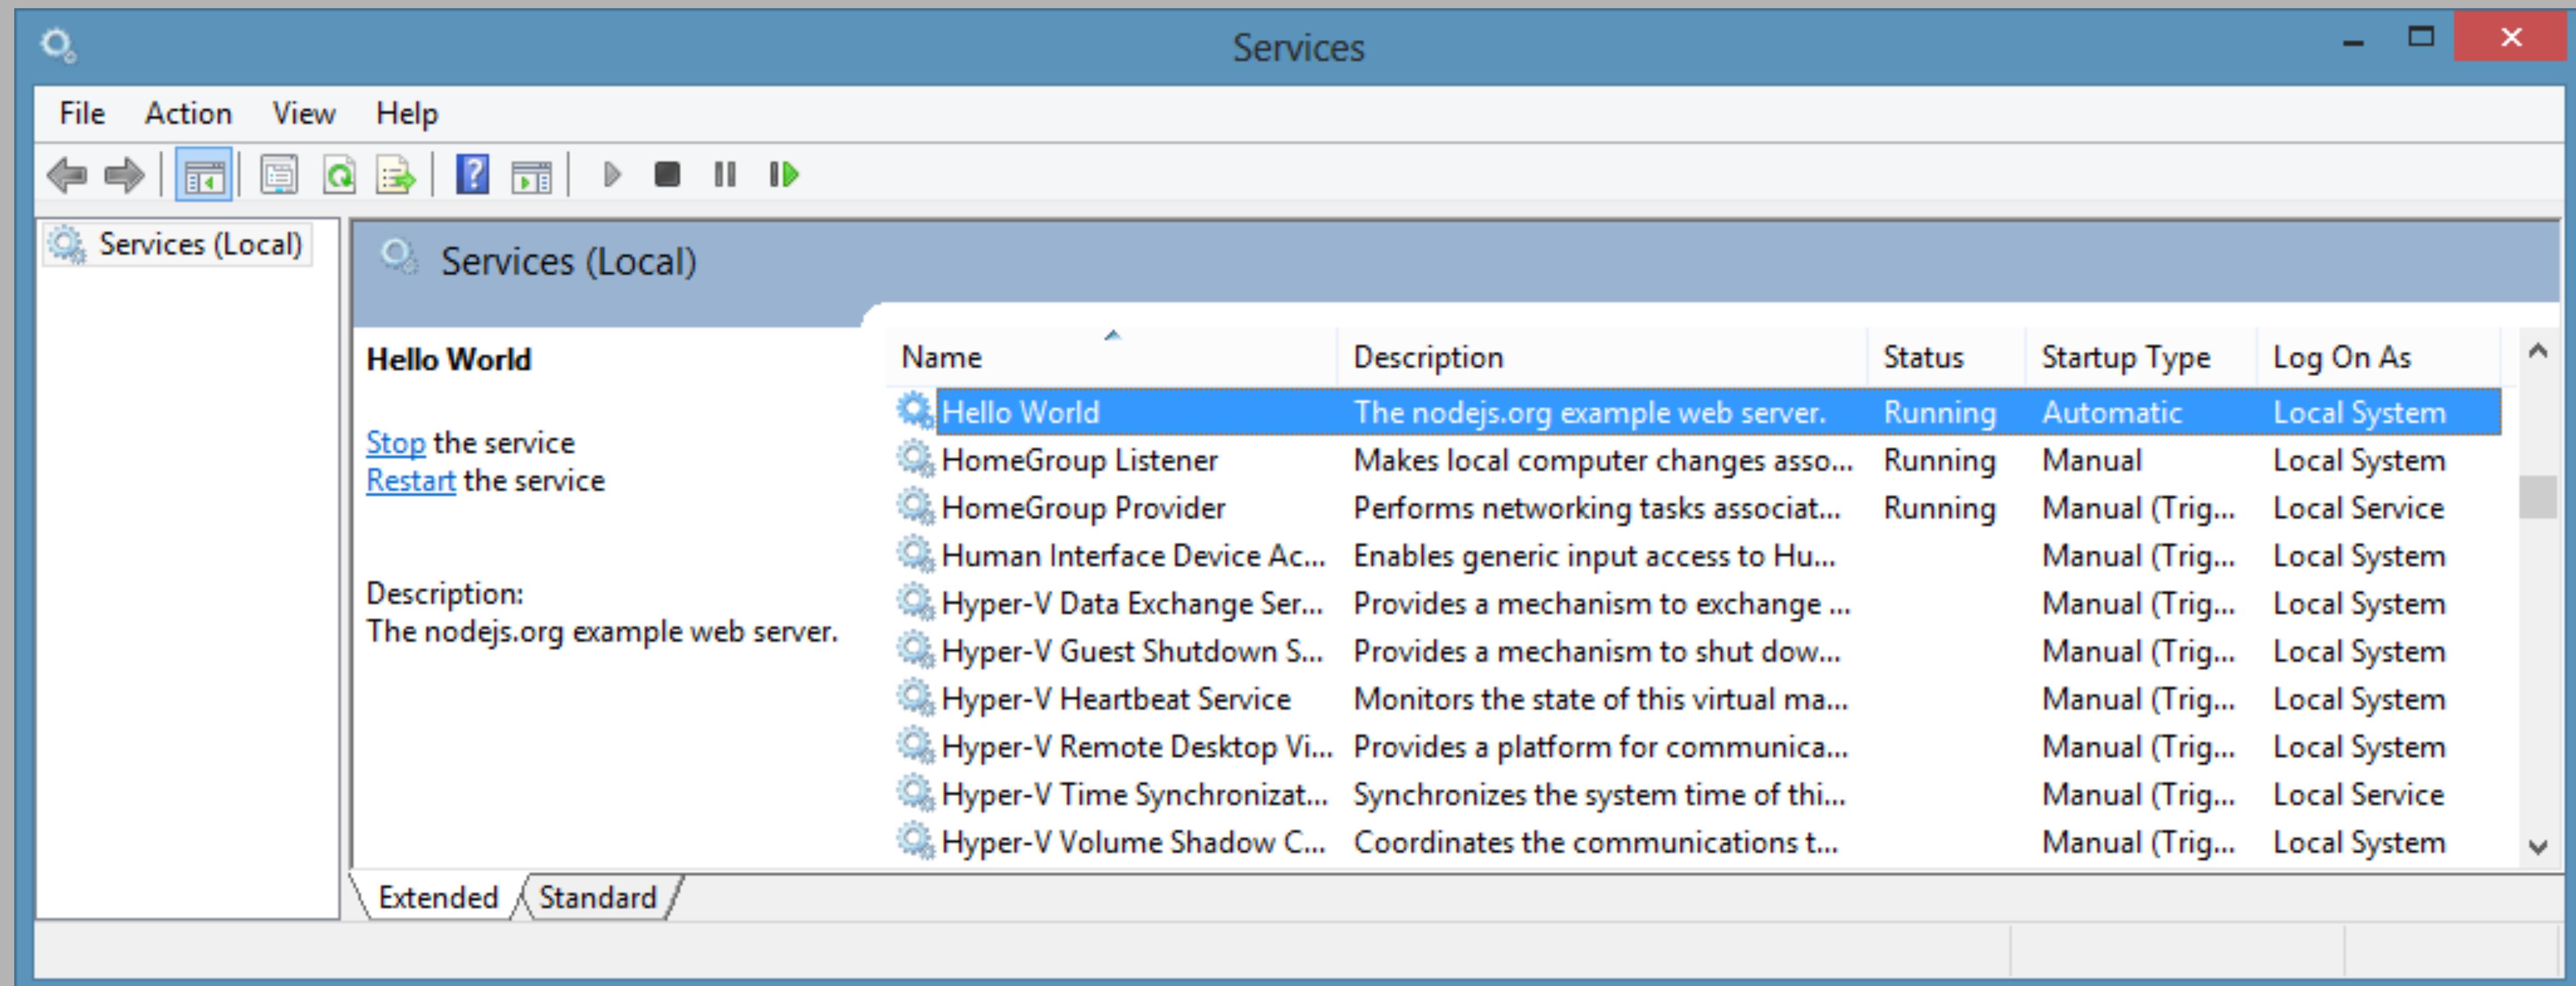
INSTALL SERVICE

```
var Service = require('node-windows').Service;

// Create a new service object
var svc = new Service({
  name:'Hello World',
  description: 'The nodejs.org example web server.',
  script: 'C:\\path\\to\\helloworld.js'
});

// Listen for the "install" event, which indicates the
// process is available as a service.
svc.on('install',function(){
  svc.start();
});
```

INSTALL SERVICE



NODE-WINDOWS EVENTS

- `install` - Fired when the script is installed as a service.
- `alreadyinstalled` - Fired if the script is already known to be a service.
- `invalidinstallation` - Fired if an installation is detected but missing required files.
- `uninstall` - Fired when an uninstallation is complete.
- `start` - Fired when the new service is started.
- `stop` - Fired when the service is stopped.
- `error` - Fired in some instances when an error occurs.

ENVIRONMENT VARIABLE

```
var svc = new Service({
  name:'Hello World',
  description: 'The nodejs.org example web server',
  script: 'C:\\path\\to\\helloworld.js',
  env: [
    {
      name: "HOME",
      value: process.env["USERPROFILE"]
    },
    {
      name: "TEMP",
      value: path.join(process.env["USERPROFILE"], "/temp")
    }
  );
}

svc.install();
```

LOGGING

```
var EventLogger = require('node-windows').EventLogger;  
  
var log = new EventLogger('Hello World');  
  
log.info('Basic information.');//  
log.warn('Watch out!');//  
log.error('Something went wrong.');//
```

LOGGING

The screenshot shows the Windows Event Viewer interface. The left pane displays a navigation tree with 'Event Viewer (Local)', 'Custom Views', and 'Windows Logs' expanded, showing 'Application', 'Security', 'Setup', 'System', 'Forwarded Events', 'Applications and Services Log', and 'Subscriptions'. The right pane shows the 'Application' log with 115 events. A specific event is selected, showing details in a modal window.

Event Viewer

File Action View Help

Event Viewer (Local)

Custom Views

Windows Logs

- Application
- Security
- Setup
- System
- Forwarded Events

Applications and Services Log

Subscriptions

Application Number of events: 115

Level	Date and Time	Source	Event ID	Task Category
! Error	3/24/2013 2:16:40 PM	Hello World	1000	None
! Warning	3/24/2013 2:16:40 PM	Hello World	1000	None
! Information	3/24/2013 2:16:25 PM	Hello World	1000	None
! Warning	3/24/2013 2:16:24 PM	Hello World	1000	None
! Information	3/24/2013 2:16:08 PM	Hello World	1000	None
! Warning	3/24/2013 2:16:07 PM	Hello World	1000	None
! Information	3/24/2013 2:15:52 PM	Hello World	1000	None
! Information	3/24/2013 2:15:51 PM	helloworld.exe	0	None
! Information	3/24/2013 2:15:51 PM	helloworld.exe	0	None

Event 1000, Hello World

General Details

Too many restarts within the last 60 seconds. Please check the script.

UNINSTALL SERVICE

```
var Service = require('node-windows').Service;

// Create a new service object
var svc = new Service({
  name:'Hello World',
  script: require('path').join(__dirname,'helloworld.js')
});

// Listen for the "uninstall" event so we know when it's done.
svc.on('uninstall',function(){
  console.log('Uninstall complete.');
  console.log('The service exists: ',svc.exists);
});

// Uninstall the service.
svc.uninstall();
```



LAB 2-3

pm2

<http://pm2.keymetrics.io/>

INSTALLATION

```
$ npm install pm2 -g
```

OR

```
$ npm install pm2@latest -g
```

START NODE.JS APP PROCESS

```
$ pm2 start app.js --name "app-name"
```

```
$ pm2 start app.js -i 0
```

```
$ pm2 start app.js -i max
```

LIST NODE.JS APP PROCESS

```
$ pm2 list
```

```
$ pm2 jlist
```

```
$ pm2 prettylist
```

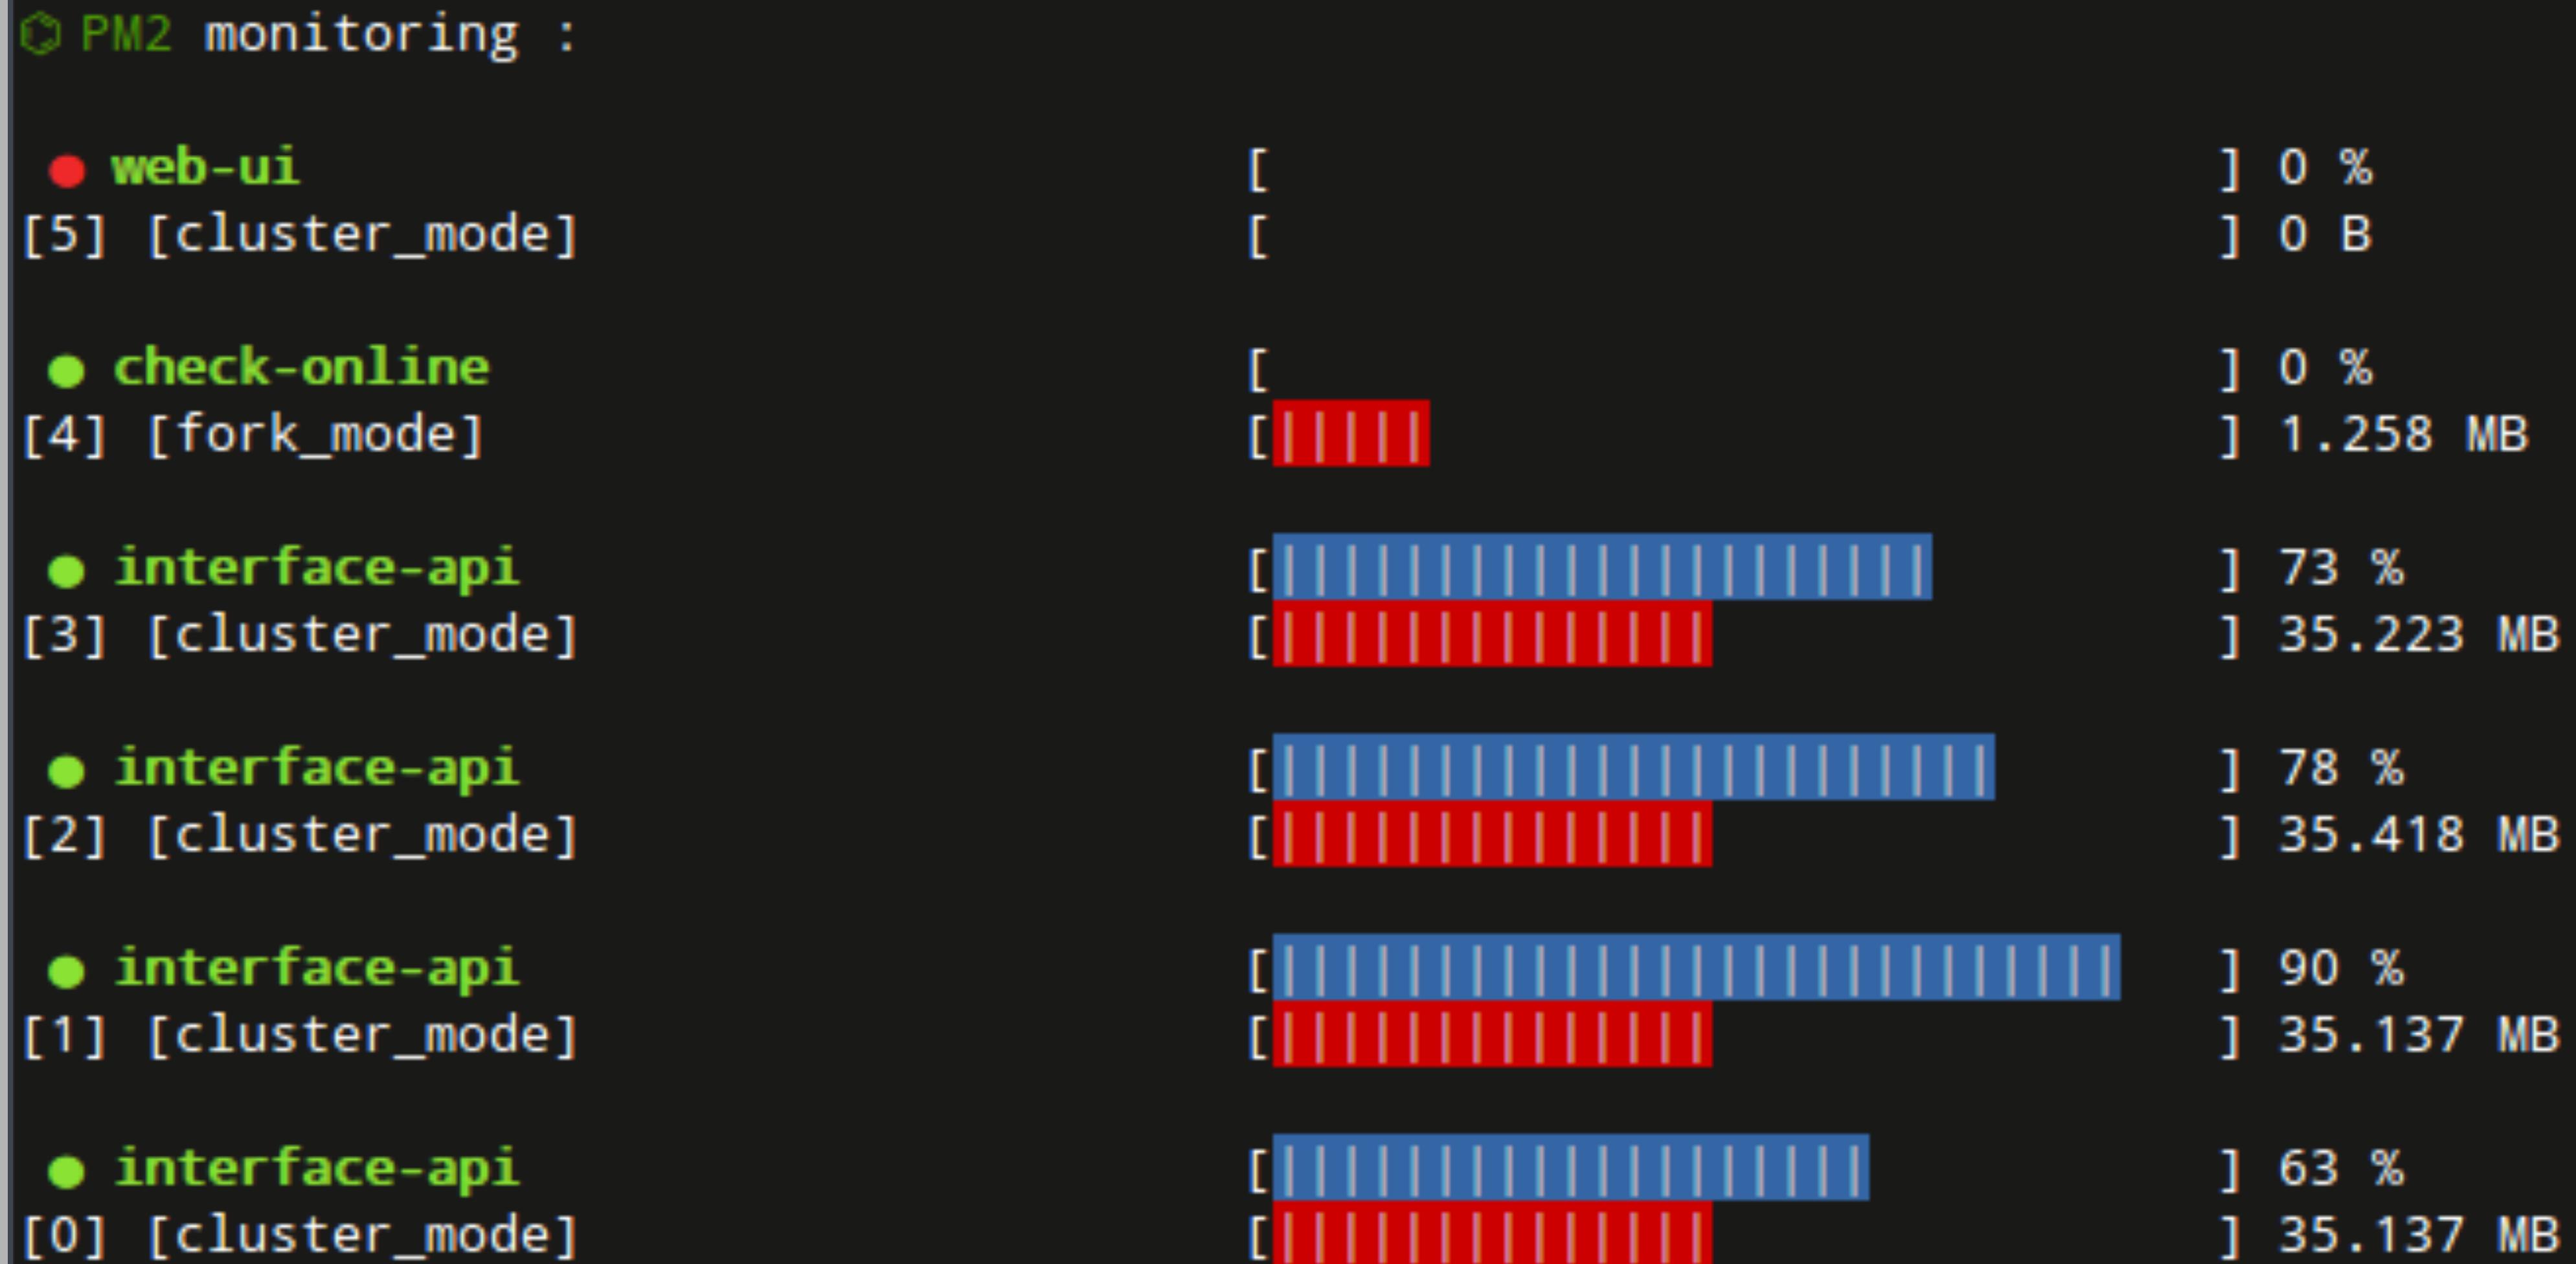
```
$ pm2 describe 0
```

```
$ pm2 updatePM2          # Update in memory pm2
```

```
$ pm2 ping                # Ensure pm2 daemon has been launched
```

MONITORING NODE.JS

```
$ pm2 monit
```



NODE.JS APP LOGS

```
$ pm2 logs [—raw]
```

```
$ pm2 flush
```

```
$ pm2 reloadLogs
```

STOP/RESTART NODE.JS APP PROCESS

```
$ pm2 stop all          # Stop all processes  
$ pm2 restart all       # Restart all processes  
  
$ pm2 reload all        # Will 0s downtime reload (for NETWORKED apps)  
$ pm2 gracefulReload all # Send exit message then reload (for networked apps)  
  
$ pm2 stop 0             # Stop specific process id  
$ pm2 restart 0          # Restart specific process id
```

DELETE NODE.JS APP PROCESS

```
$ pm2 delete 0
```

```
$ pm2 delete all
```

```
# Will remove process from pm2 list
```

```
# Will remove all processes from pm2 list
```

AUTOSTART NODE.JS APP

สำหรับ Windows Platform

```
$ npm install pm2-windows-startup -g
```

```
$ pm2-startup install
```

สำหรับ Unix/Linux Platform

```
$ pm2 startup
```

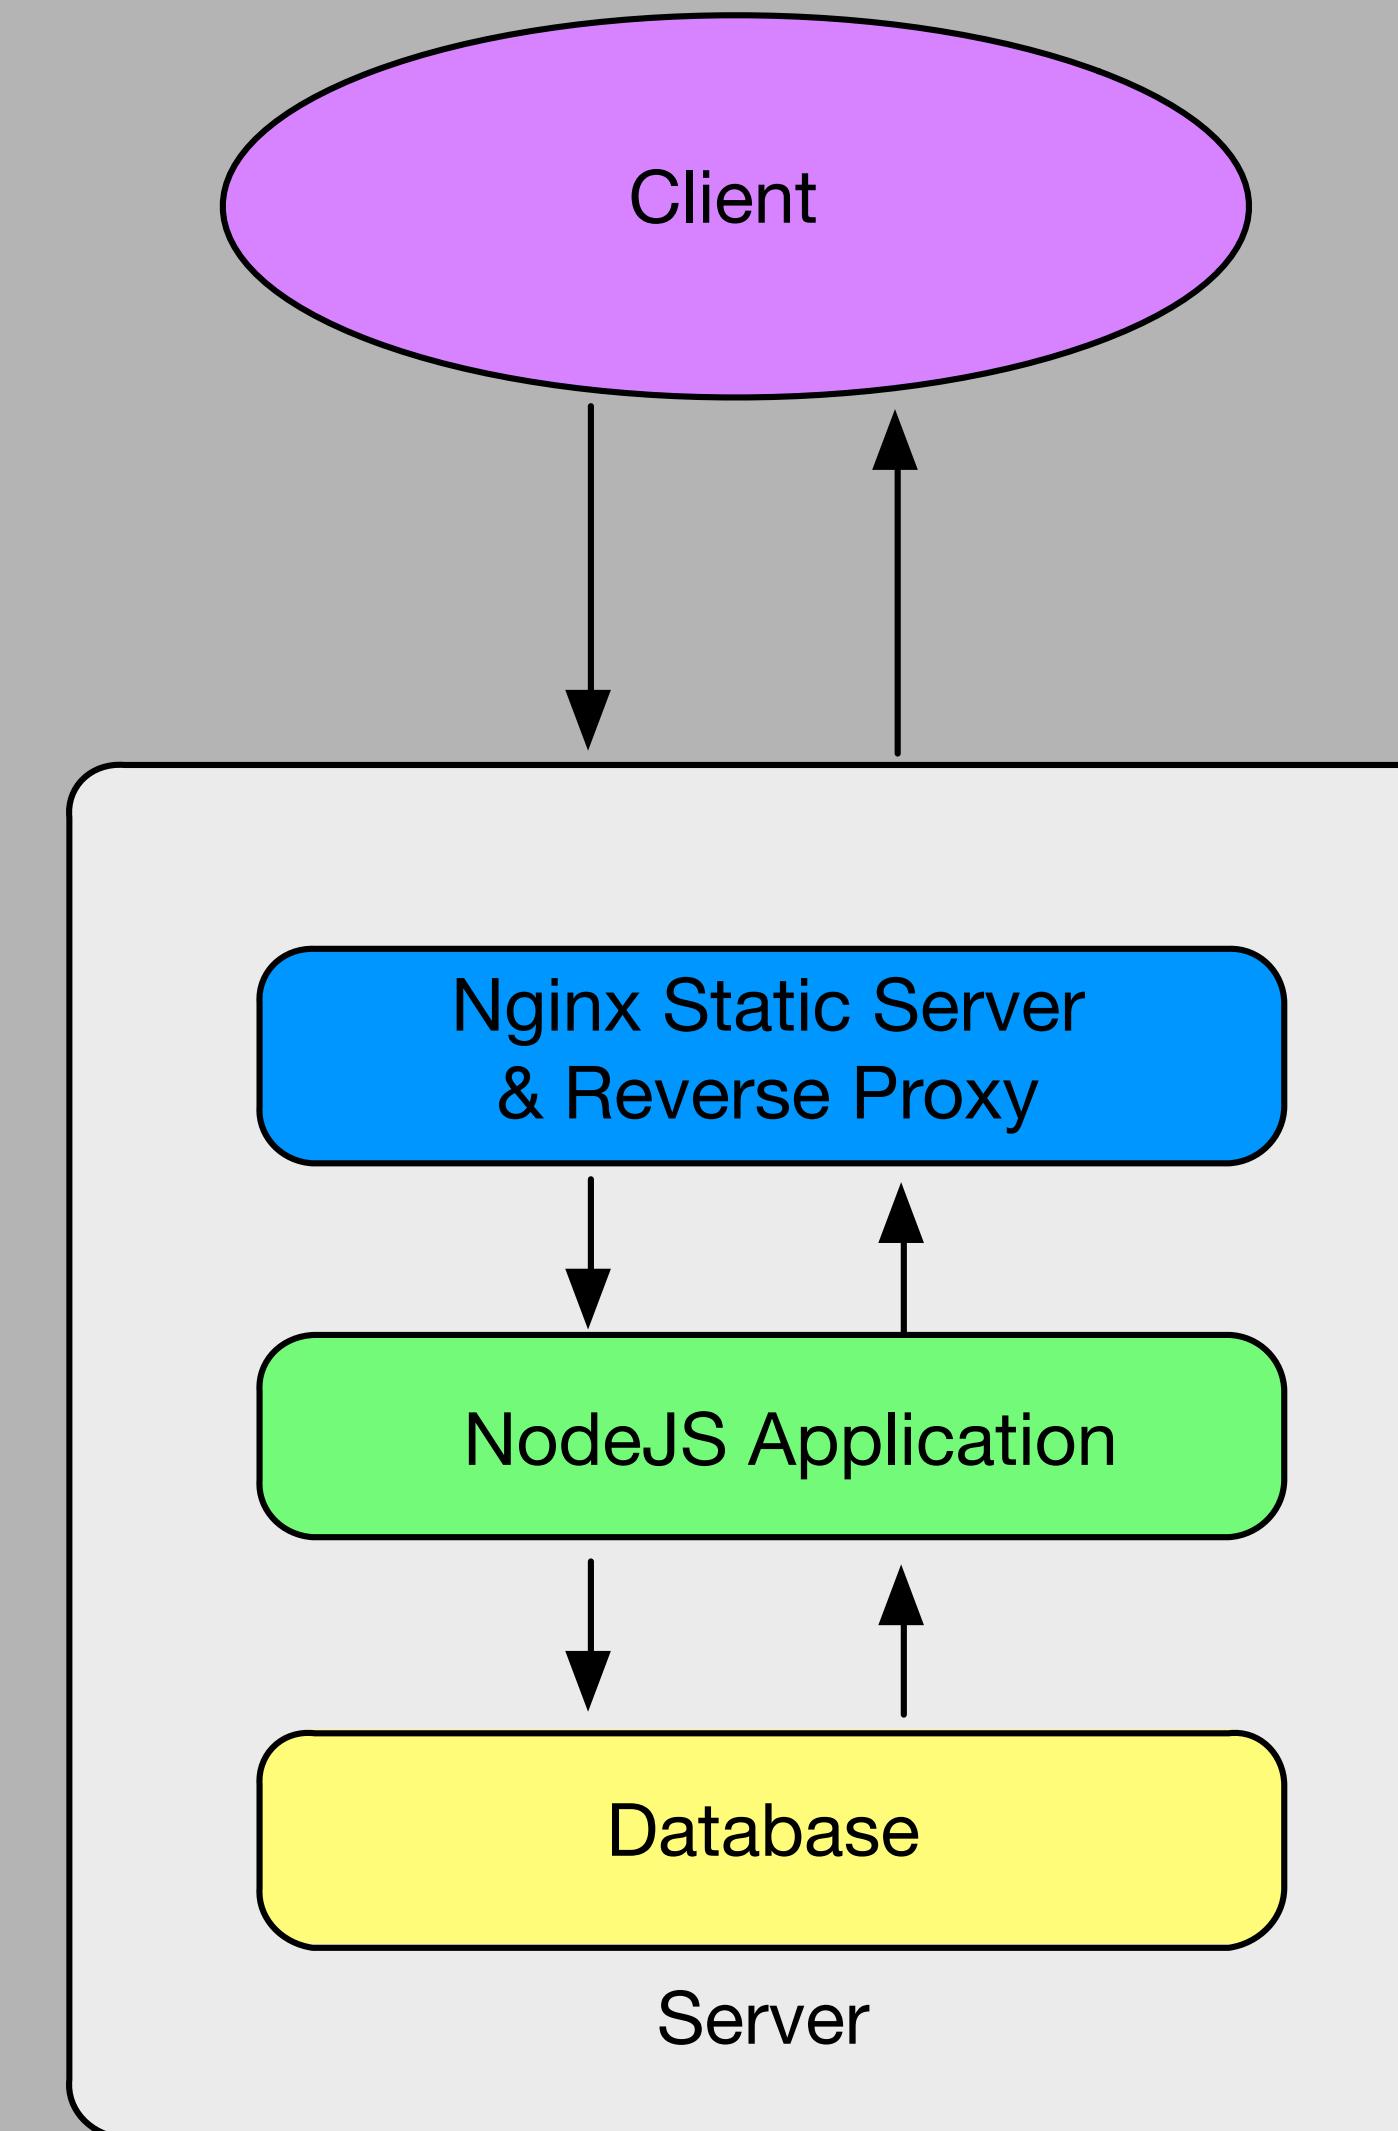
Process Autostart จะทำงานหลังจากสิ้นบันทึก Process list

```
$ pm2 save
```

SCALING NODEJS

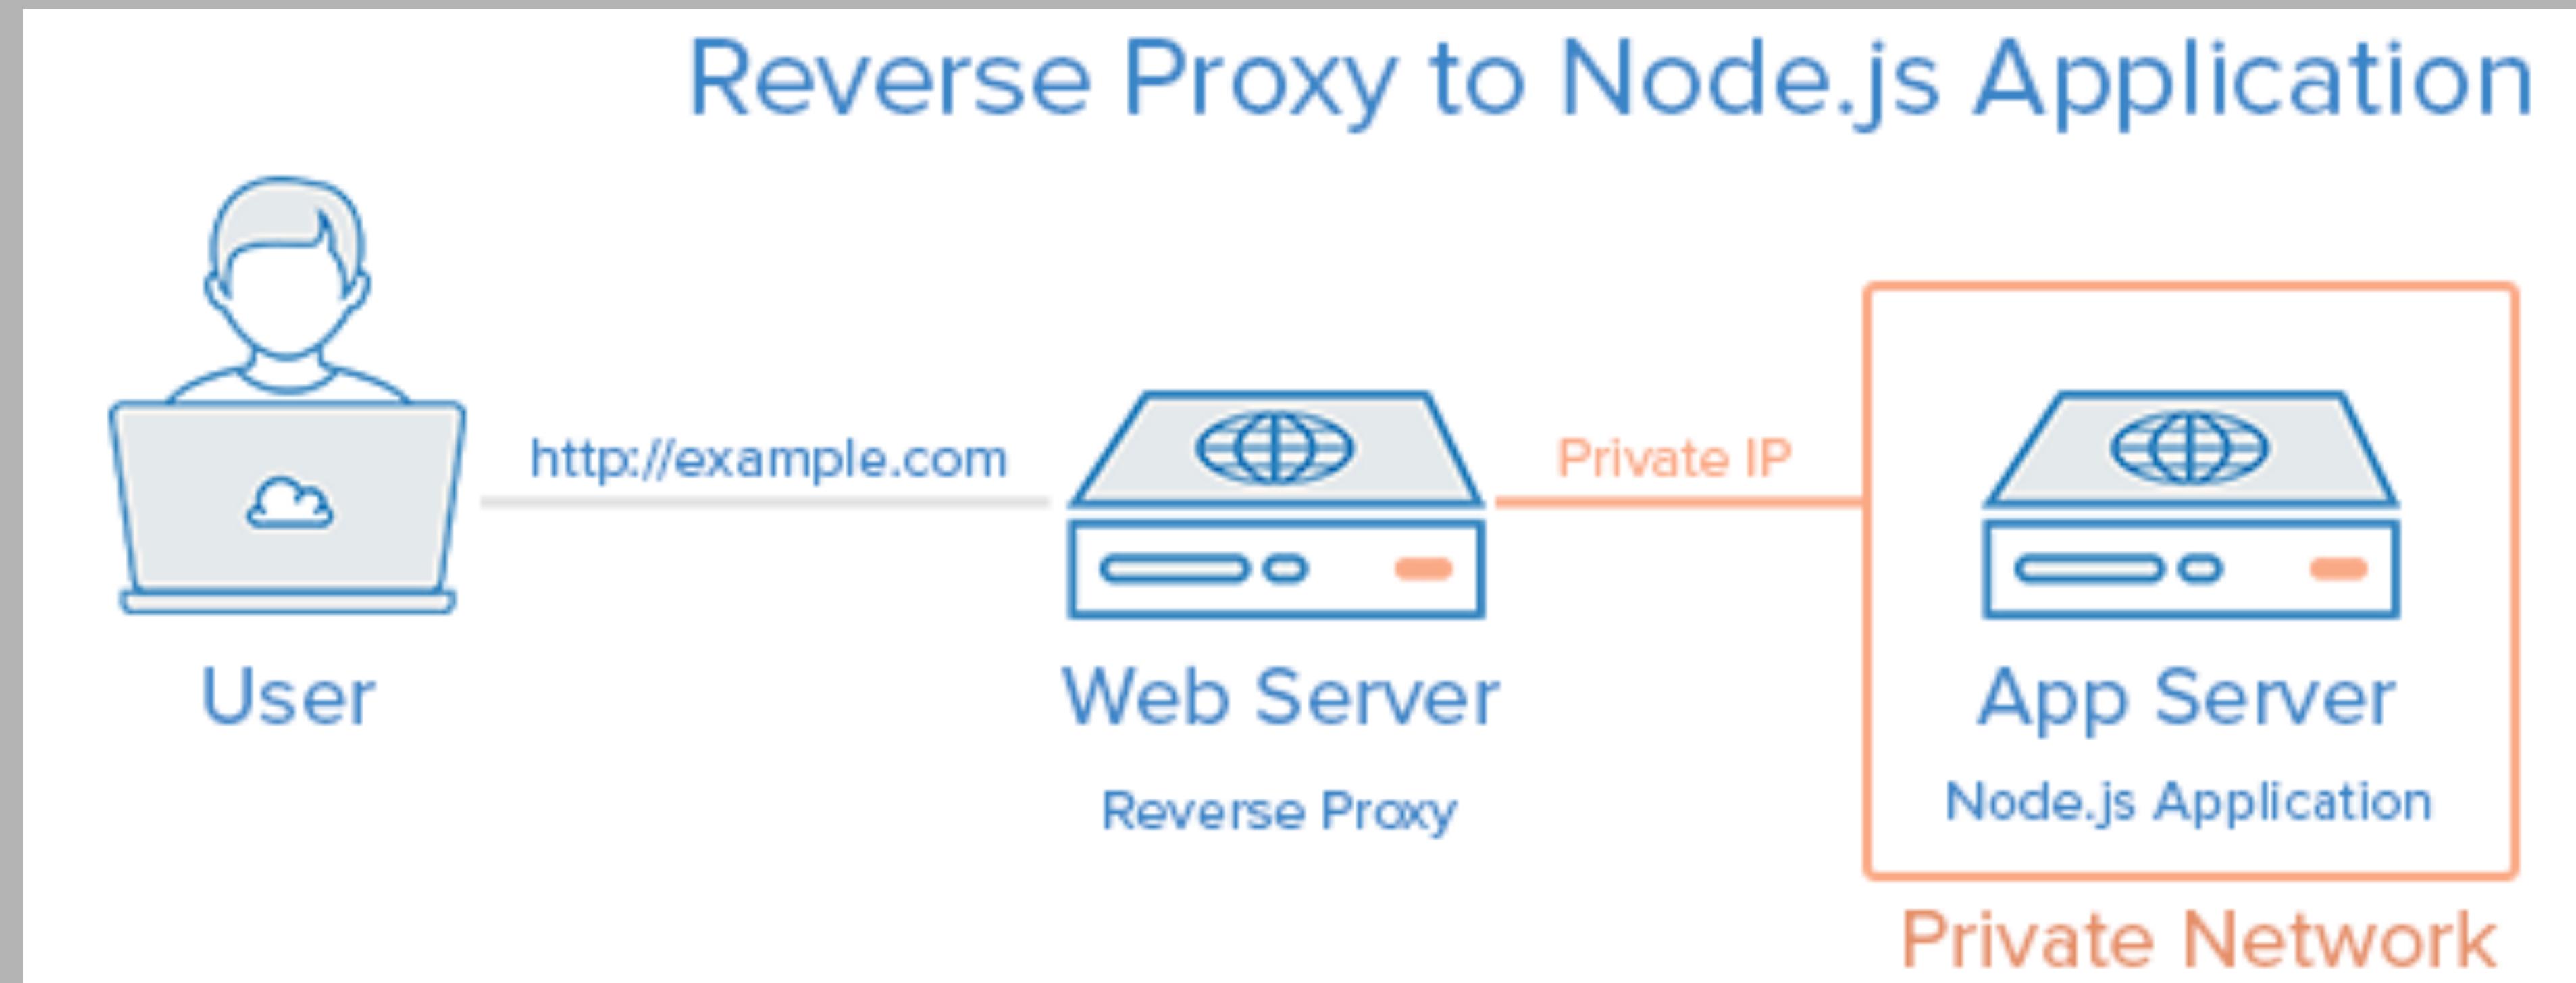
Use nginx as proxy server as well to load balance requests.

The number of node.js app instance depends on how many CPU cores on the machine.

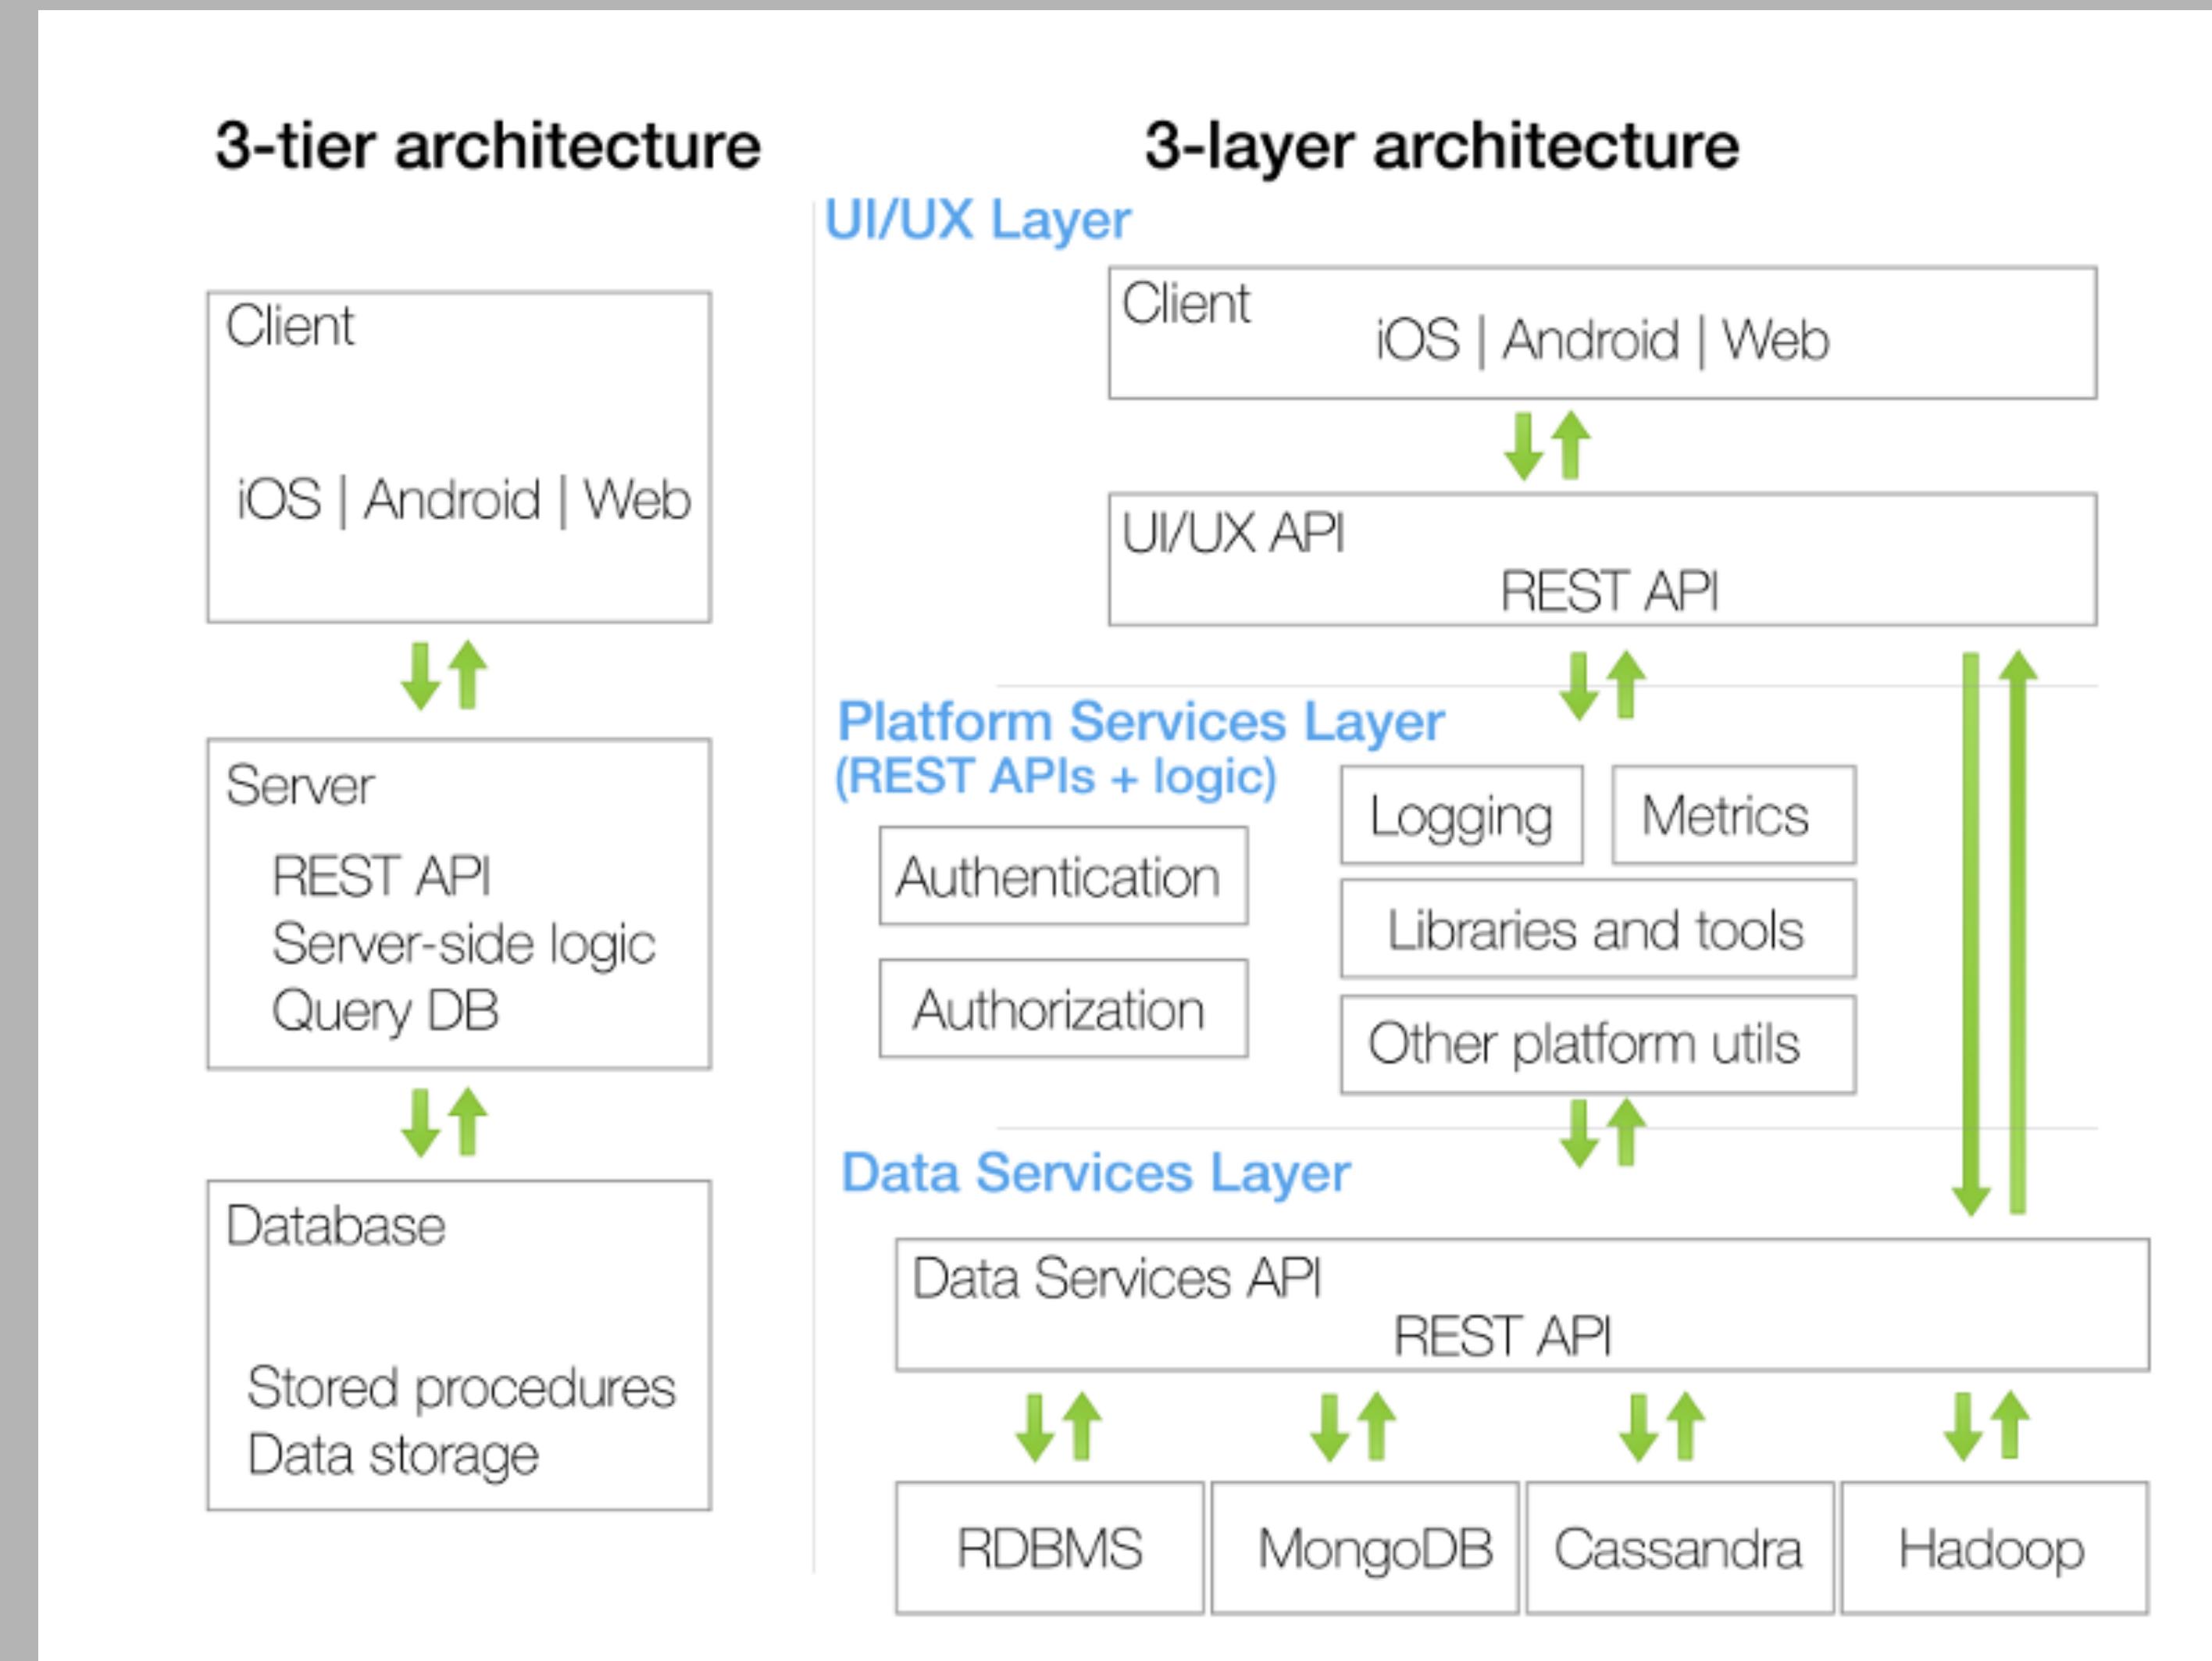


NODE.JS BEST PRACTICE

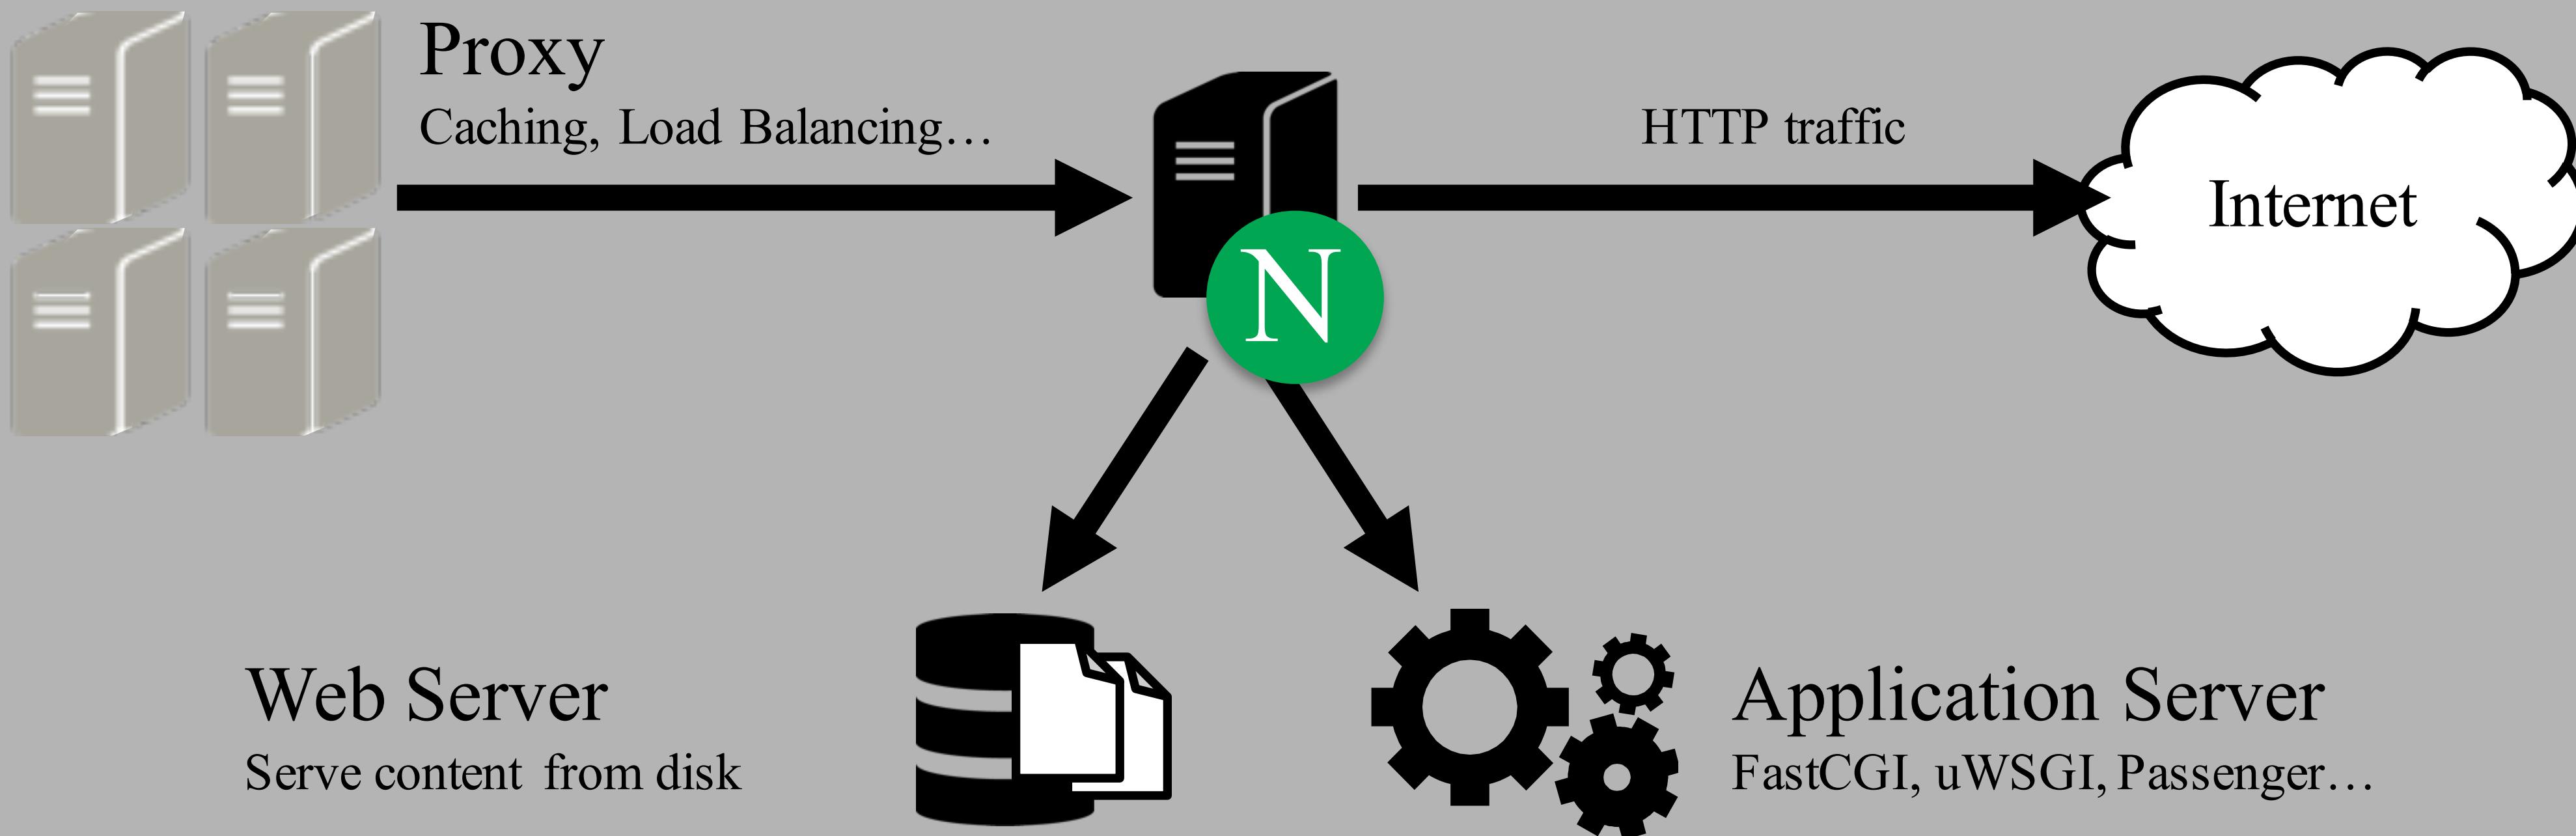
REVERSE PROXY



3-LAYER ARCHITECTURE



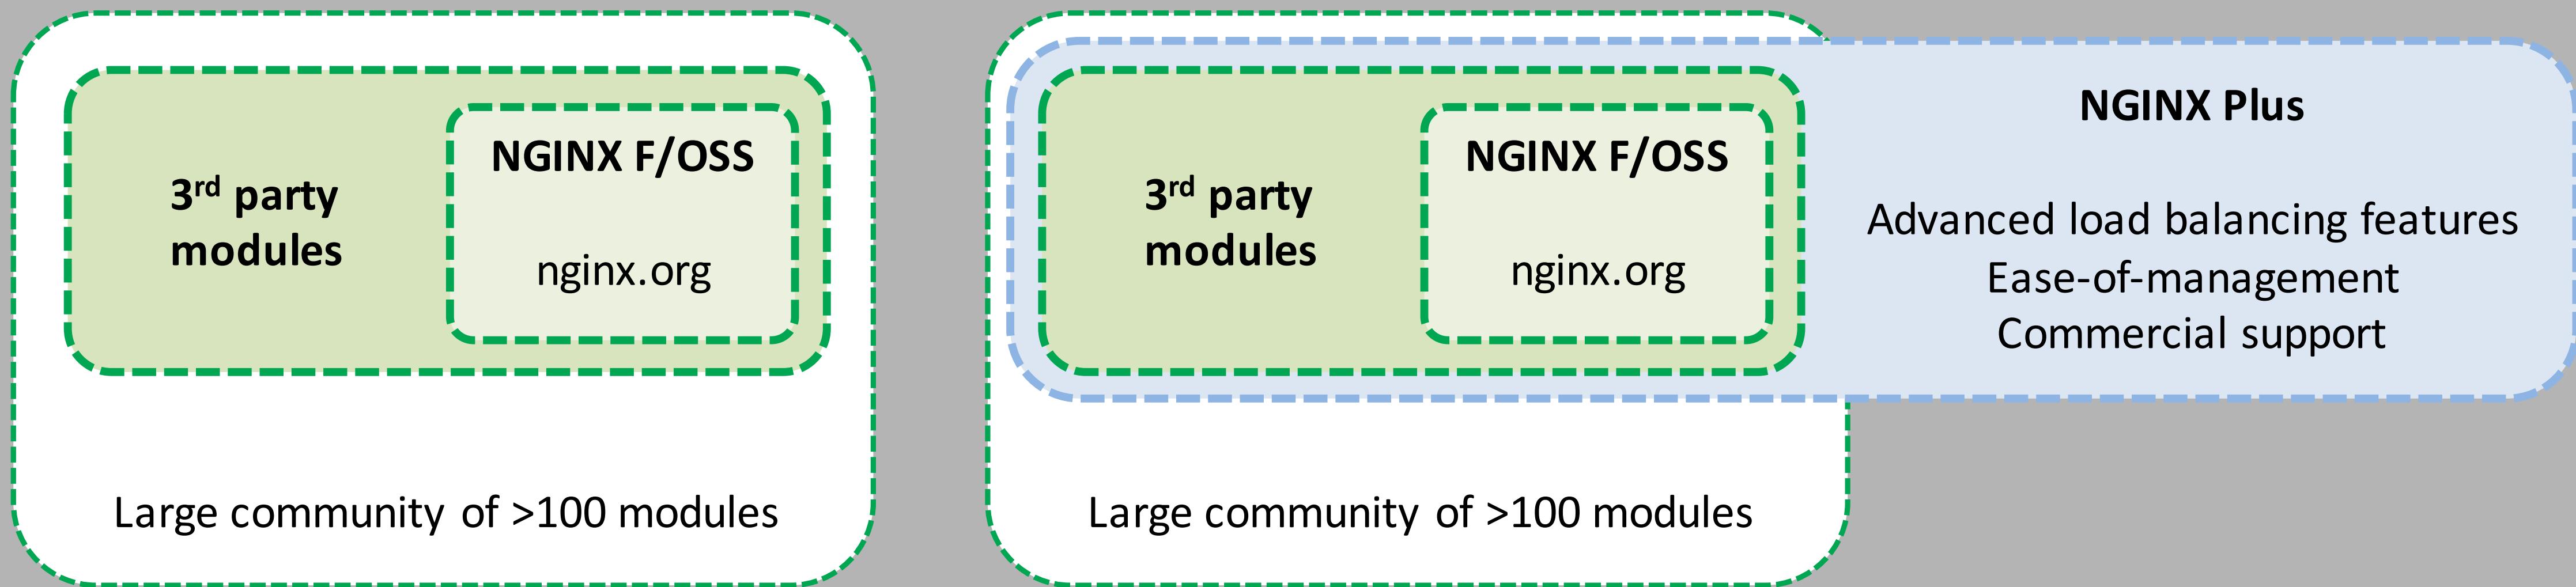
WHAT IS NGINX



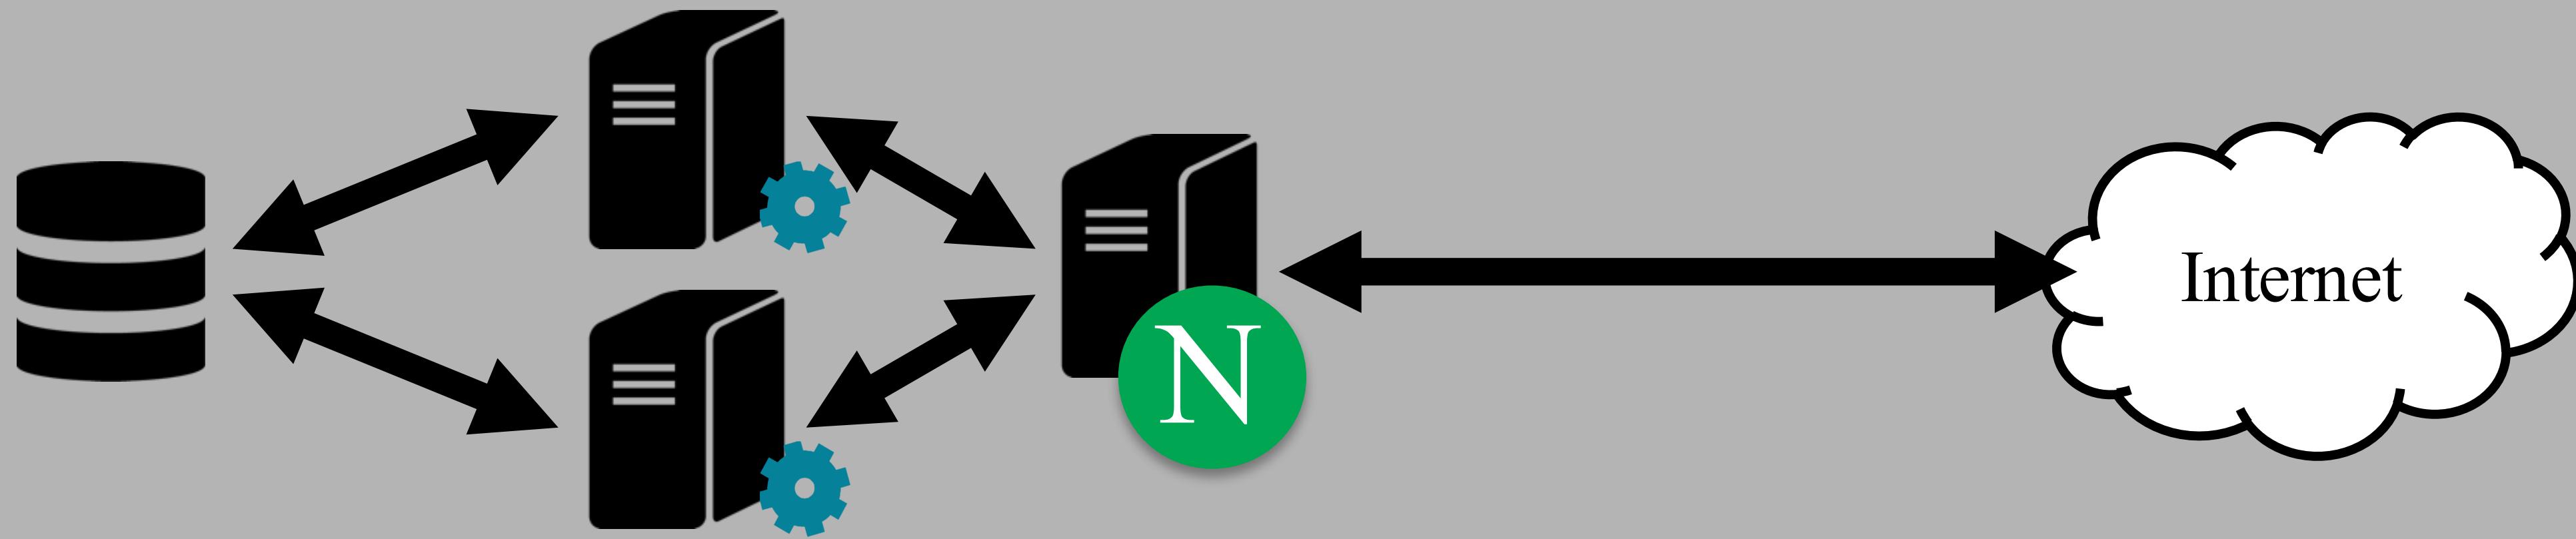
Advanced Features:

- | | | |
|--|--|--|
| <ul style="list-style-type: none"><input checked="" type="checkbox"/> Application Acceleration<input checked="" type="checkbox"/> SSL and SPDY termination<input checked="" type="checkbox"/> Performance Monitoring<input checked="" type="checkbox"/> High Availability | <ul style="list-style-type: none"><input checked="" type="checkbox"/> Bandwidth Management<input checked="" type="checkbox"/> Content-based Routing<input checked="" type="checkbox"/> Request Manipulation<input checked="" type="checkbox"/> Response Rewriting | <ul style="list-style-type: none"><input checked="" type="checkbox"/> Authentication<input checked="" type="checkbox"/> Video Delivery<input checked="" type="checkbox"/> Mail Proxy<input checked="" type="checkbox"/> GeoLocation |
|--|--|--|

NGINX AND NGINX PLUS



LOAD BALANCING



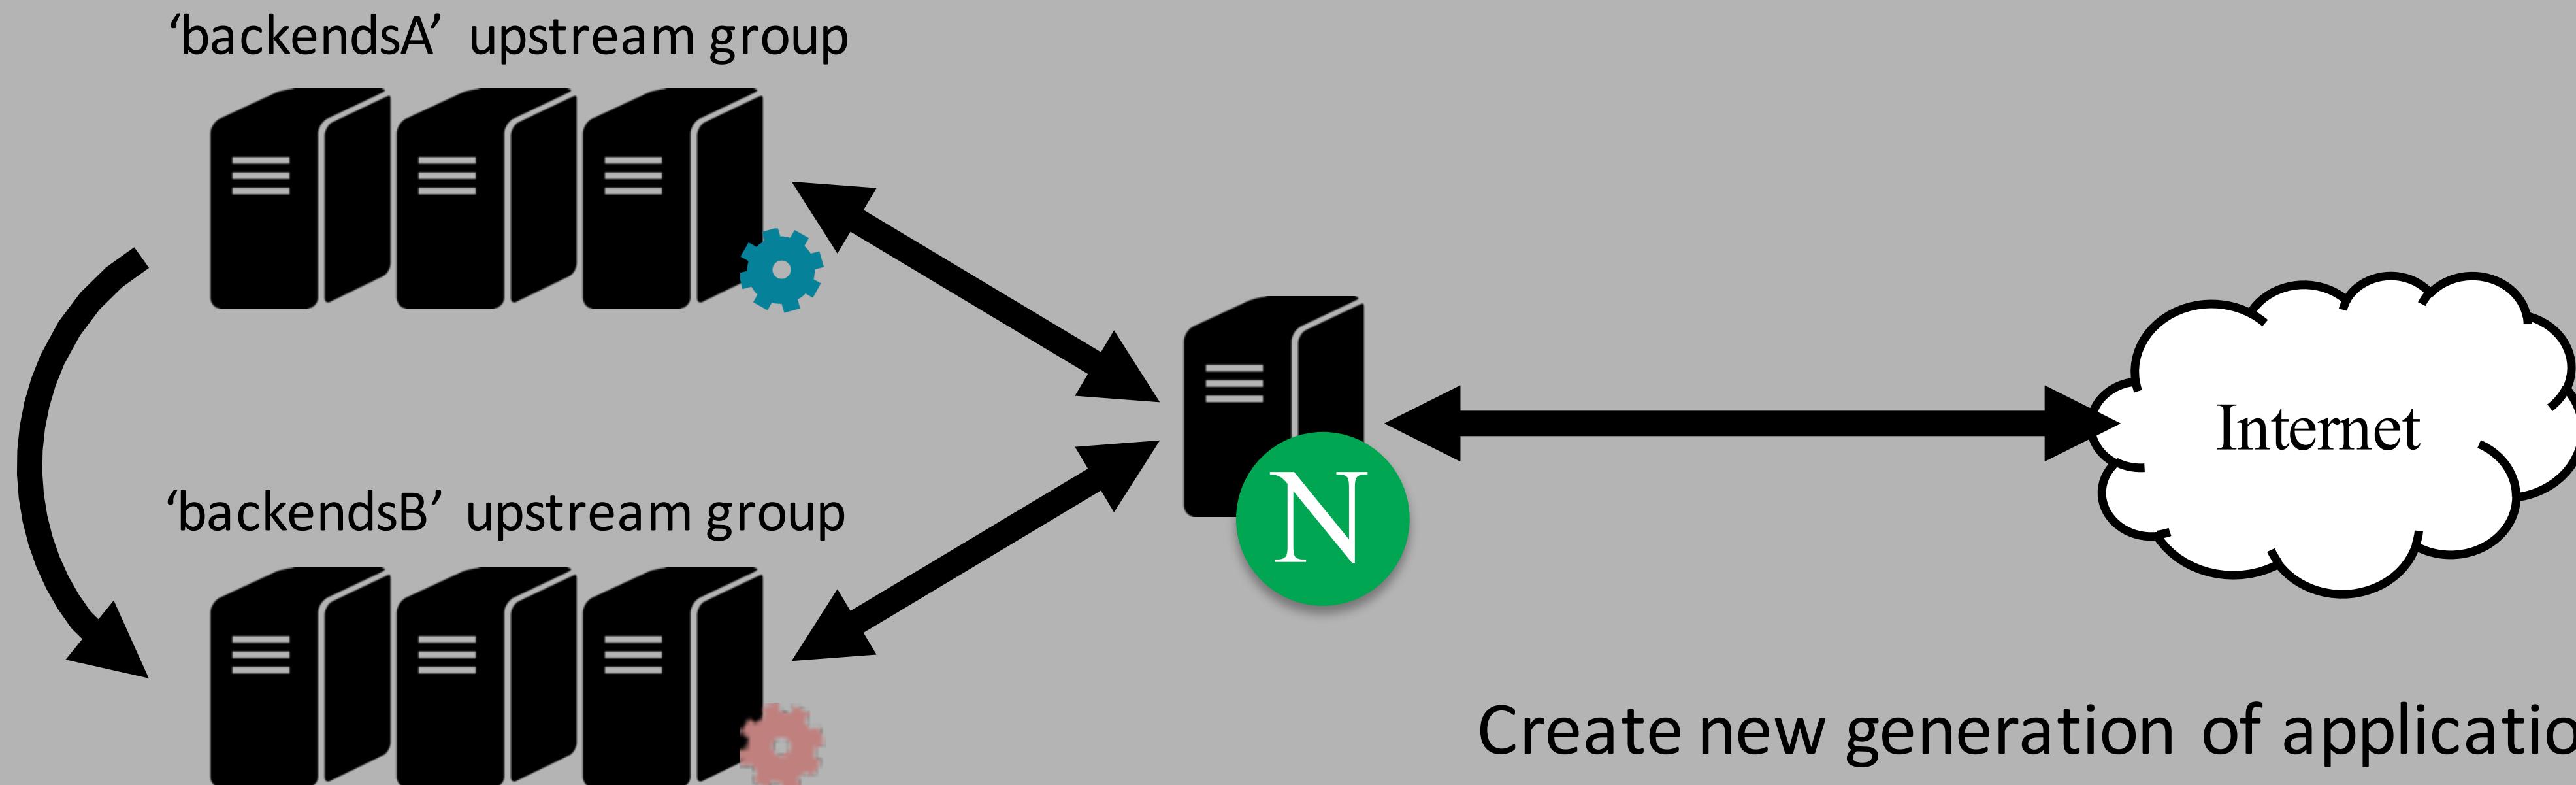
Why?

- Improved Application Availability
- Management
- Increased Capacity
- Advanced techniques e.g. A|B testing

How?

- DNS Round Robin
- Hardware L4 load balancer
- Software Reverse Proxy LB
- Cloud solution

LOAD BALANCING



Create new generation of application
Migrate users from old to new
Preserve sessions, no interruptions



LAB 2-4

nginx reverse proxy

ROOT CONTEXT

Unix/Linux

/etc/nginx/sites-available/default

```
server {  
    listen 80;  
  
    server_name example.com;  
  
    location / {  
        proxy_pass http://APP_PRIVATE_IP_ADDRESS:8080;  
        proxy_http_version 1.1;  
        proxy_set_header Upgrade $http_upgrade;  
        proxy_set_header Connection 'upgrade';  
        proxy_set_header Host $host;  
        proxy_cache_bypass $http_upgrade;  
    }  
}
```

CUSTOM CONTEXT

Unix/Linux

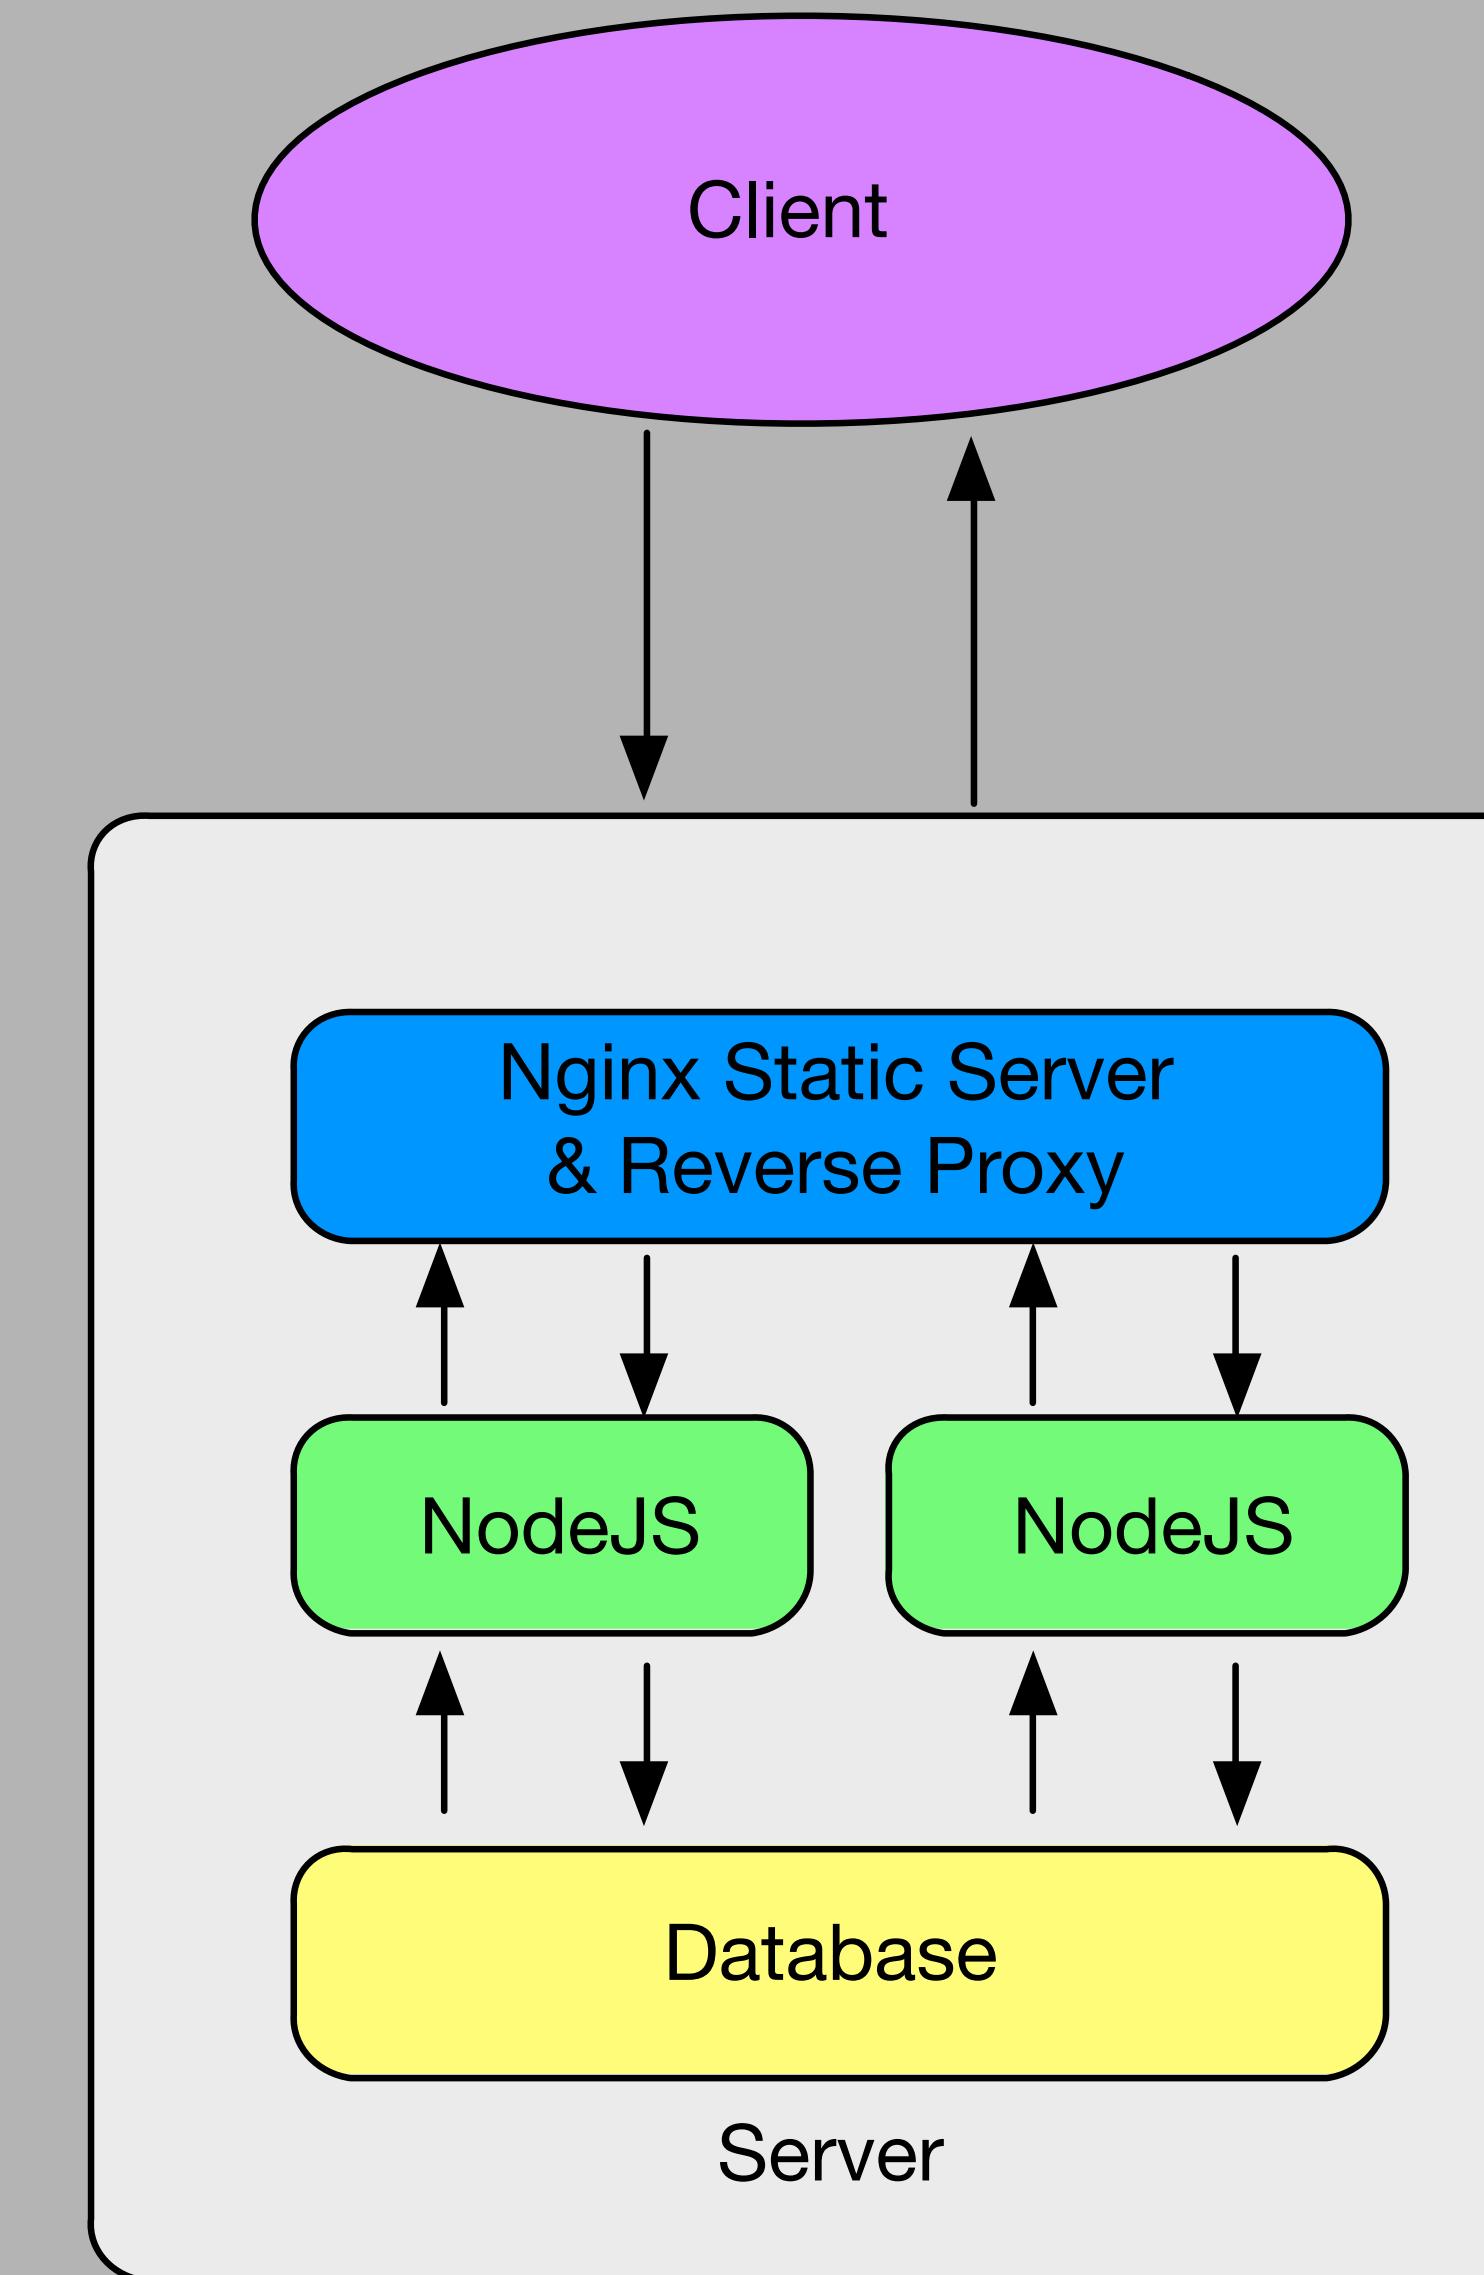
/etc/nginx/sites-available/default

```
server {  
    listen 80;  
  
    server_name example.com;  
  
    location /app2 {  
        proxy_pass http://APP_PRIVATE_IP_ADDRESS:8081;  
        proxy_http_version 1.1;  
        proxy_set_header Upgrade $http_upgrade;  
        proxy_set_header Connection 'upgrade';  
        proxy_set_header Host $host;  
        proxy_cache_bypass $http_upgrade;  
    }  
}
```

SCALING NODEJS

Use nginx as proxy server as well to load balance requests.

The number of node.js app instance depends on how many CPU cores on the machine.





LAB 2-5

nginx load balancer

UPSTREAM (ROUND ROBIN)

Unix/Linux

/etc/nginx/sites-available/default

```
upstream backend_hosts {  
    server host1.example.com weight=3;  
    server host2.example.com;  
    server host3.example.com;  
}  
  
server {  
    listen 80;  
    server_name example.com;  
  
    location /app {  
        proxy_pass http://backend_hosts;  
    }  
}
```

UPSTREAM - LEAST_CONN

Unix/Linux

/etc/nginx/sites-available/default

```
upstream backend_hosts {
    least_conn;
    server host1.example.com;
    server host2.example.com;
    server host3.example.com;
}

server {
    listen 80;
    server_name example.com;

    location /app {
        proxy_pass http://backend_hosts;
    }
}
```

UPSTREAM - HASH

Unix/Linux

/etc/nginx/sites-available/default

```
upstream backend_hosts {
    hash $remote_addr$remote_port consistent;
    server host1.example.com;
    server host2.example.com;
    server host3.example.com;
}

server {
    listen 80;
    server_name example.com;

    location /app {
        proxy_pass http://backend_hosts;
    }
}
```

UPSTREAM - A/B TESTING

Unix/Linux

/etc/nginx/sites-available/default

```
upstream backend_hosts {
    hash $remote_addr$remote_port consistent;
    server host1.example.com;
    server host2.example.com;
    server host3.example.com;
}

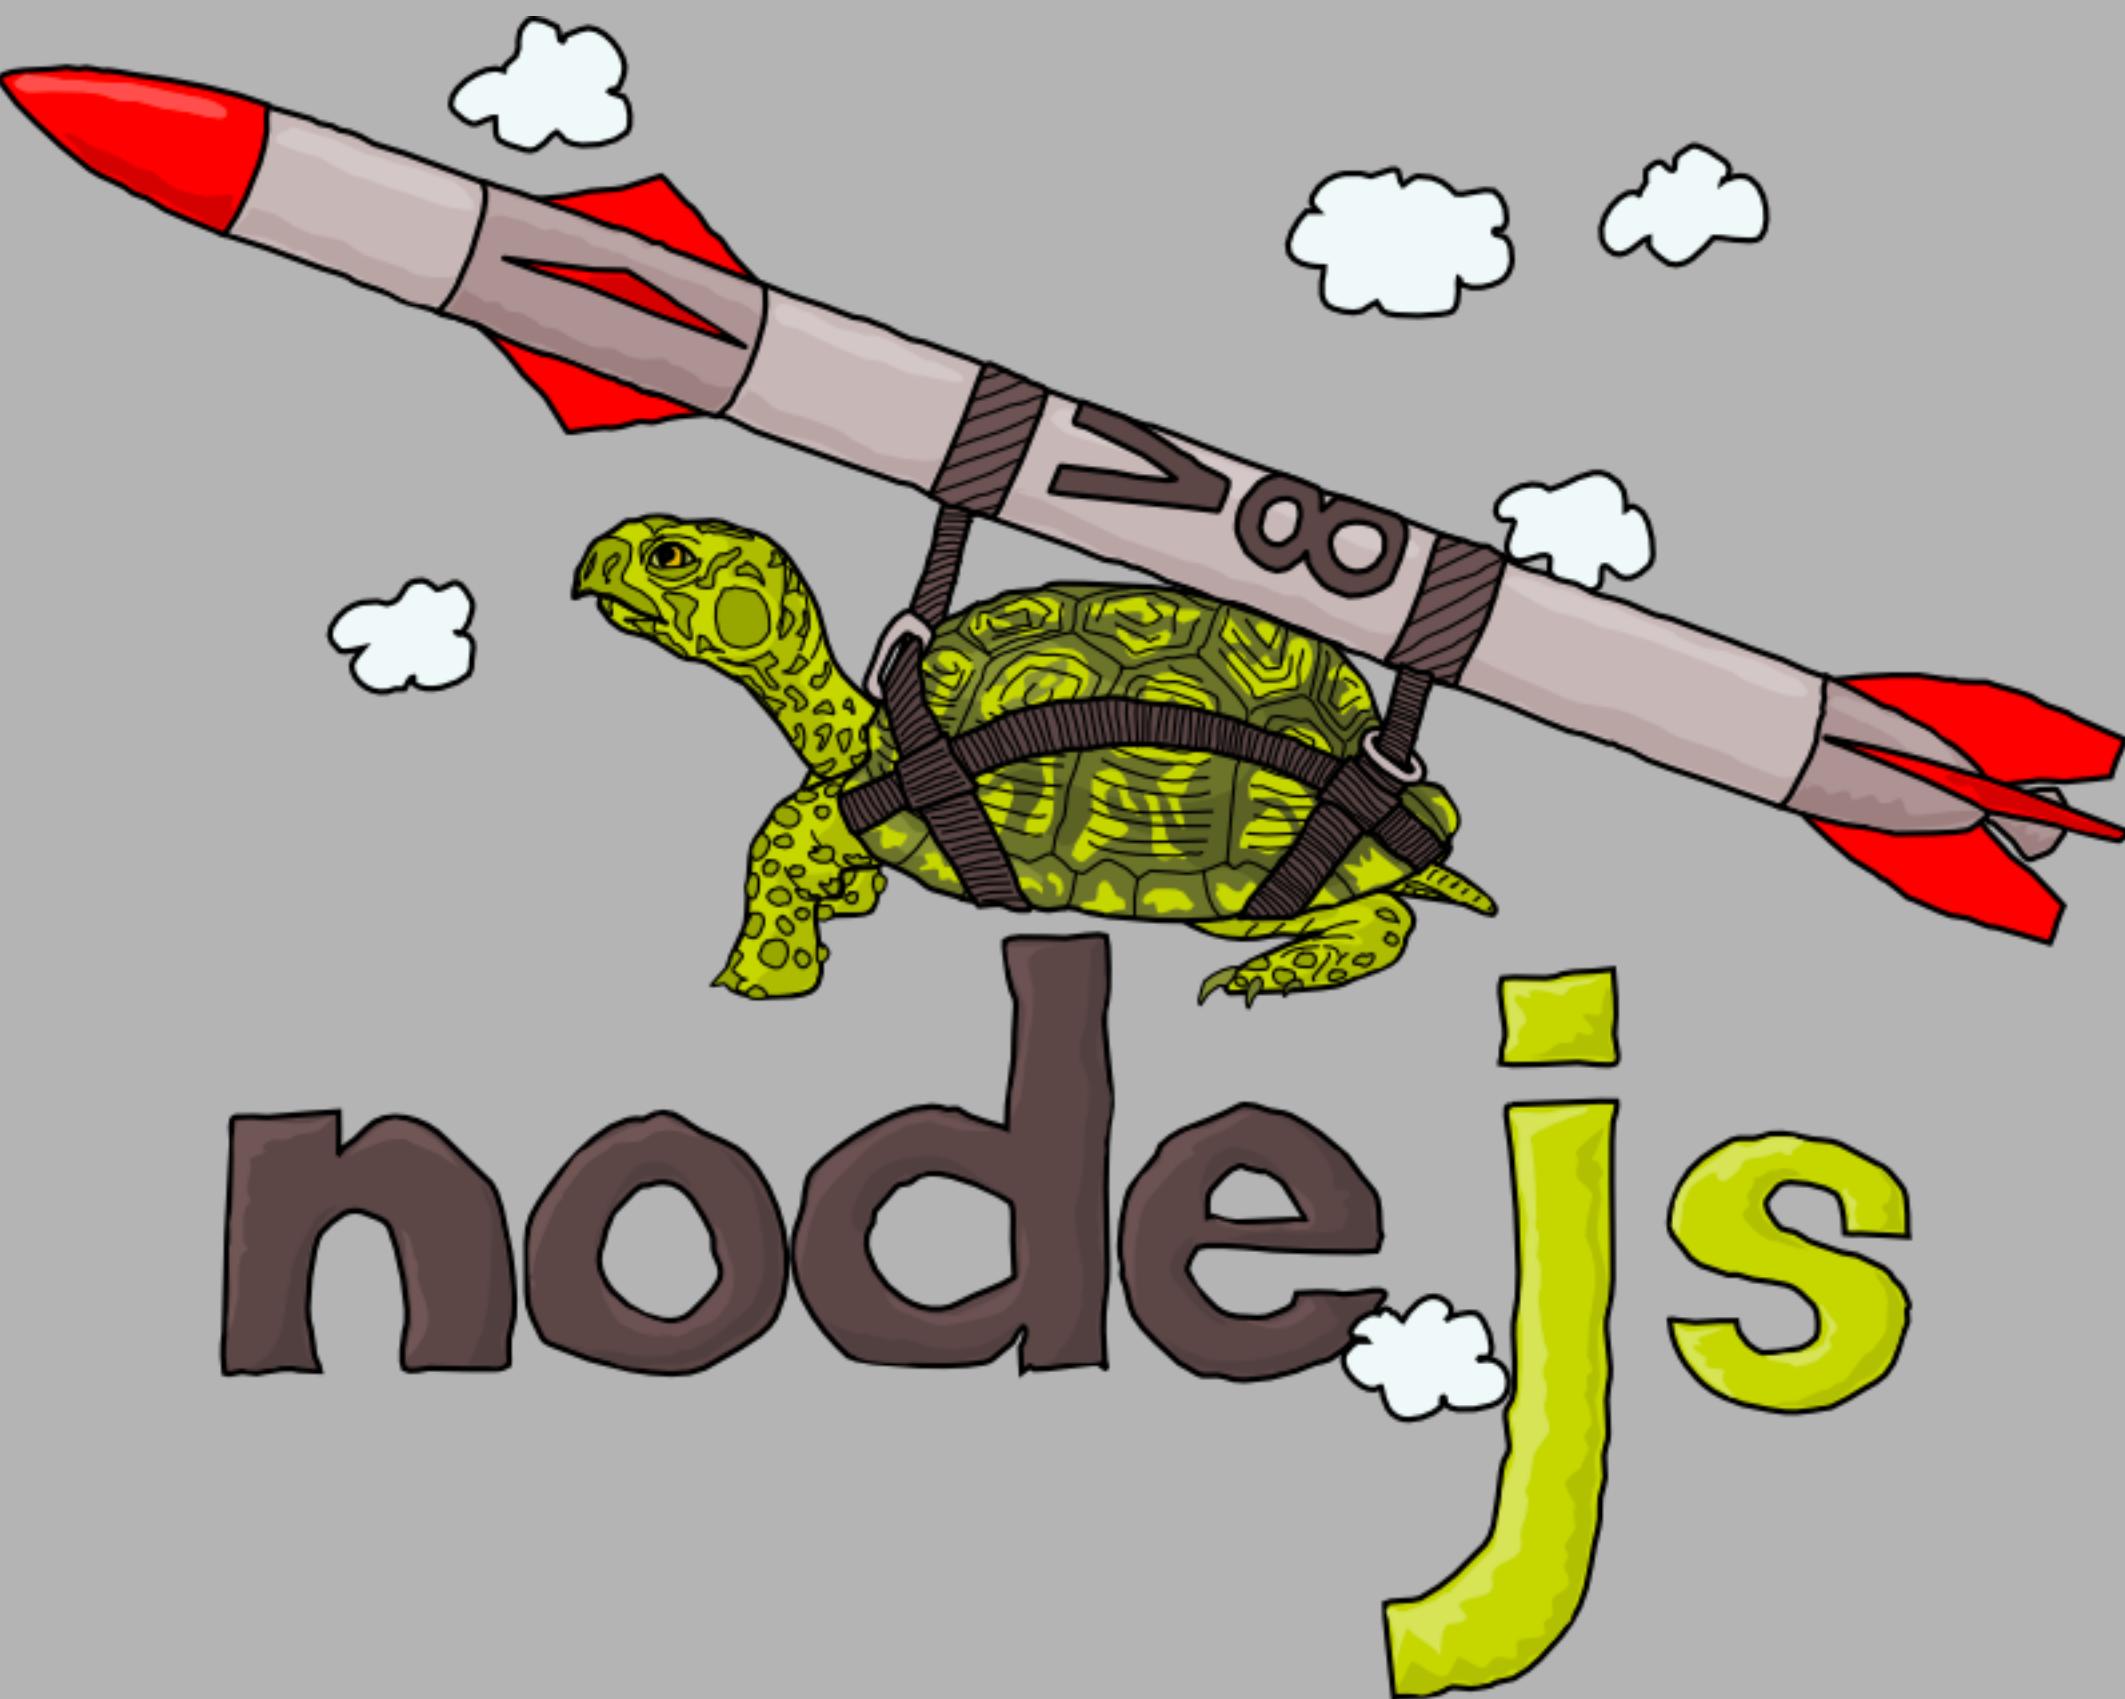
map $cookie_group $group {
    ~(?P<value>.+) $value;
    default      backendB; # The default upstream group
}

server {
    listen 80;

    location / {
        add_header Set-Cookie "group=$group; path=/"
        proxy_pass http://$group;
    }
}
```

MORE ADVANCED SCENARIOS

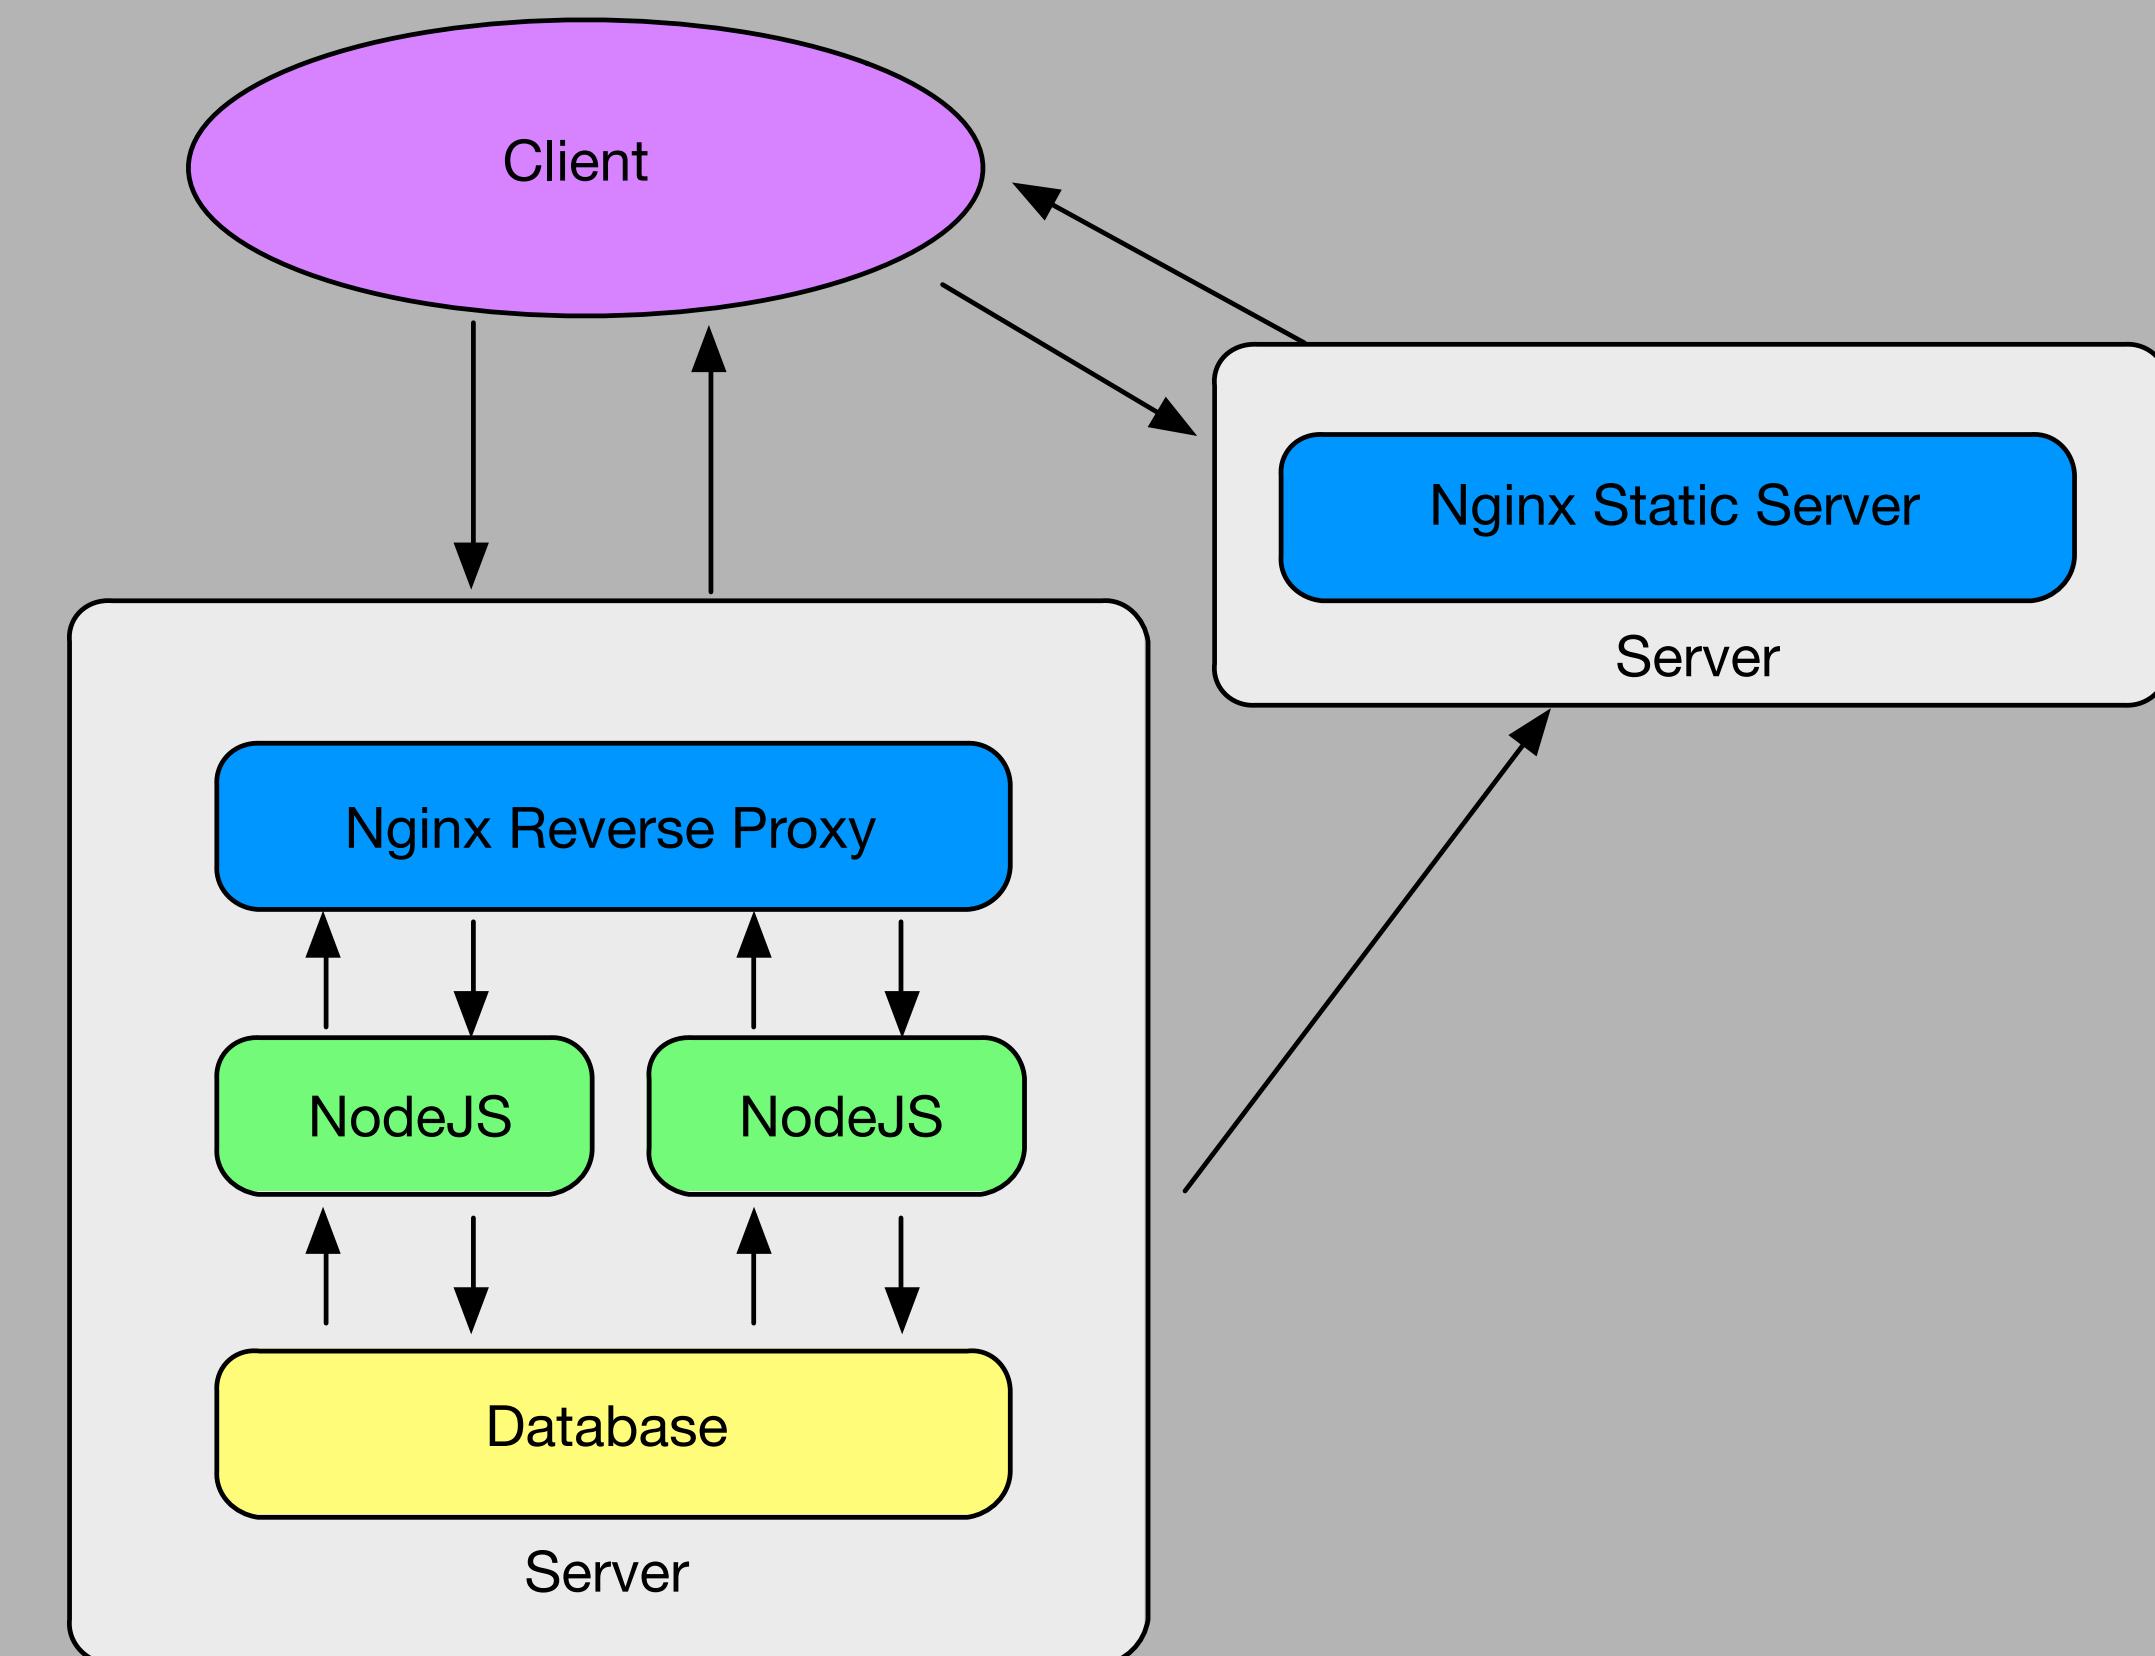
for further experiment



SCALING NODEJS

Use nginx as proxy server as well to load balance requests.

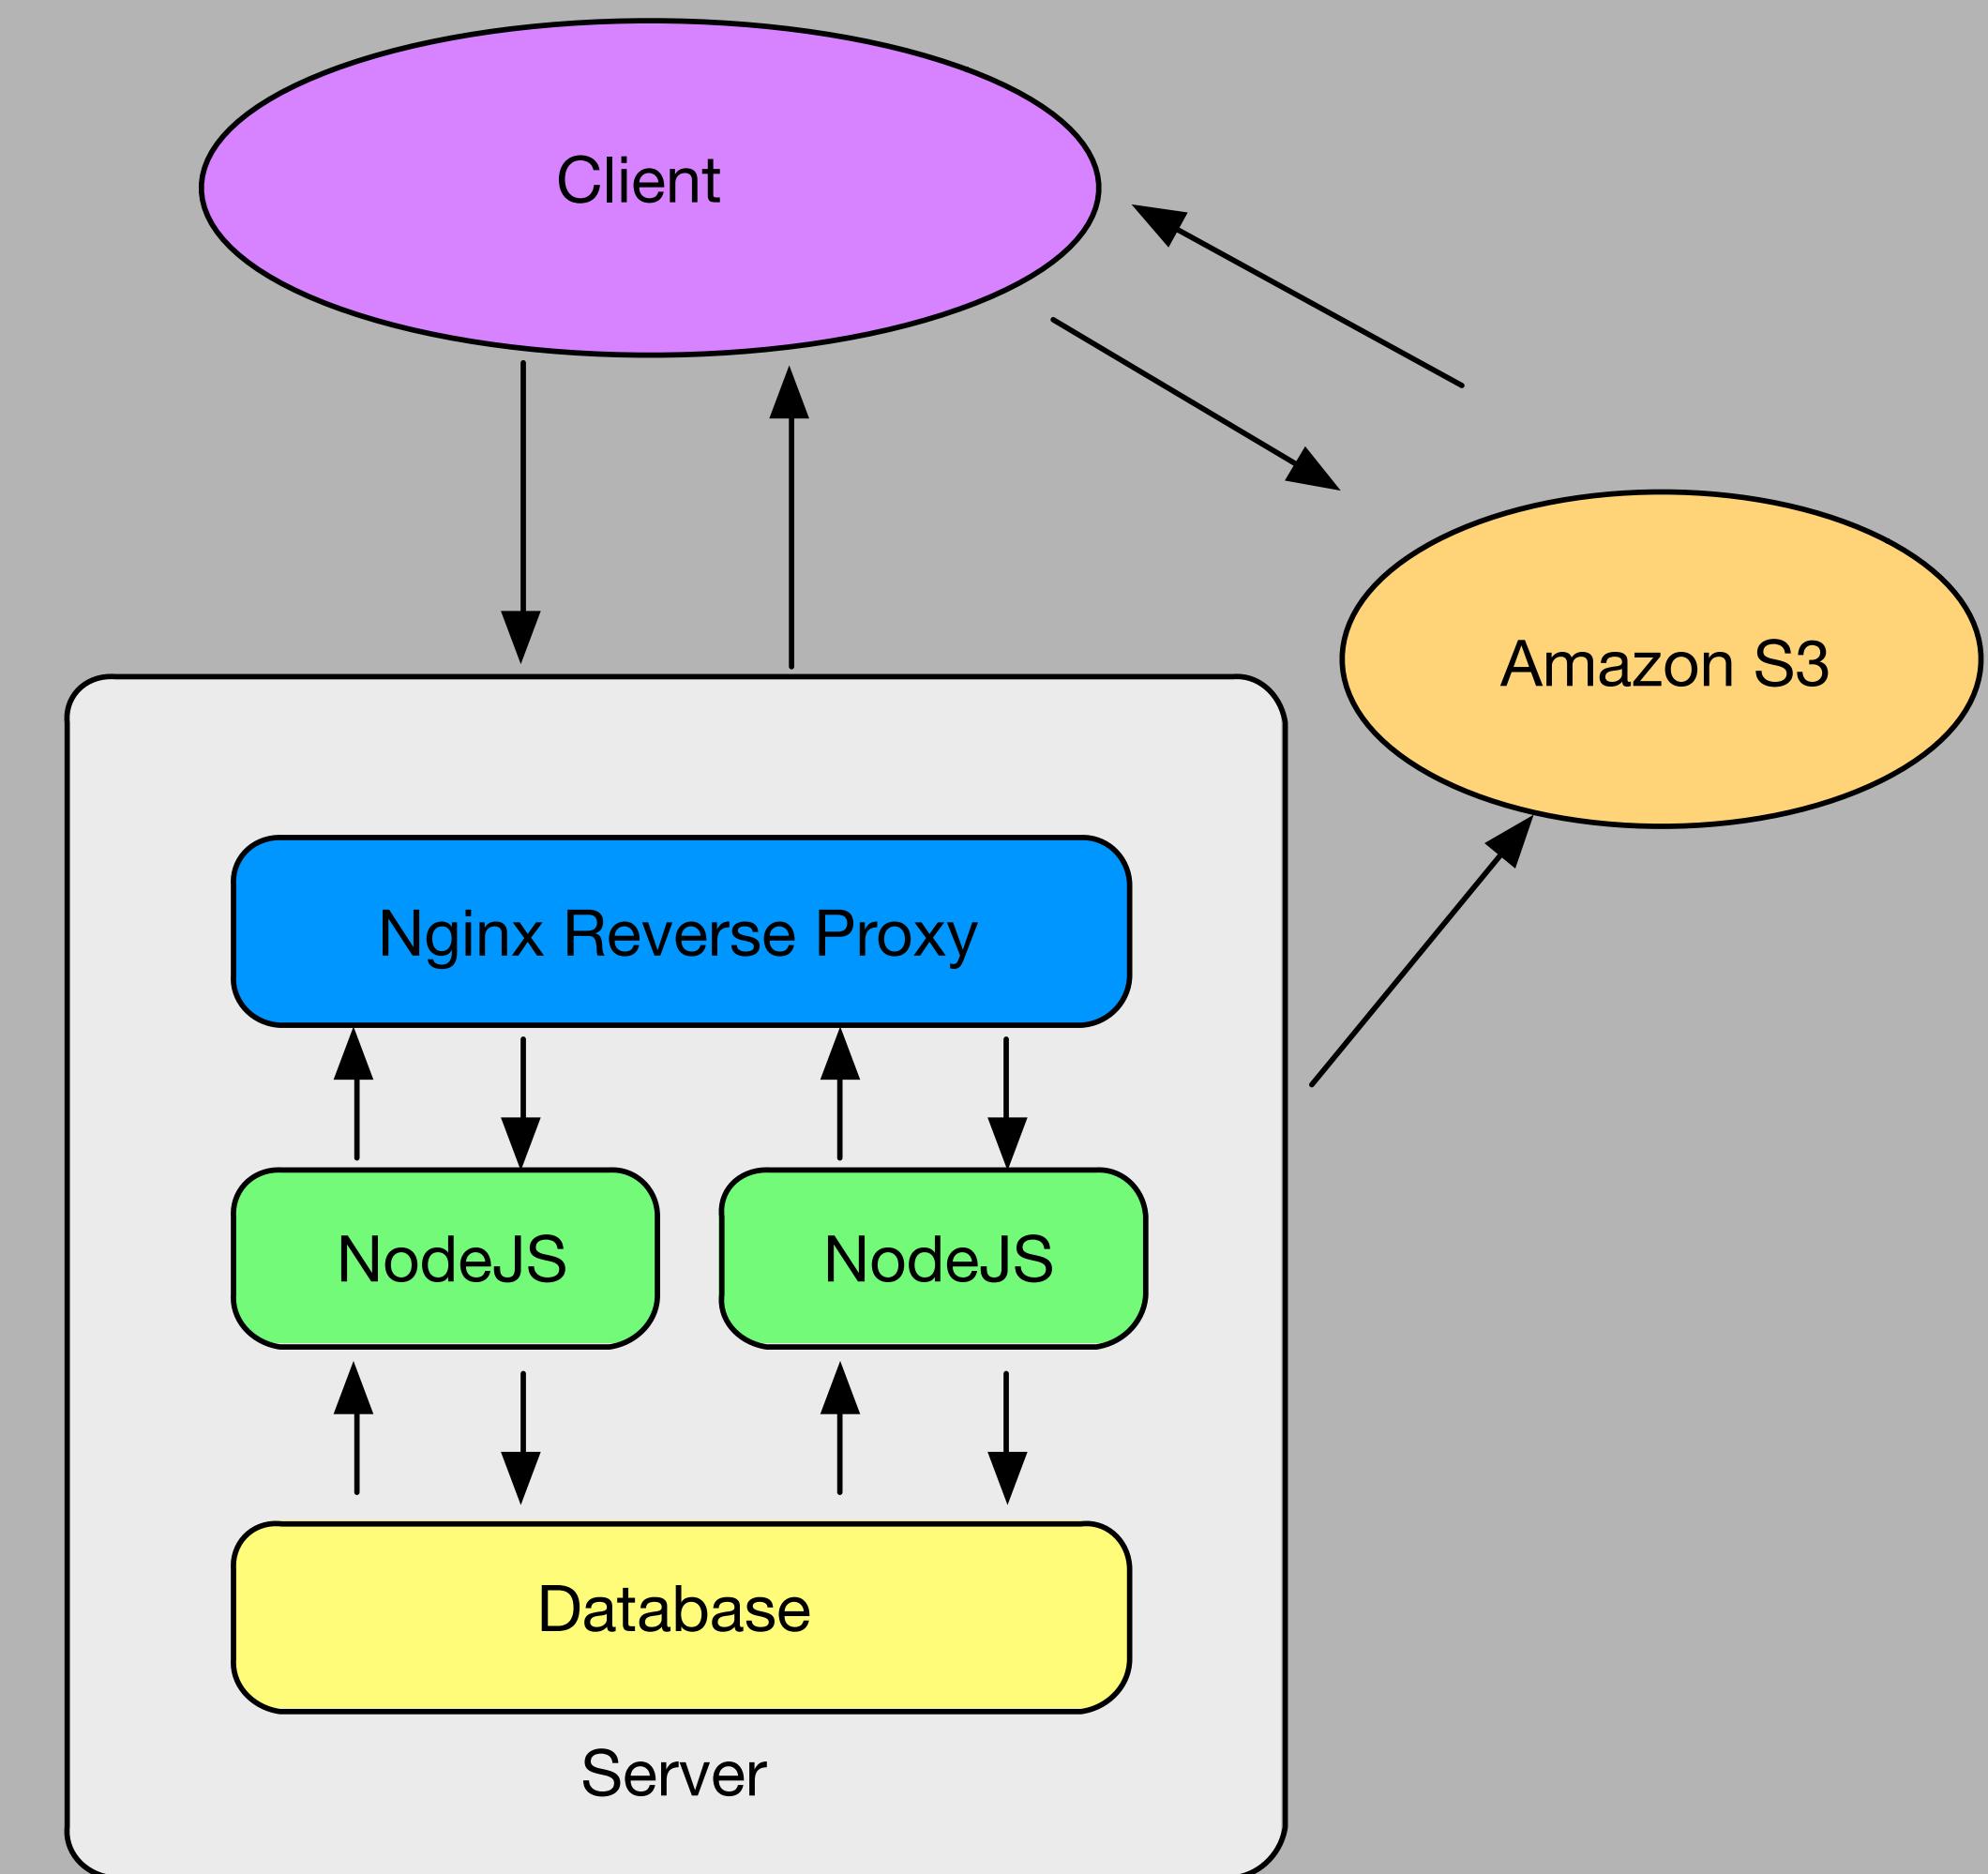
The number of node.js app instance depends on how many CPU cores on the machine.



SCALING NODEJS

Use nginx as proxy server as well to load balance requests.

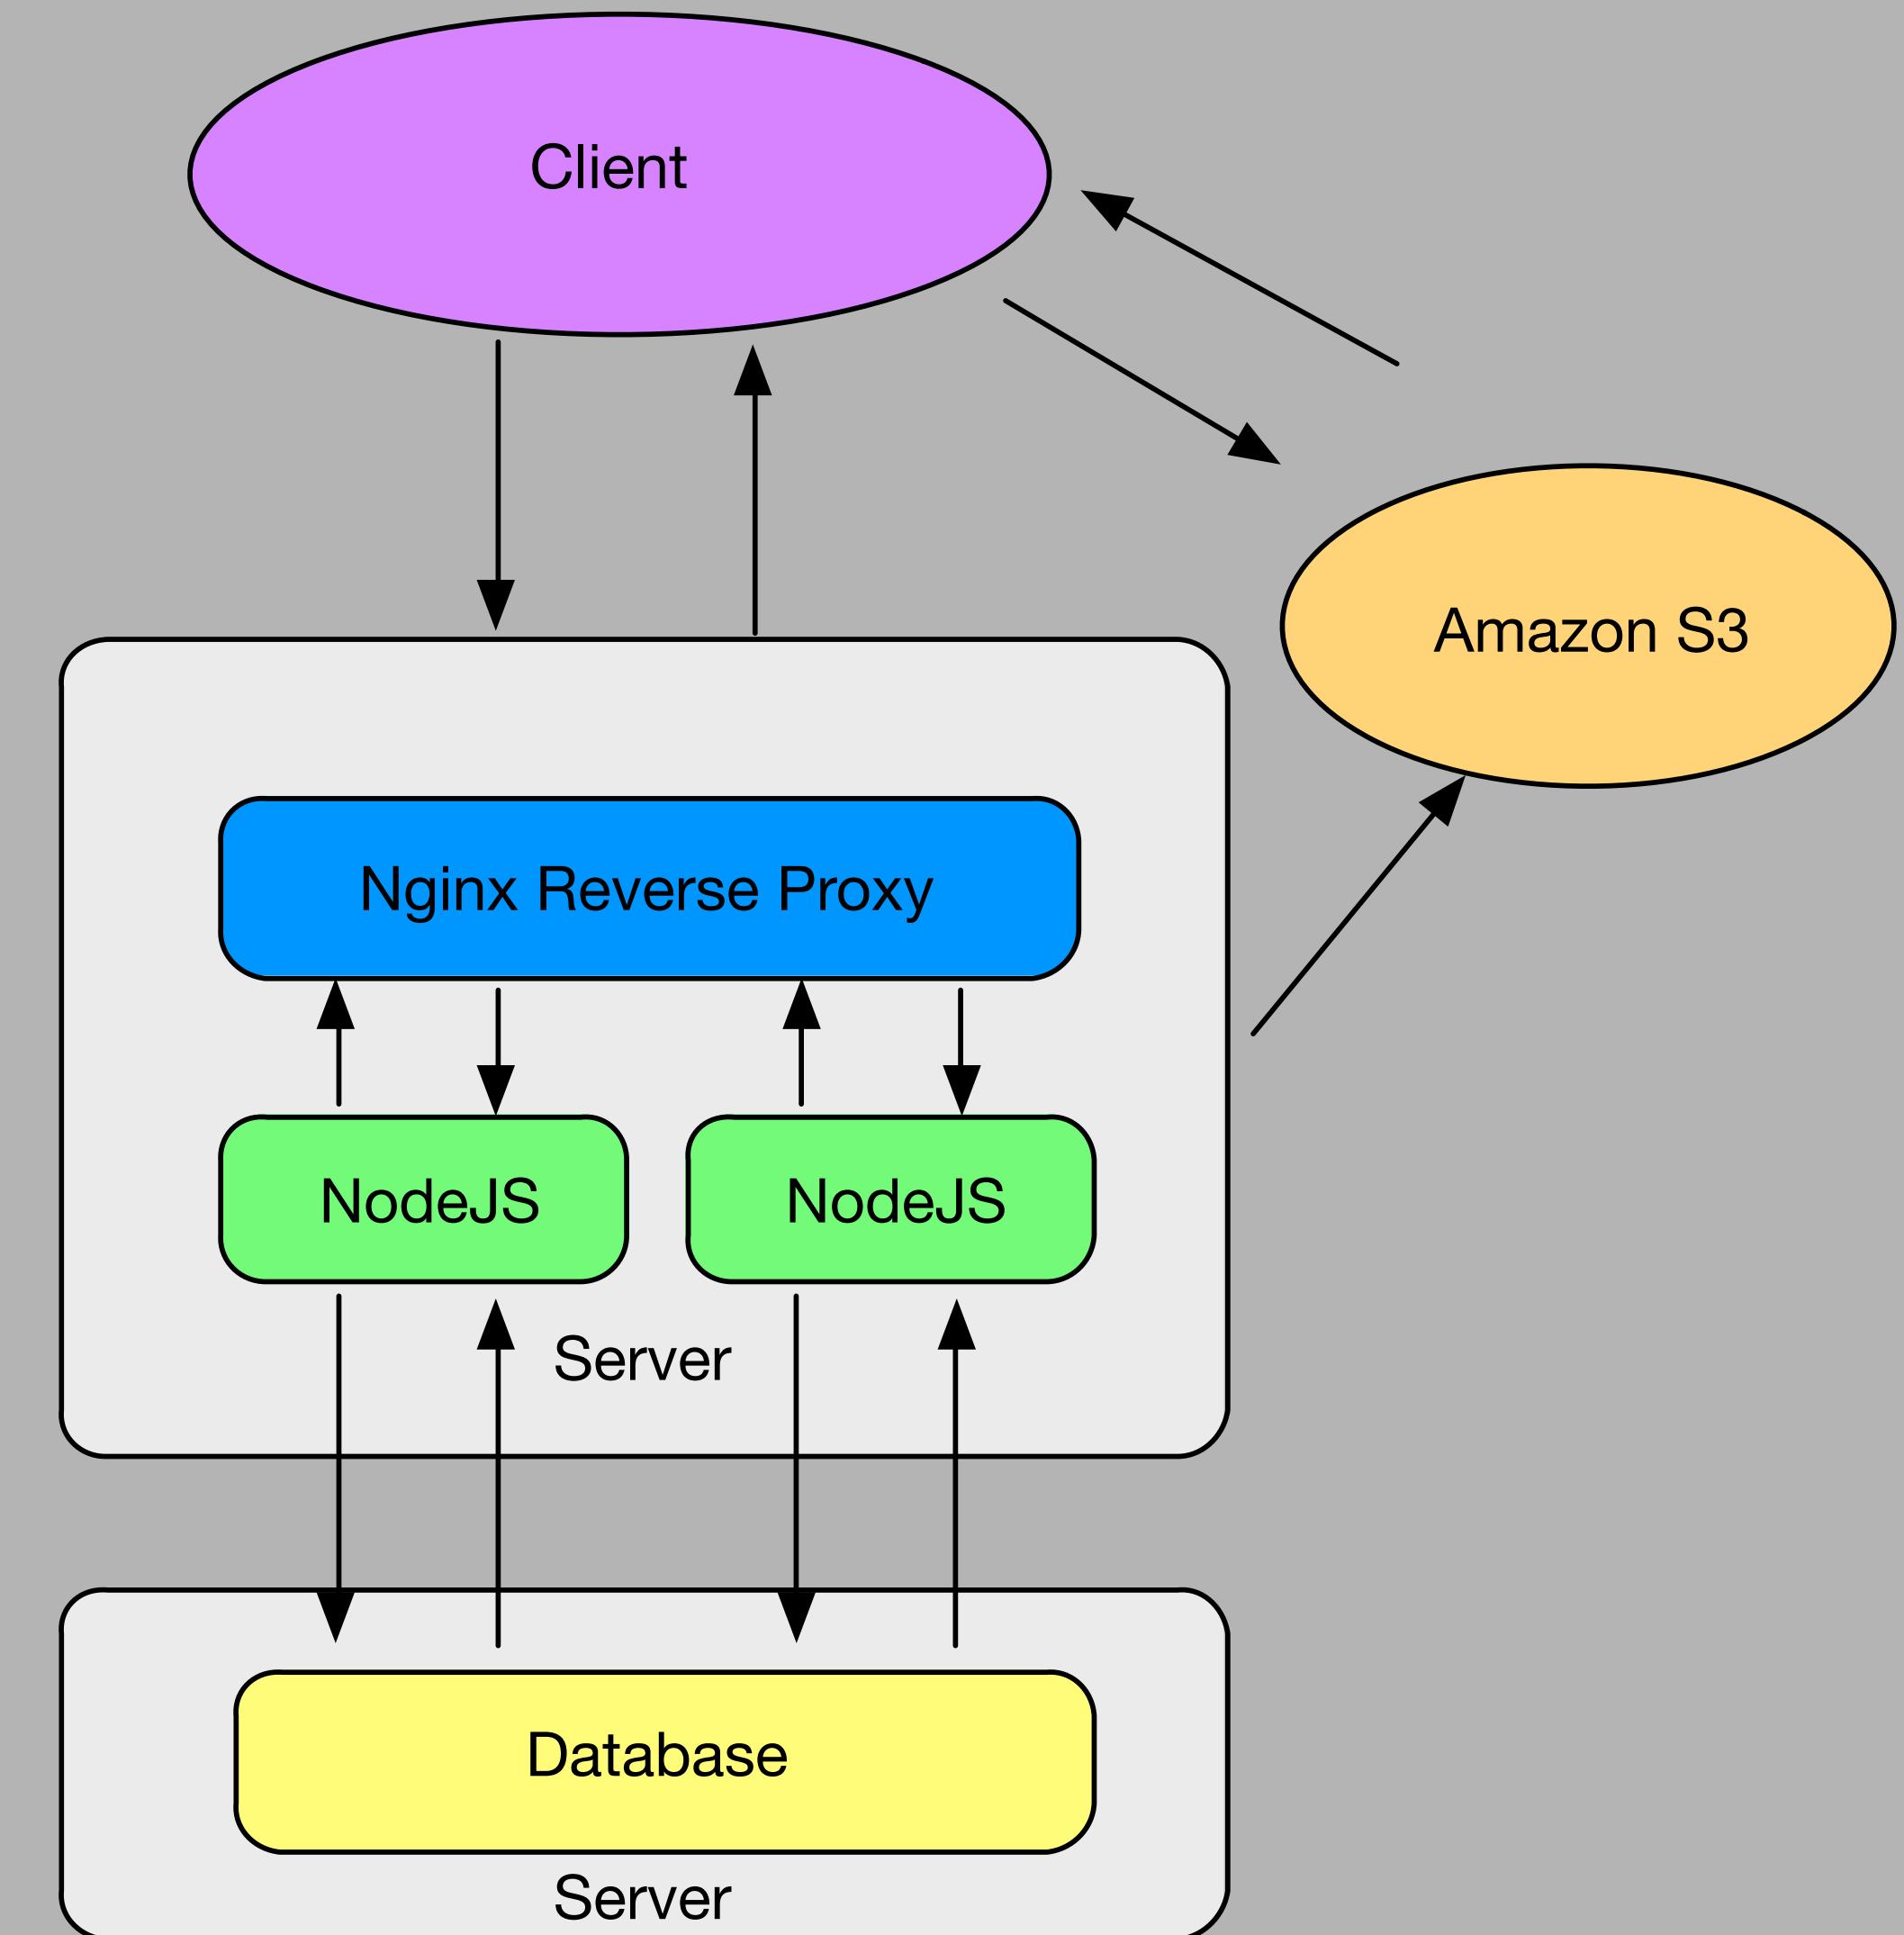
The number of node.js app instance depends on how many CPU cores on the machine.



SCALING NODEJS

Use nginx as proxy server as well to load balance requests.

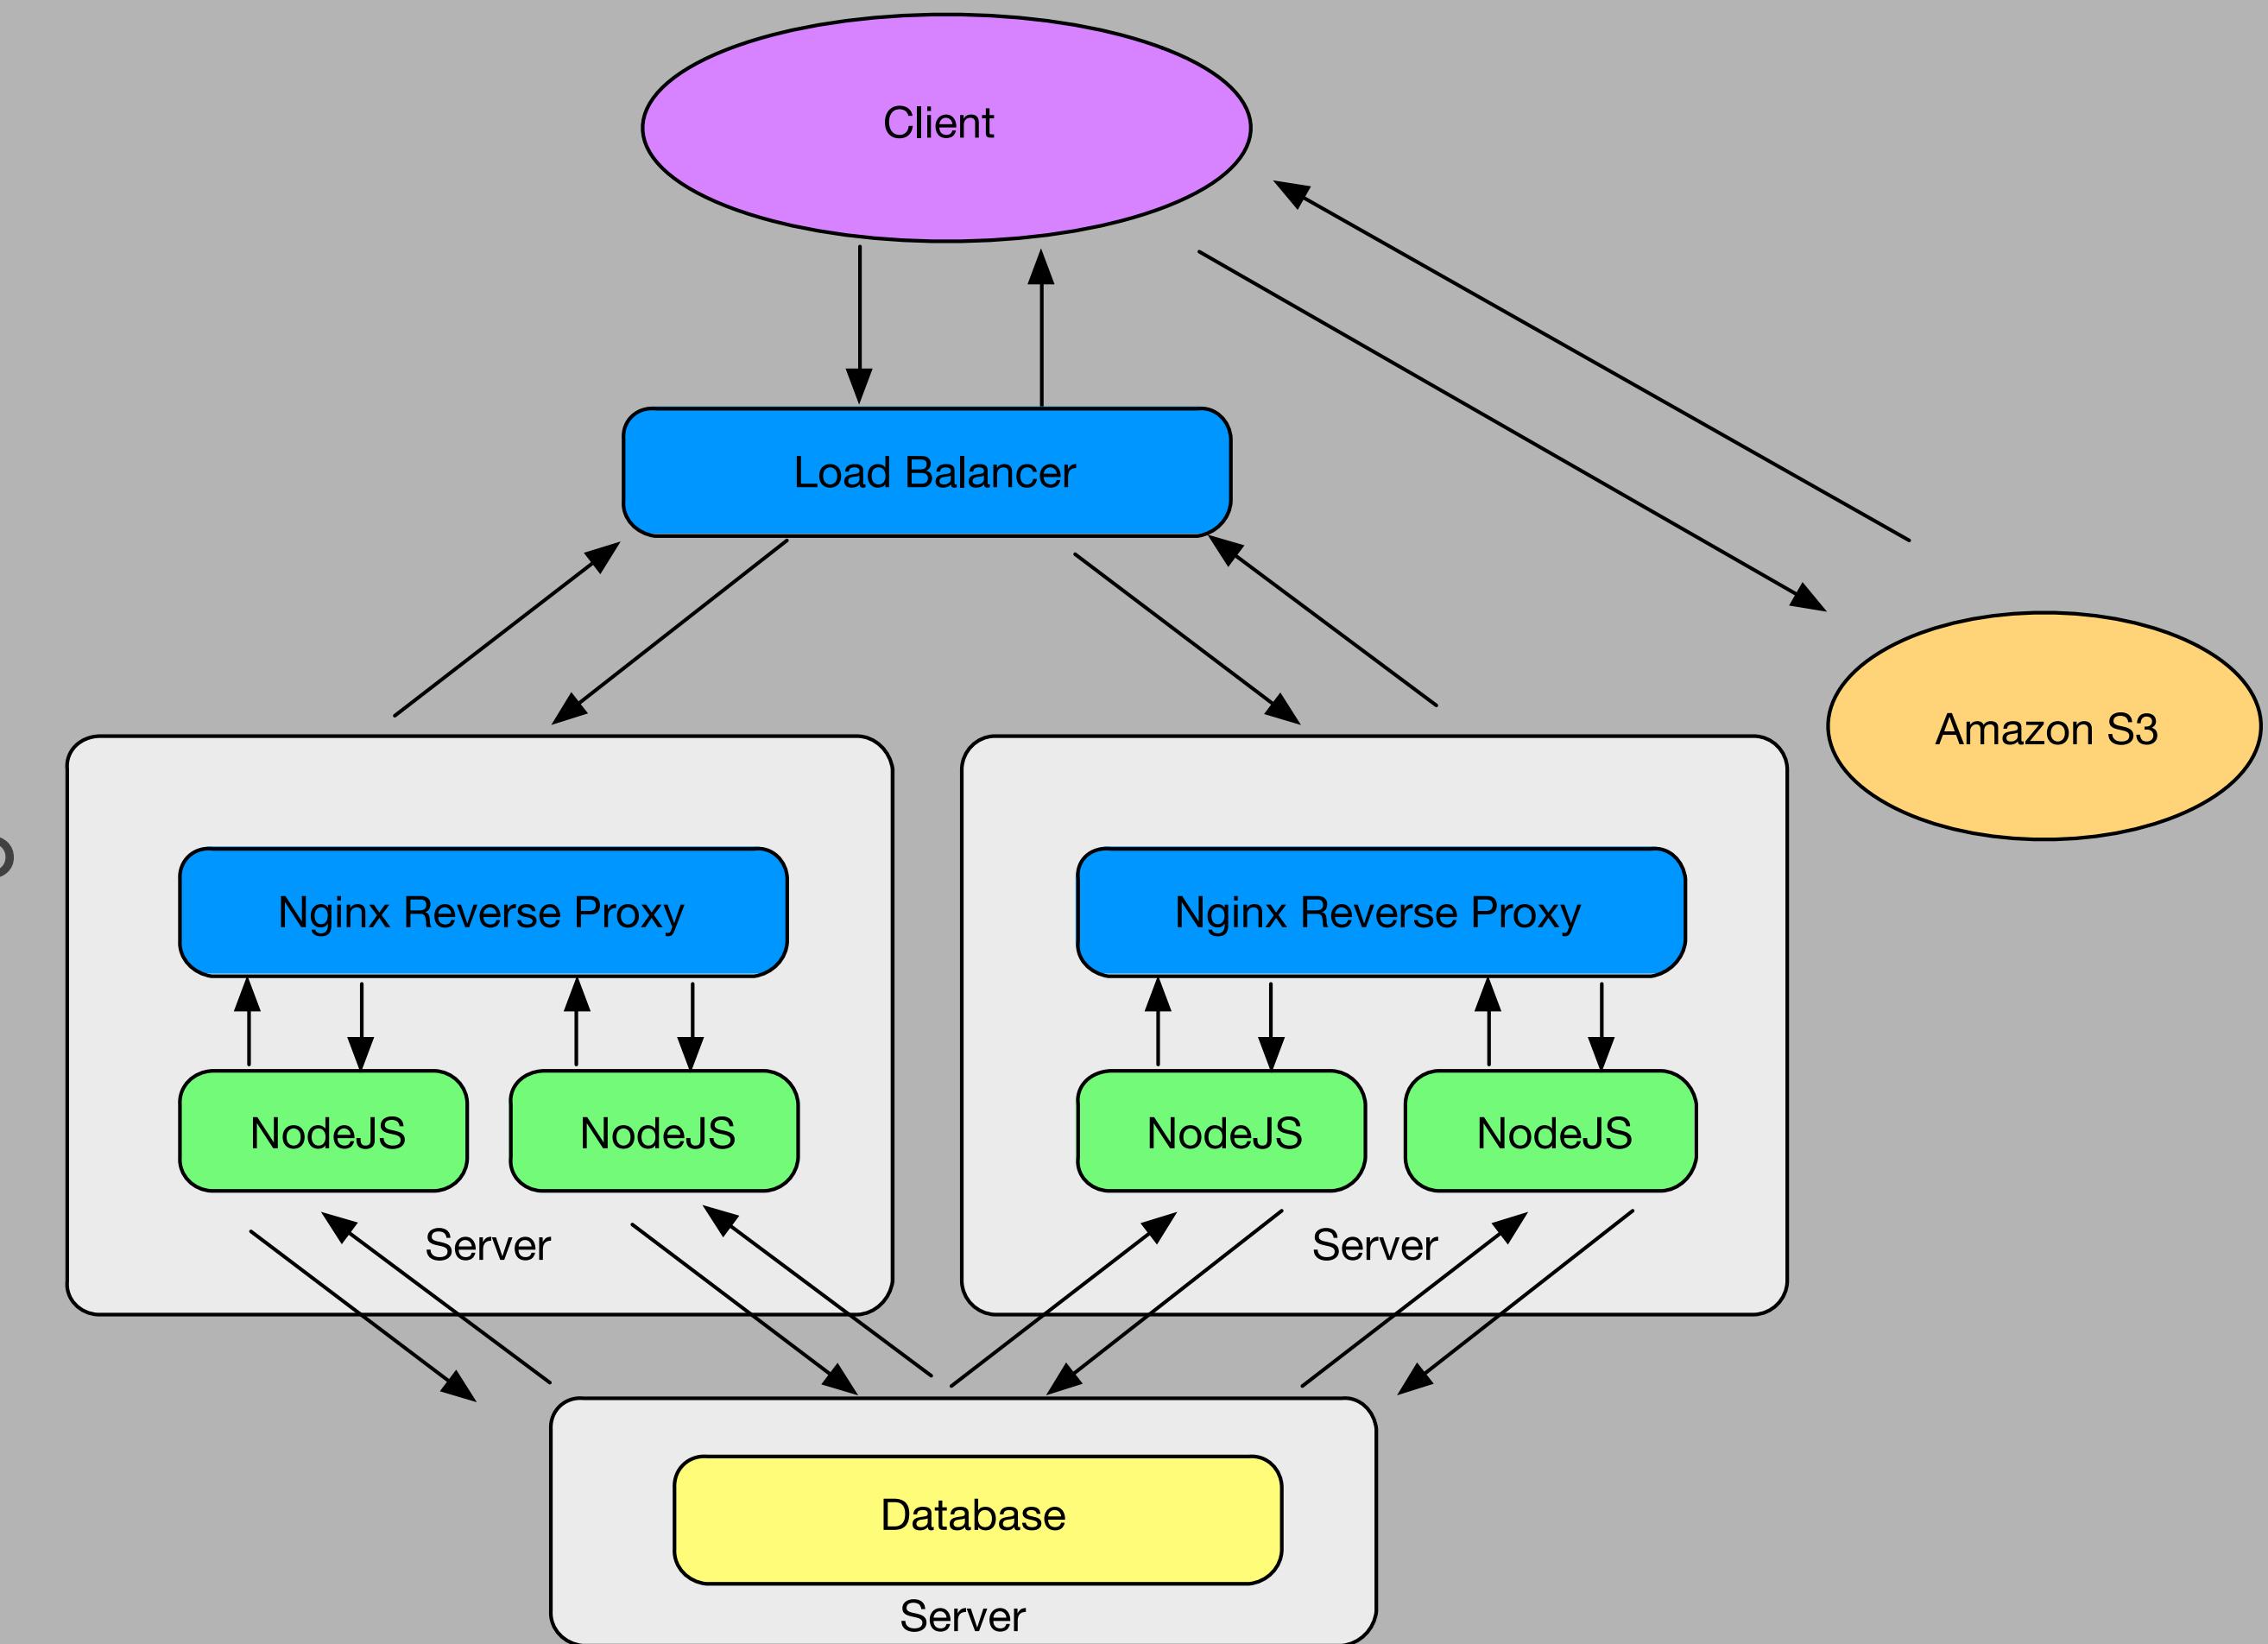
The number of node.js app instance depends on how many CPU cores on the machine.



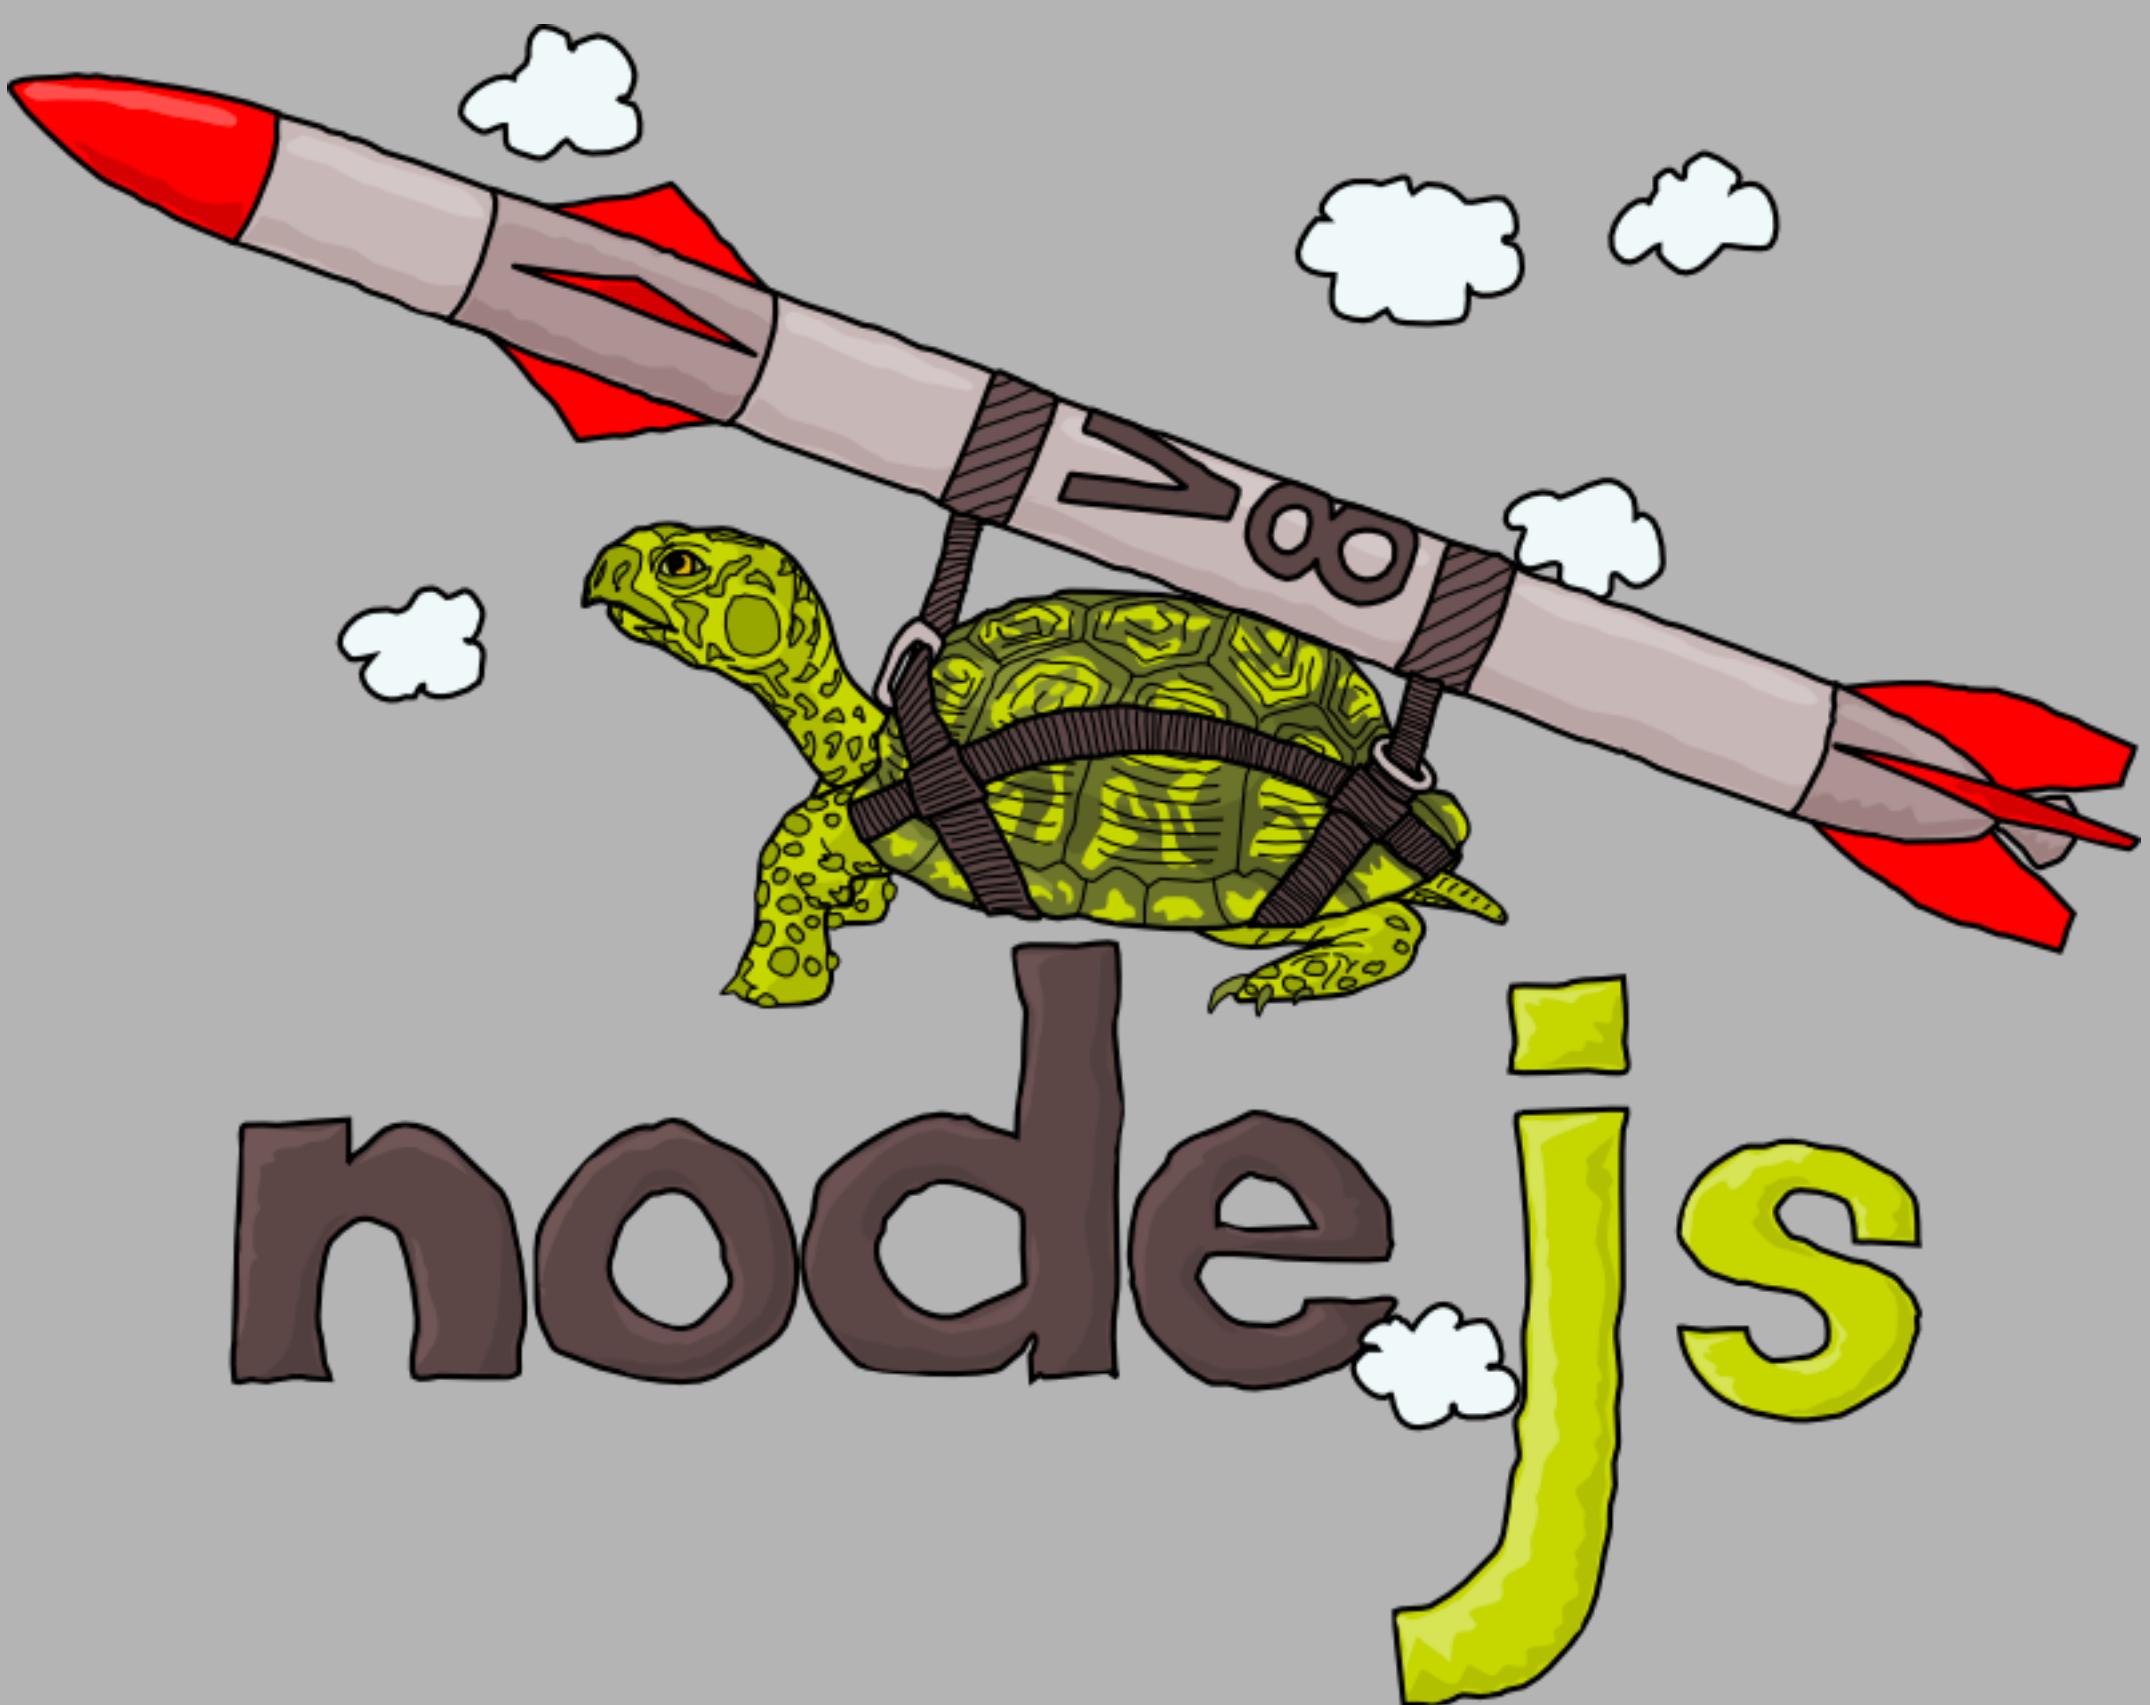
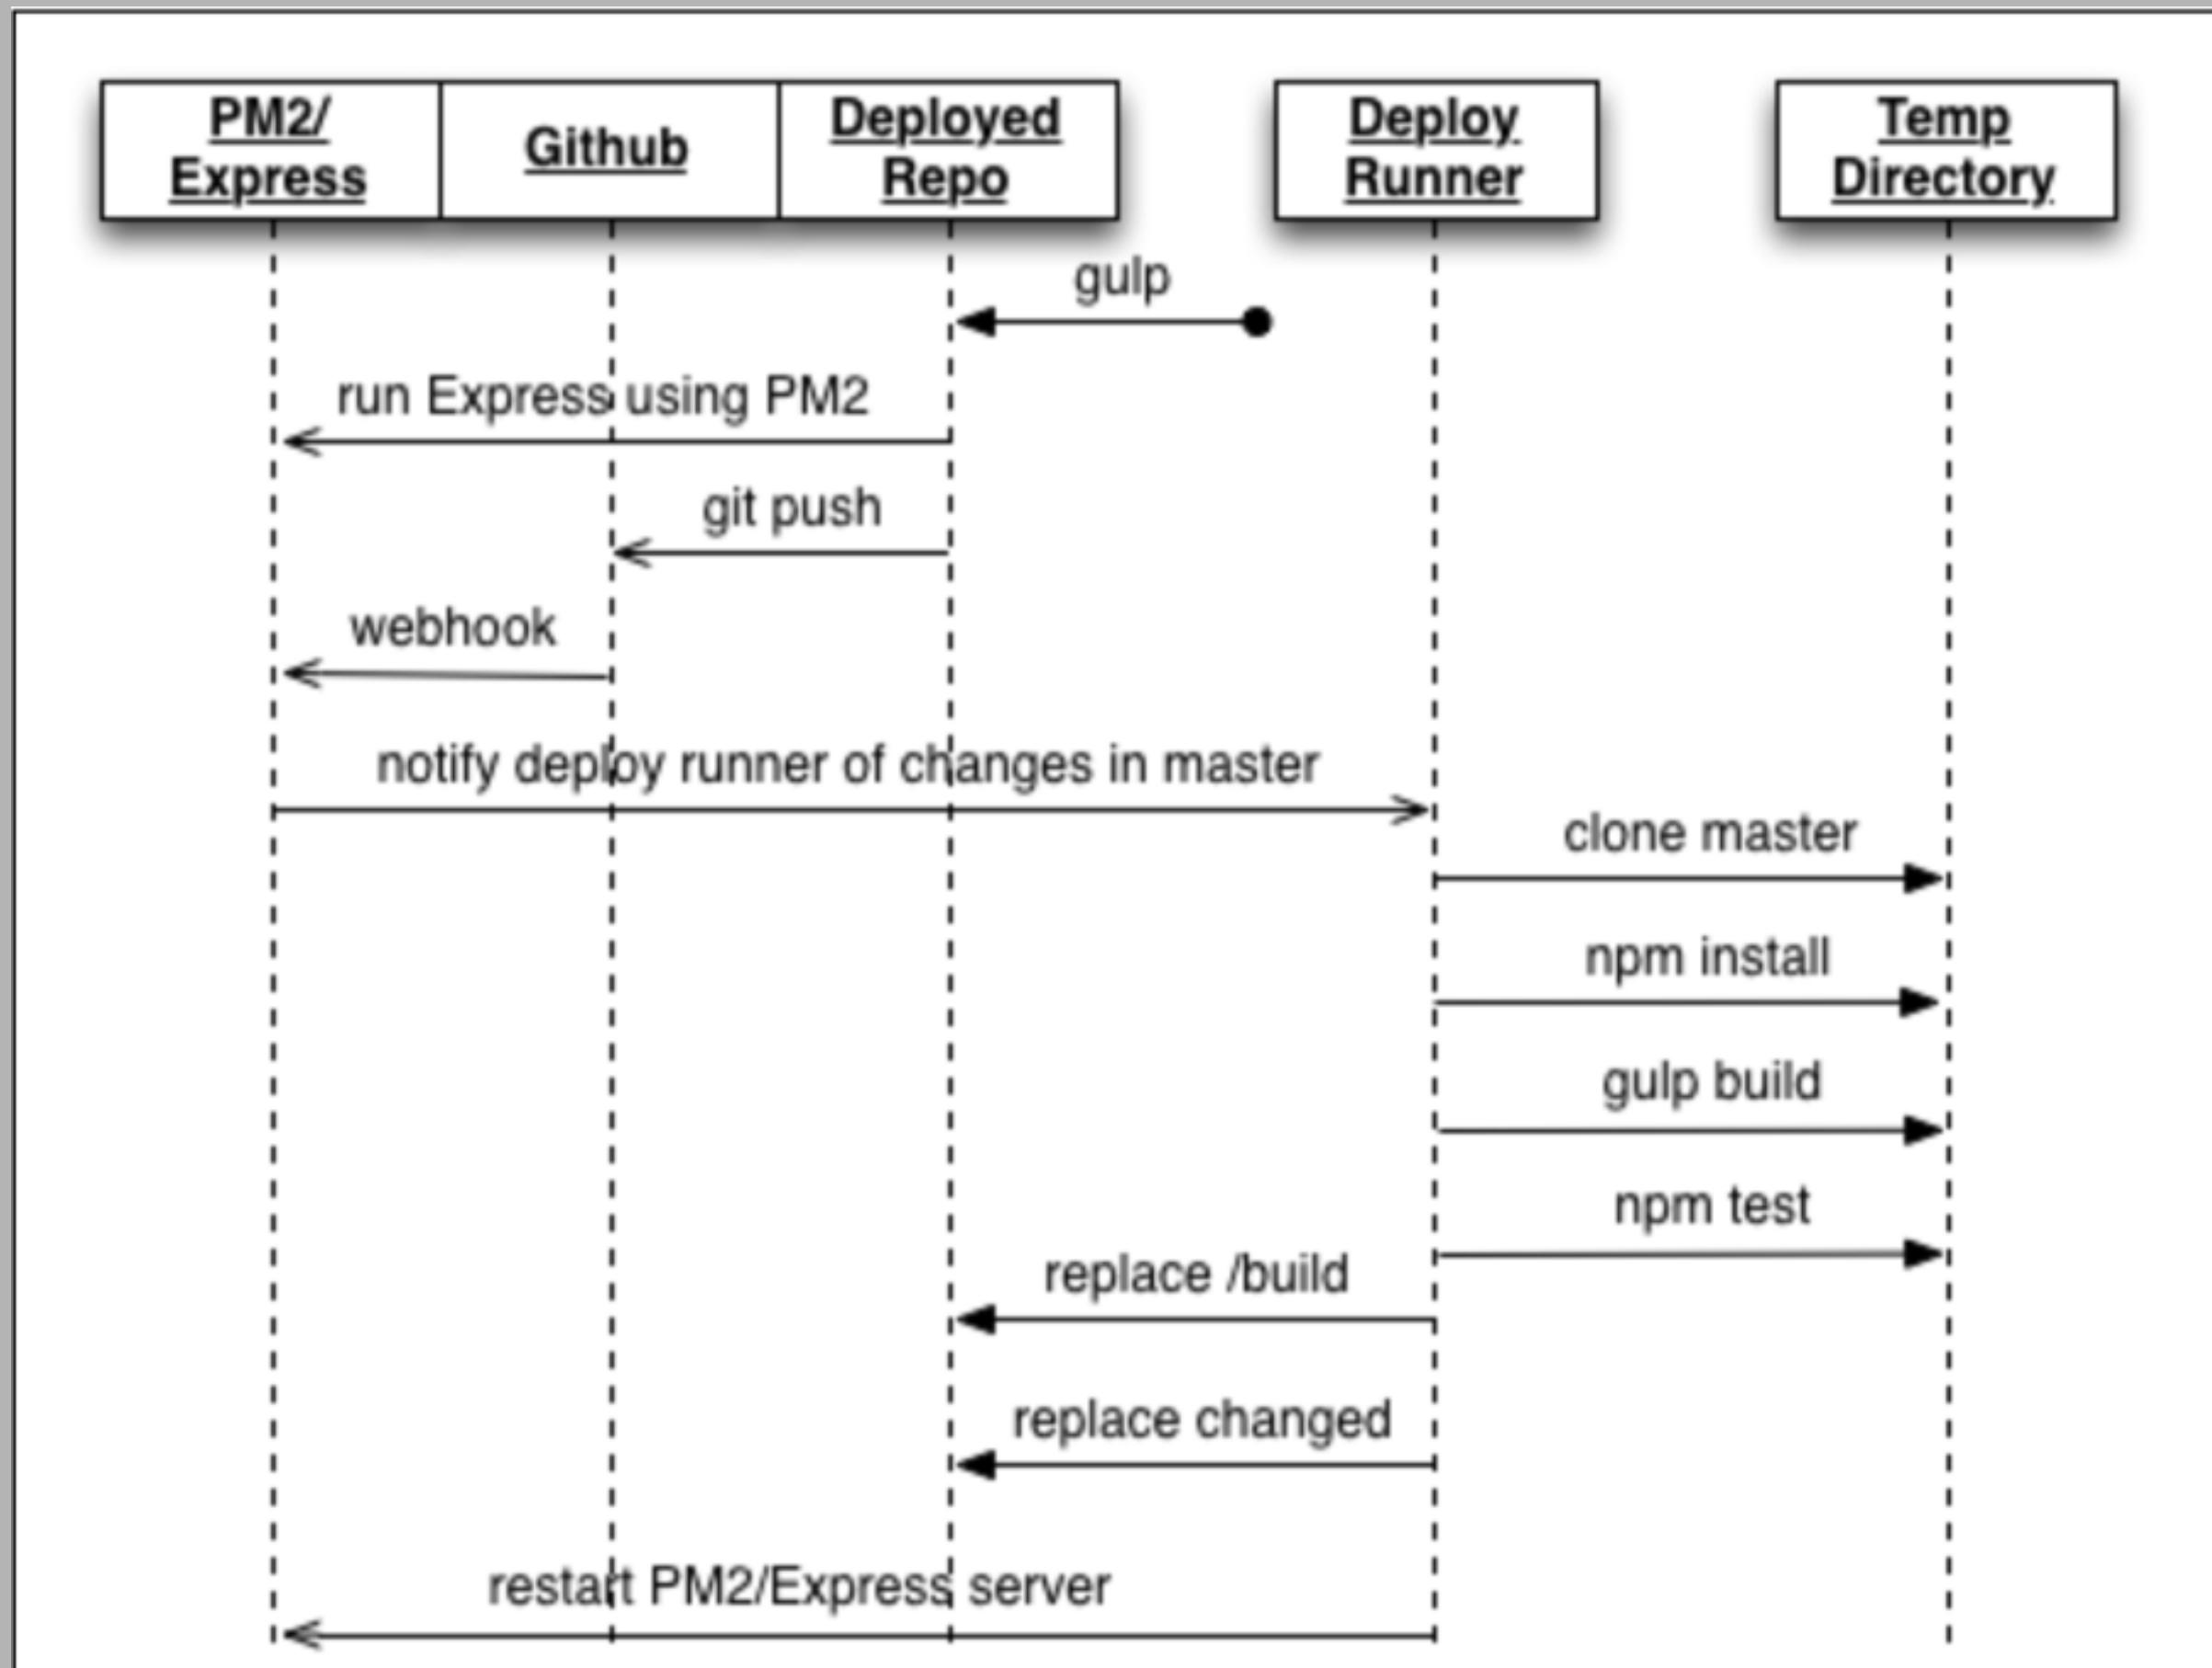
SCALING NODEJS

Use nginx as proxy server as well to load balance requests.

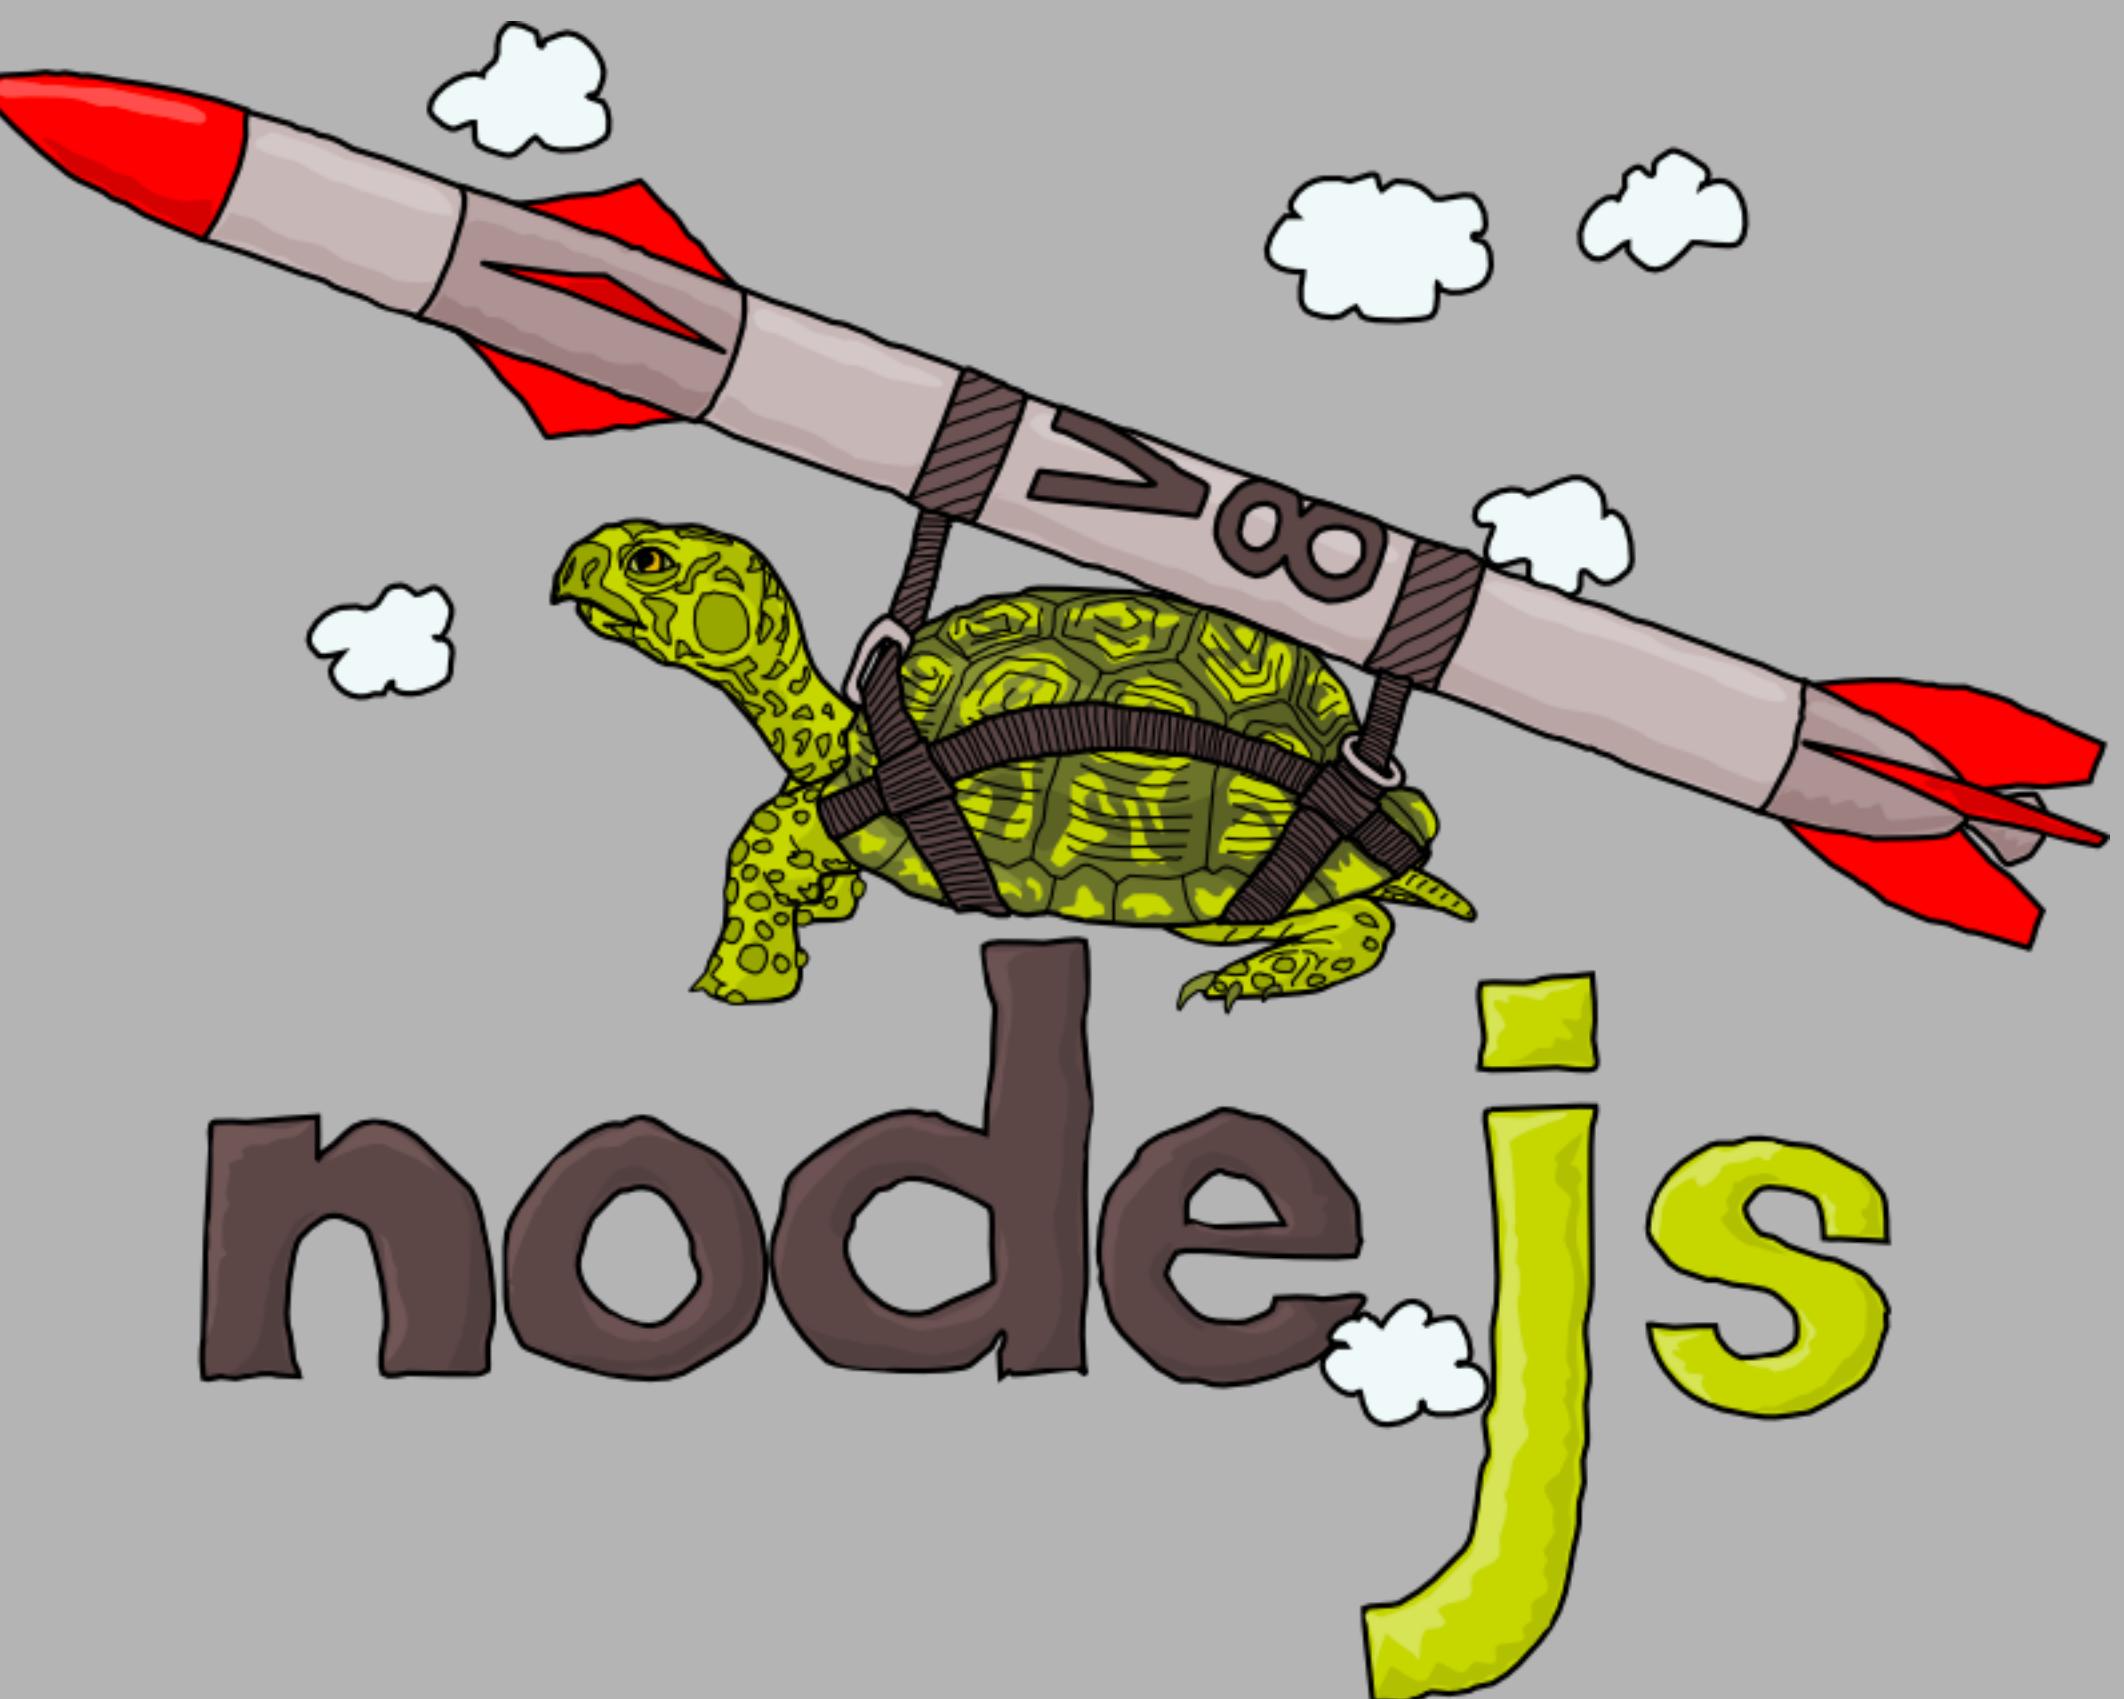
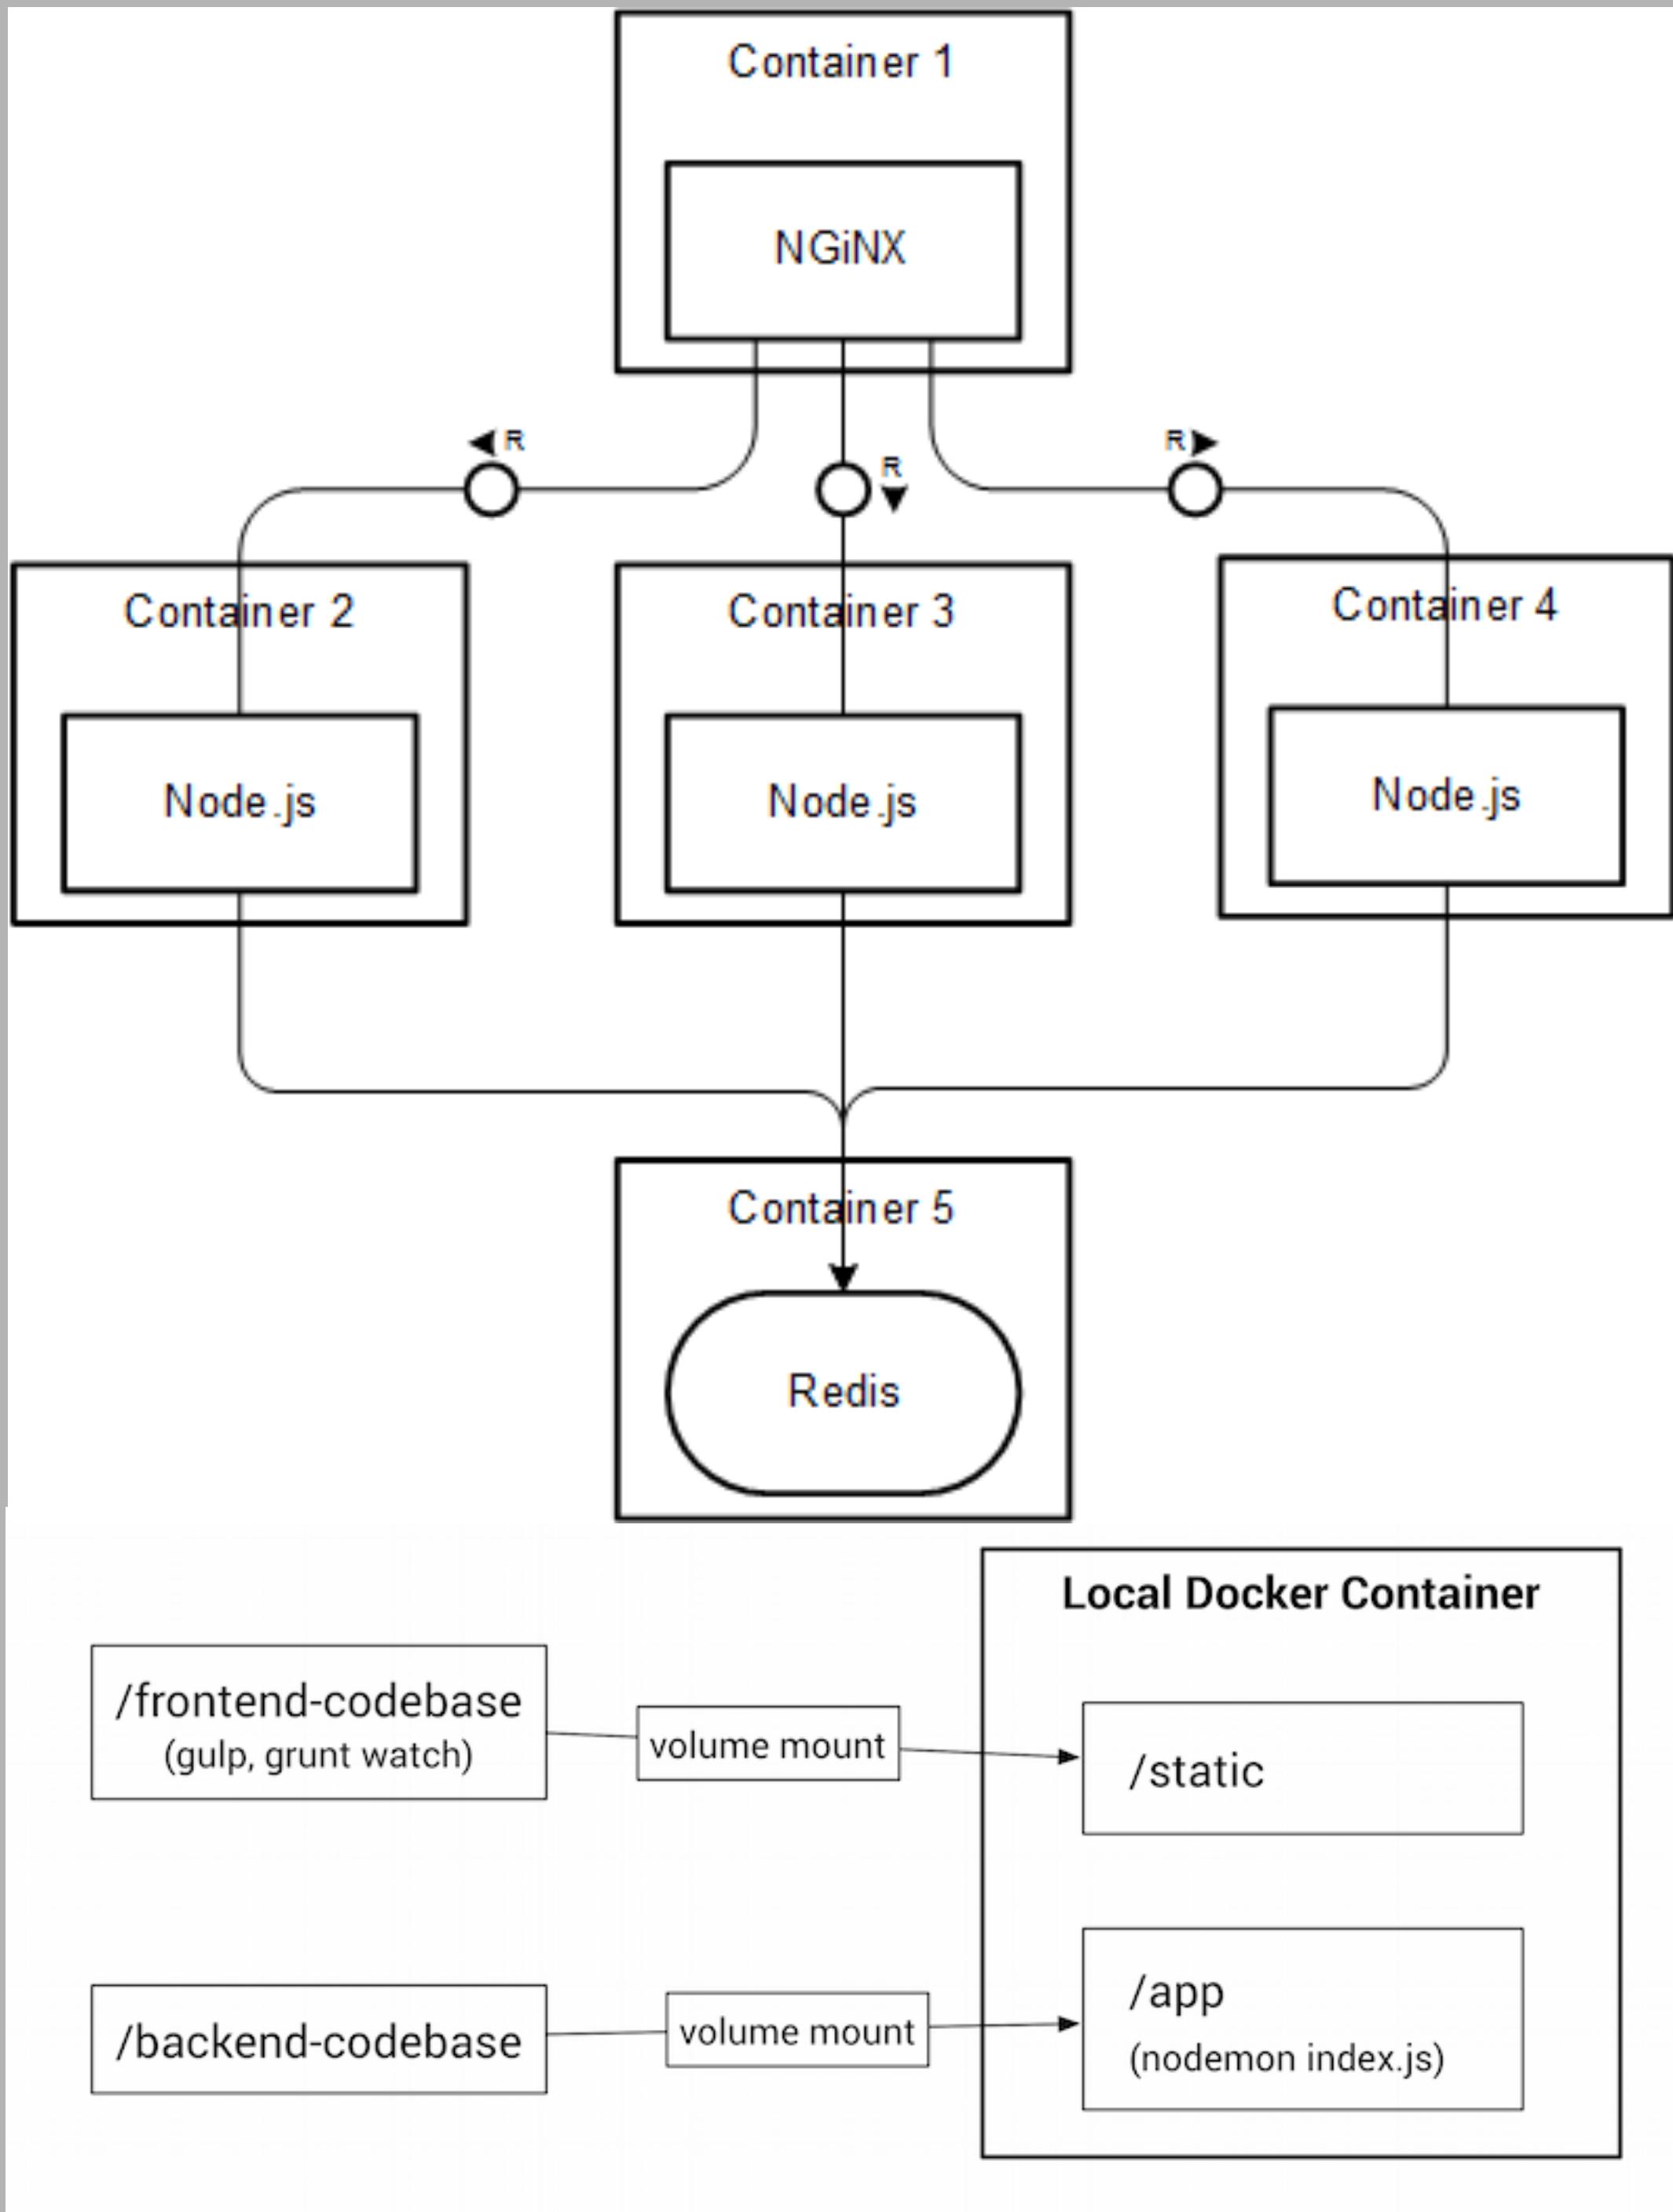
The number of node.js app instance depends on how many CPU cores on the machine.



PUSH/DEPLOY WITH GIT



DOCKER NODEJS





PERFORMANCE TESTING



WHAT IS PERFORMANCE TESTING

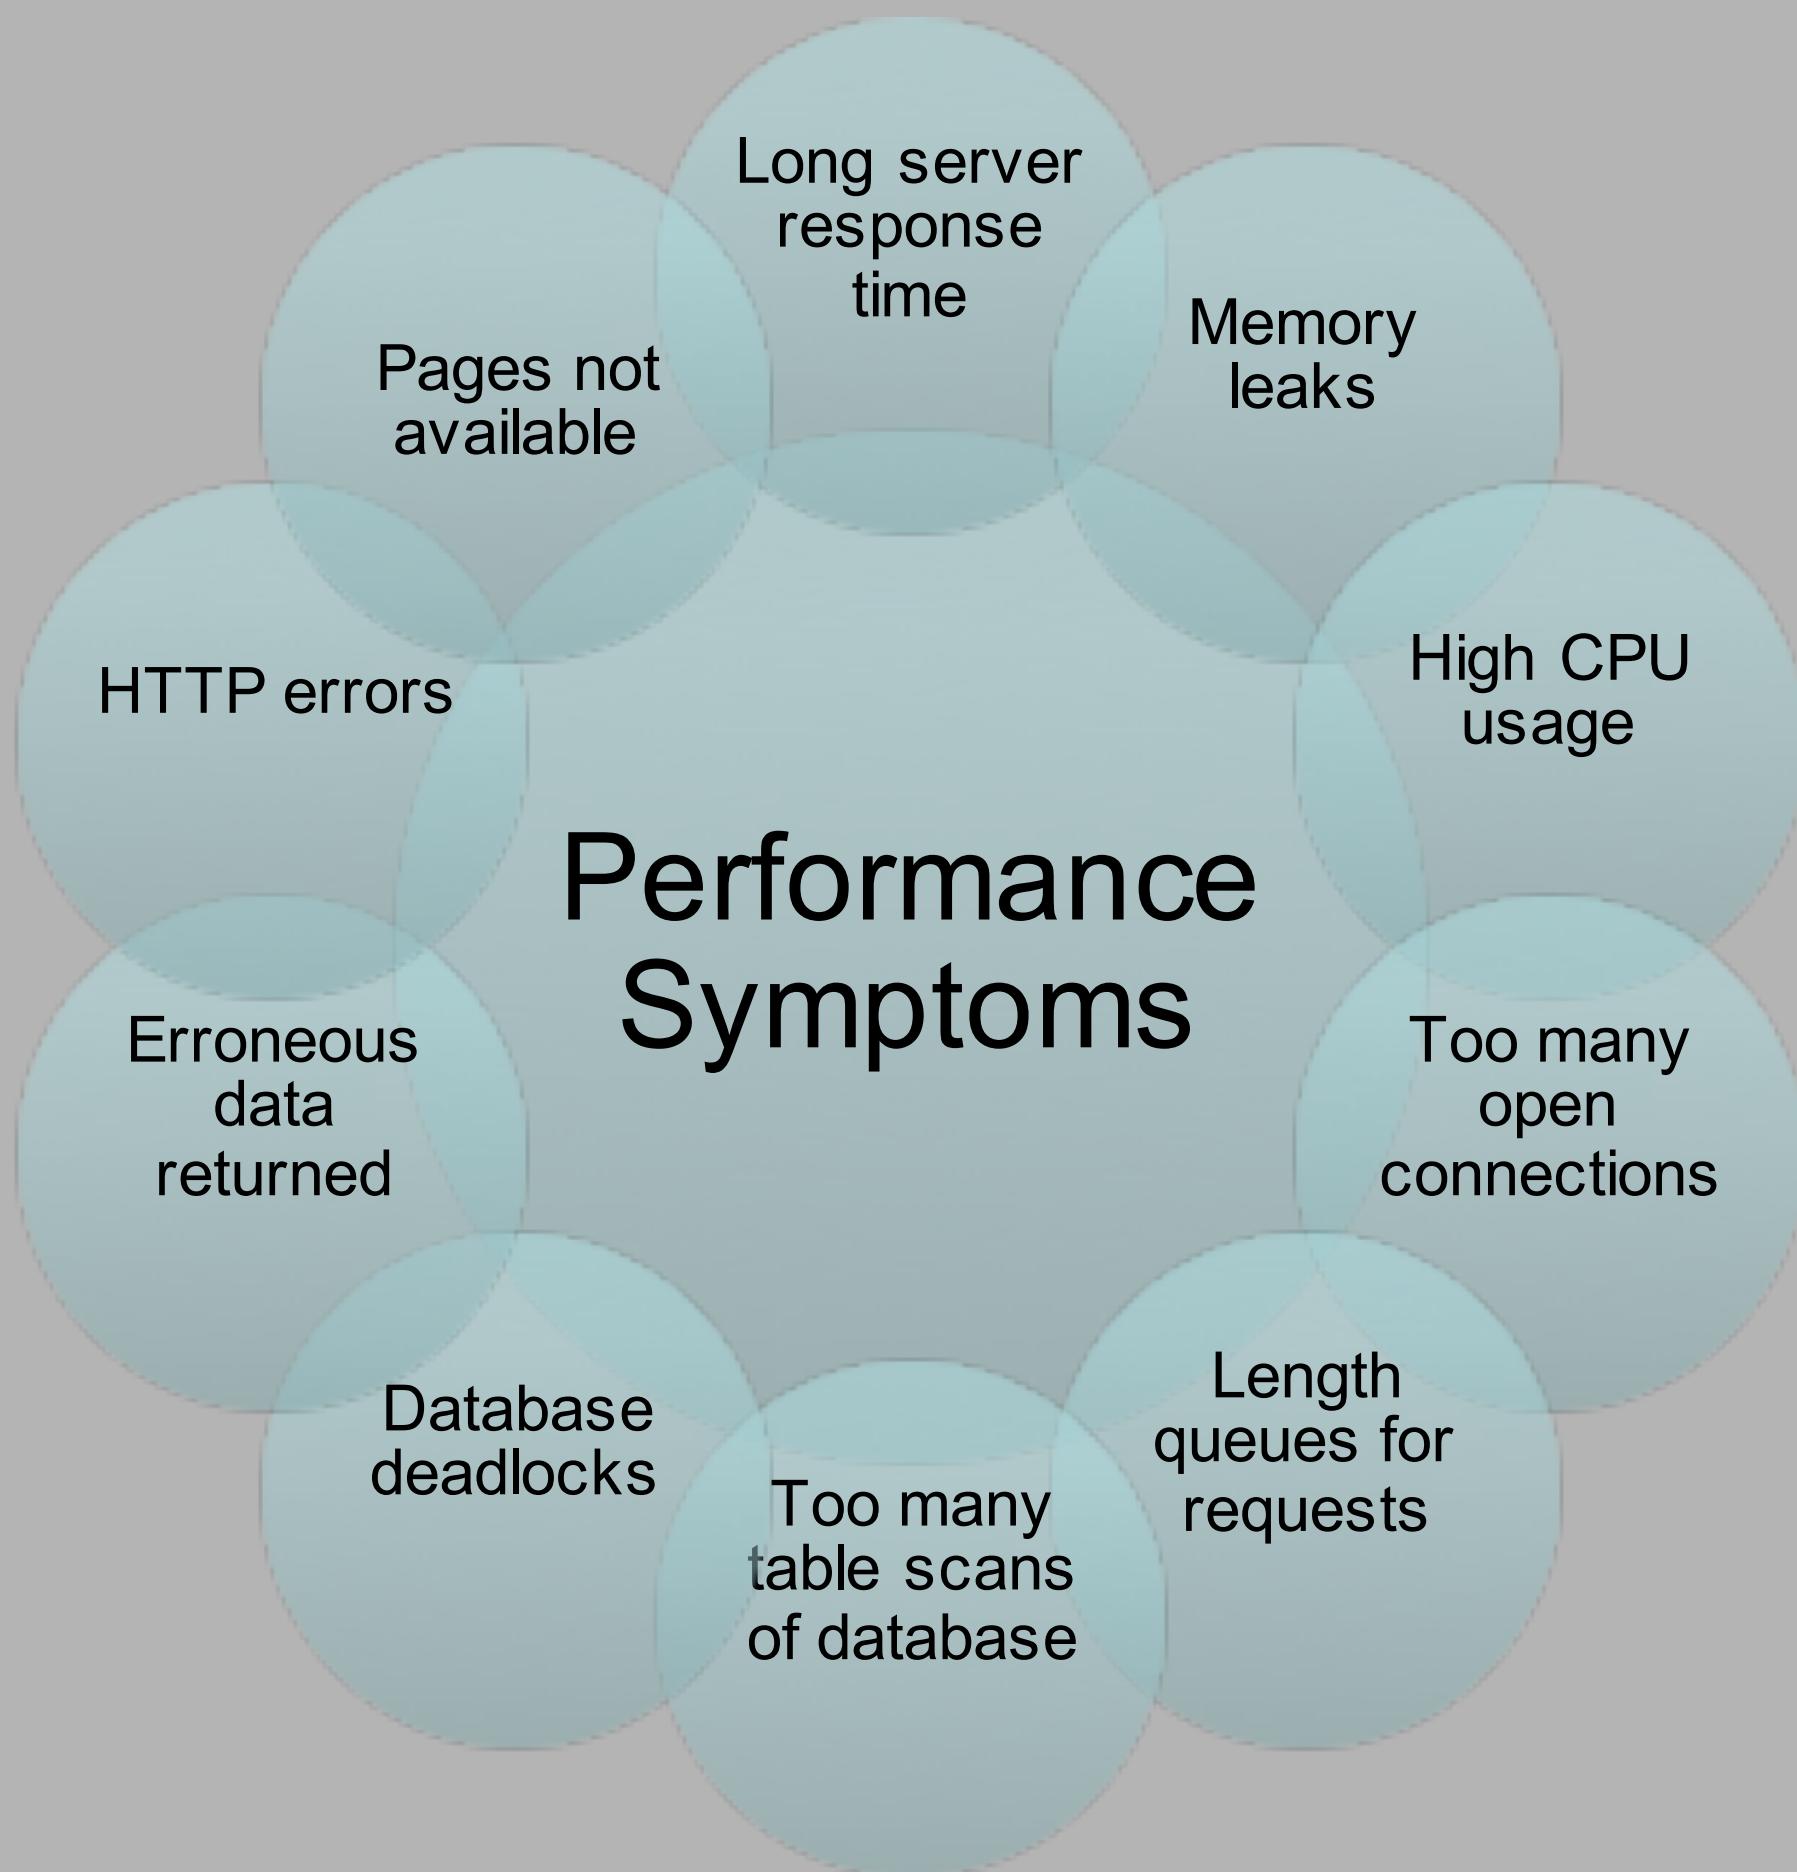
A non-functional testing performed to determine how a system performs in terms of responsiveness and stability under a particular workload. It can also serve to investigate, measure, validate or verify other quality attributes of the system, such as scalability, reliability and resource usage



GOAL/OBJECTIVE IS TO FIND THE BOTTLE NECK IN THE SYSTEM

SYMPTOMS OF PERFORMANCE PROBLEMS

Performance Symptoms



Database

- Insufficient indexing
- Fragmented database
- Out-of-date statistics
- Faulty application design

Web Server

- Poor server design
- Memory problems
- High CPU usage

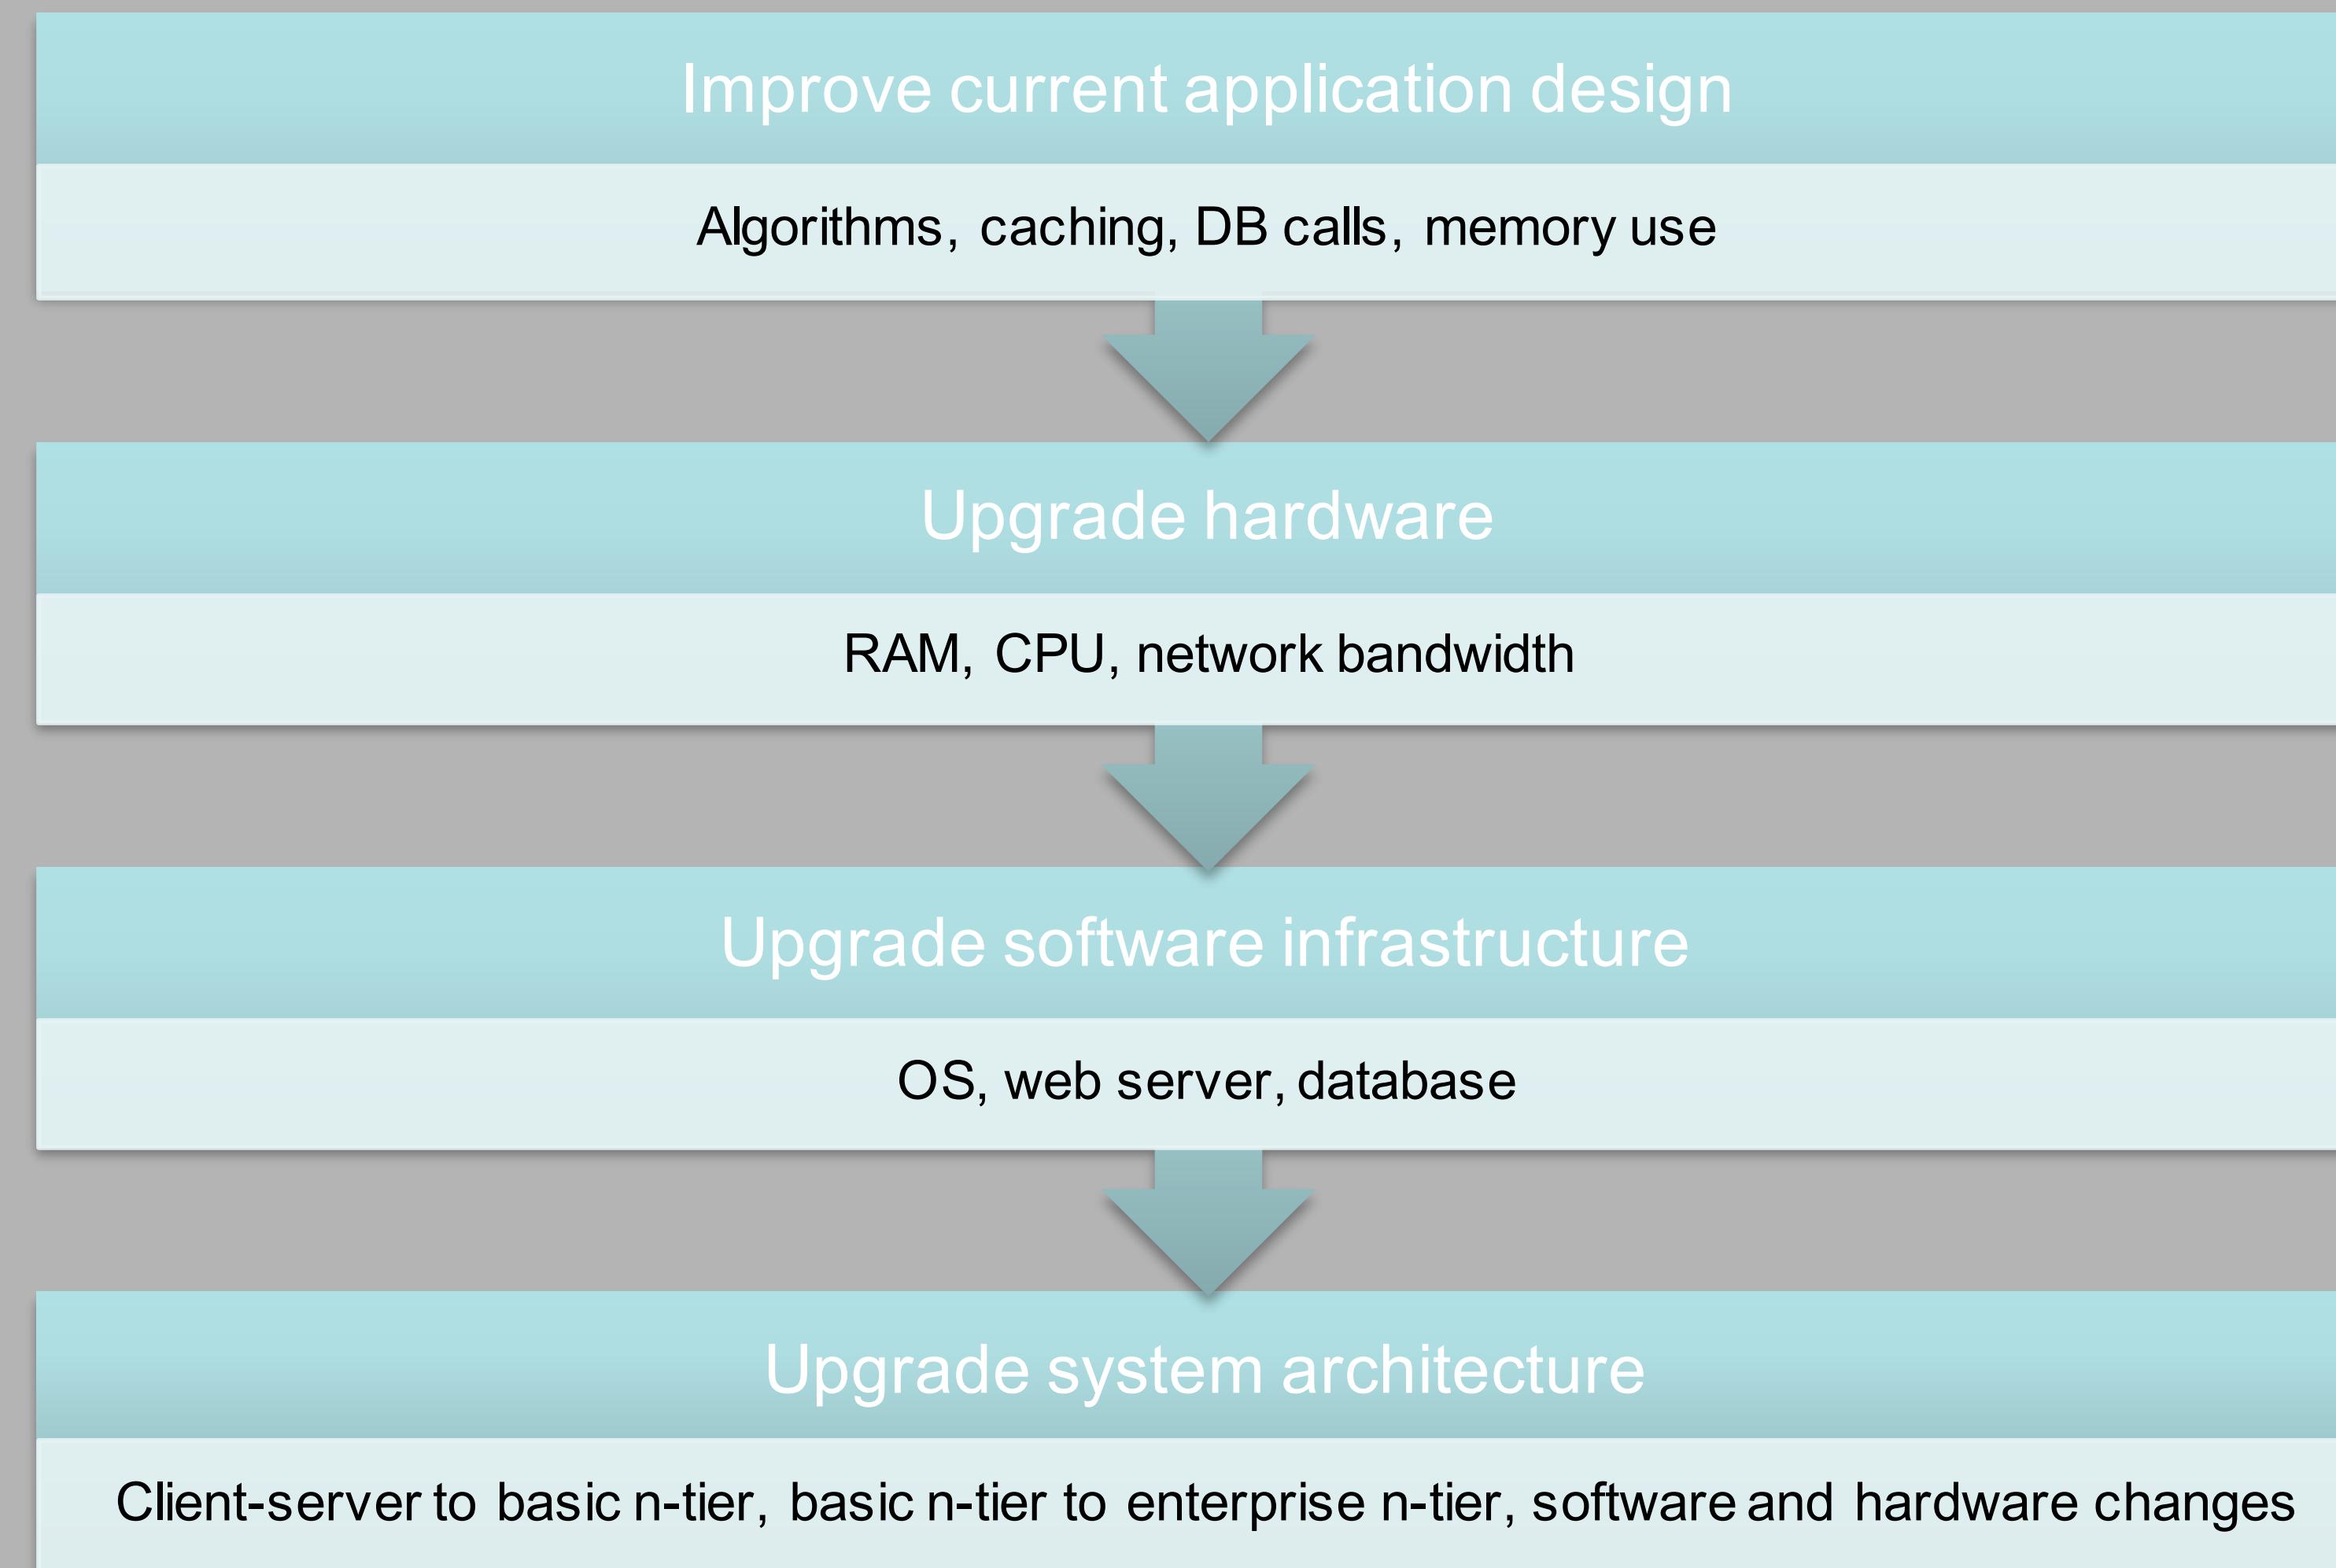
Application Server

- Poor database tuning
- Poor cache management
- Poor session management
- Poor security design

Network

- Firewall throughput
- Internet access throughput
- Load balancers, gateways, routers

SYMPTOMS OF PERFORMANCE PROBLEMS



WORKLOAD MODELING



TYPES OF PERFORMANCE TESTING

Load/Capacity Testing



Stress Testing



Volume Testing



Endurance/Soak Testing



Spike Testing



PERFORMANCE TESTING TOOLS

Load Runner, commercial load testing tool from HP

JMeter, an open source tool from Apache

RPT, commercial load test tool from IBM

NeoLoad, commercial, for Windows, Linux, Solaris

Microsoft Visual Studio Team System 2010, commercial, for Windows, which includes Load Test Analyser and Load Test Monitor tools.

OpenLoad, commercial load testing tool and hosted service

OpenSTA, an open source tool

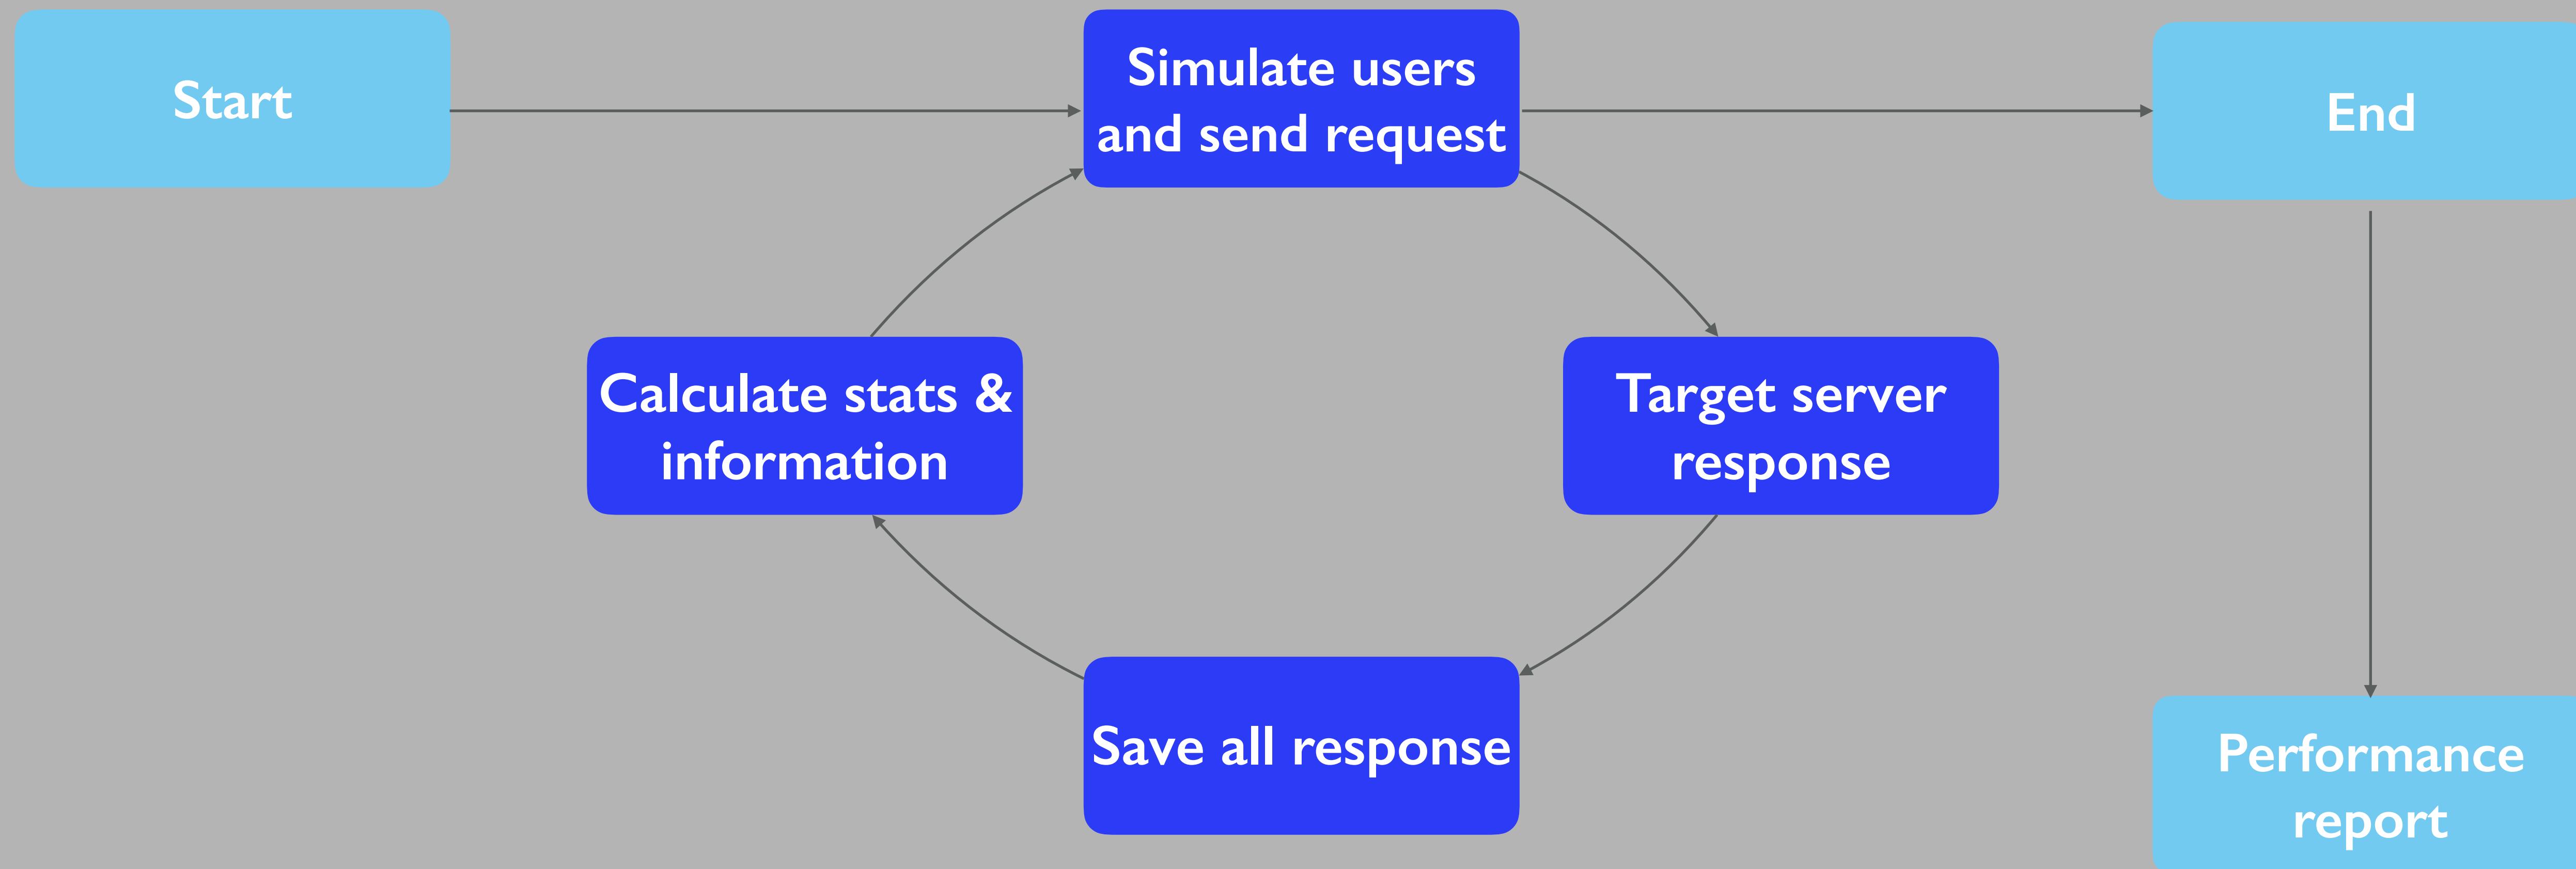
PureLoad, commercial, multi-platform load testing tool

The Grinder, an open source tool

WebServer Stress Tool, commercial and free, from Paessler



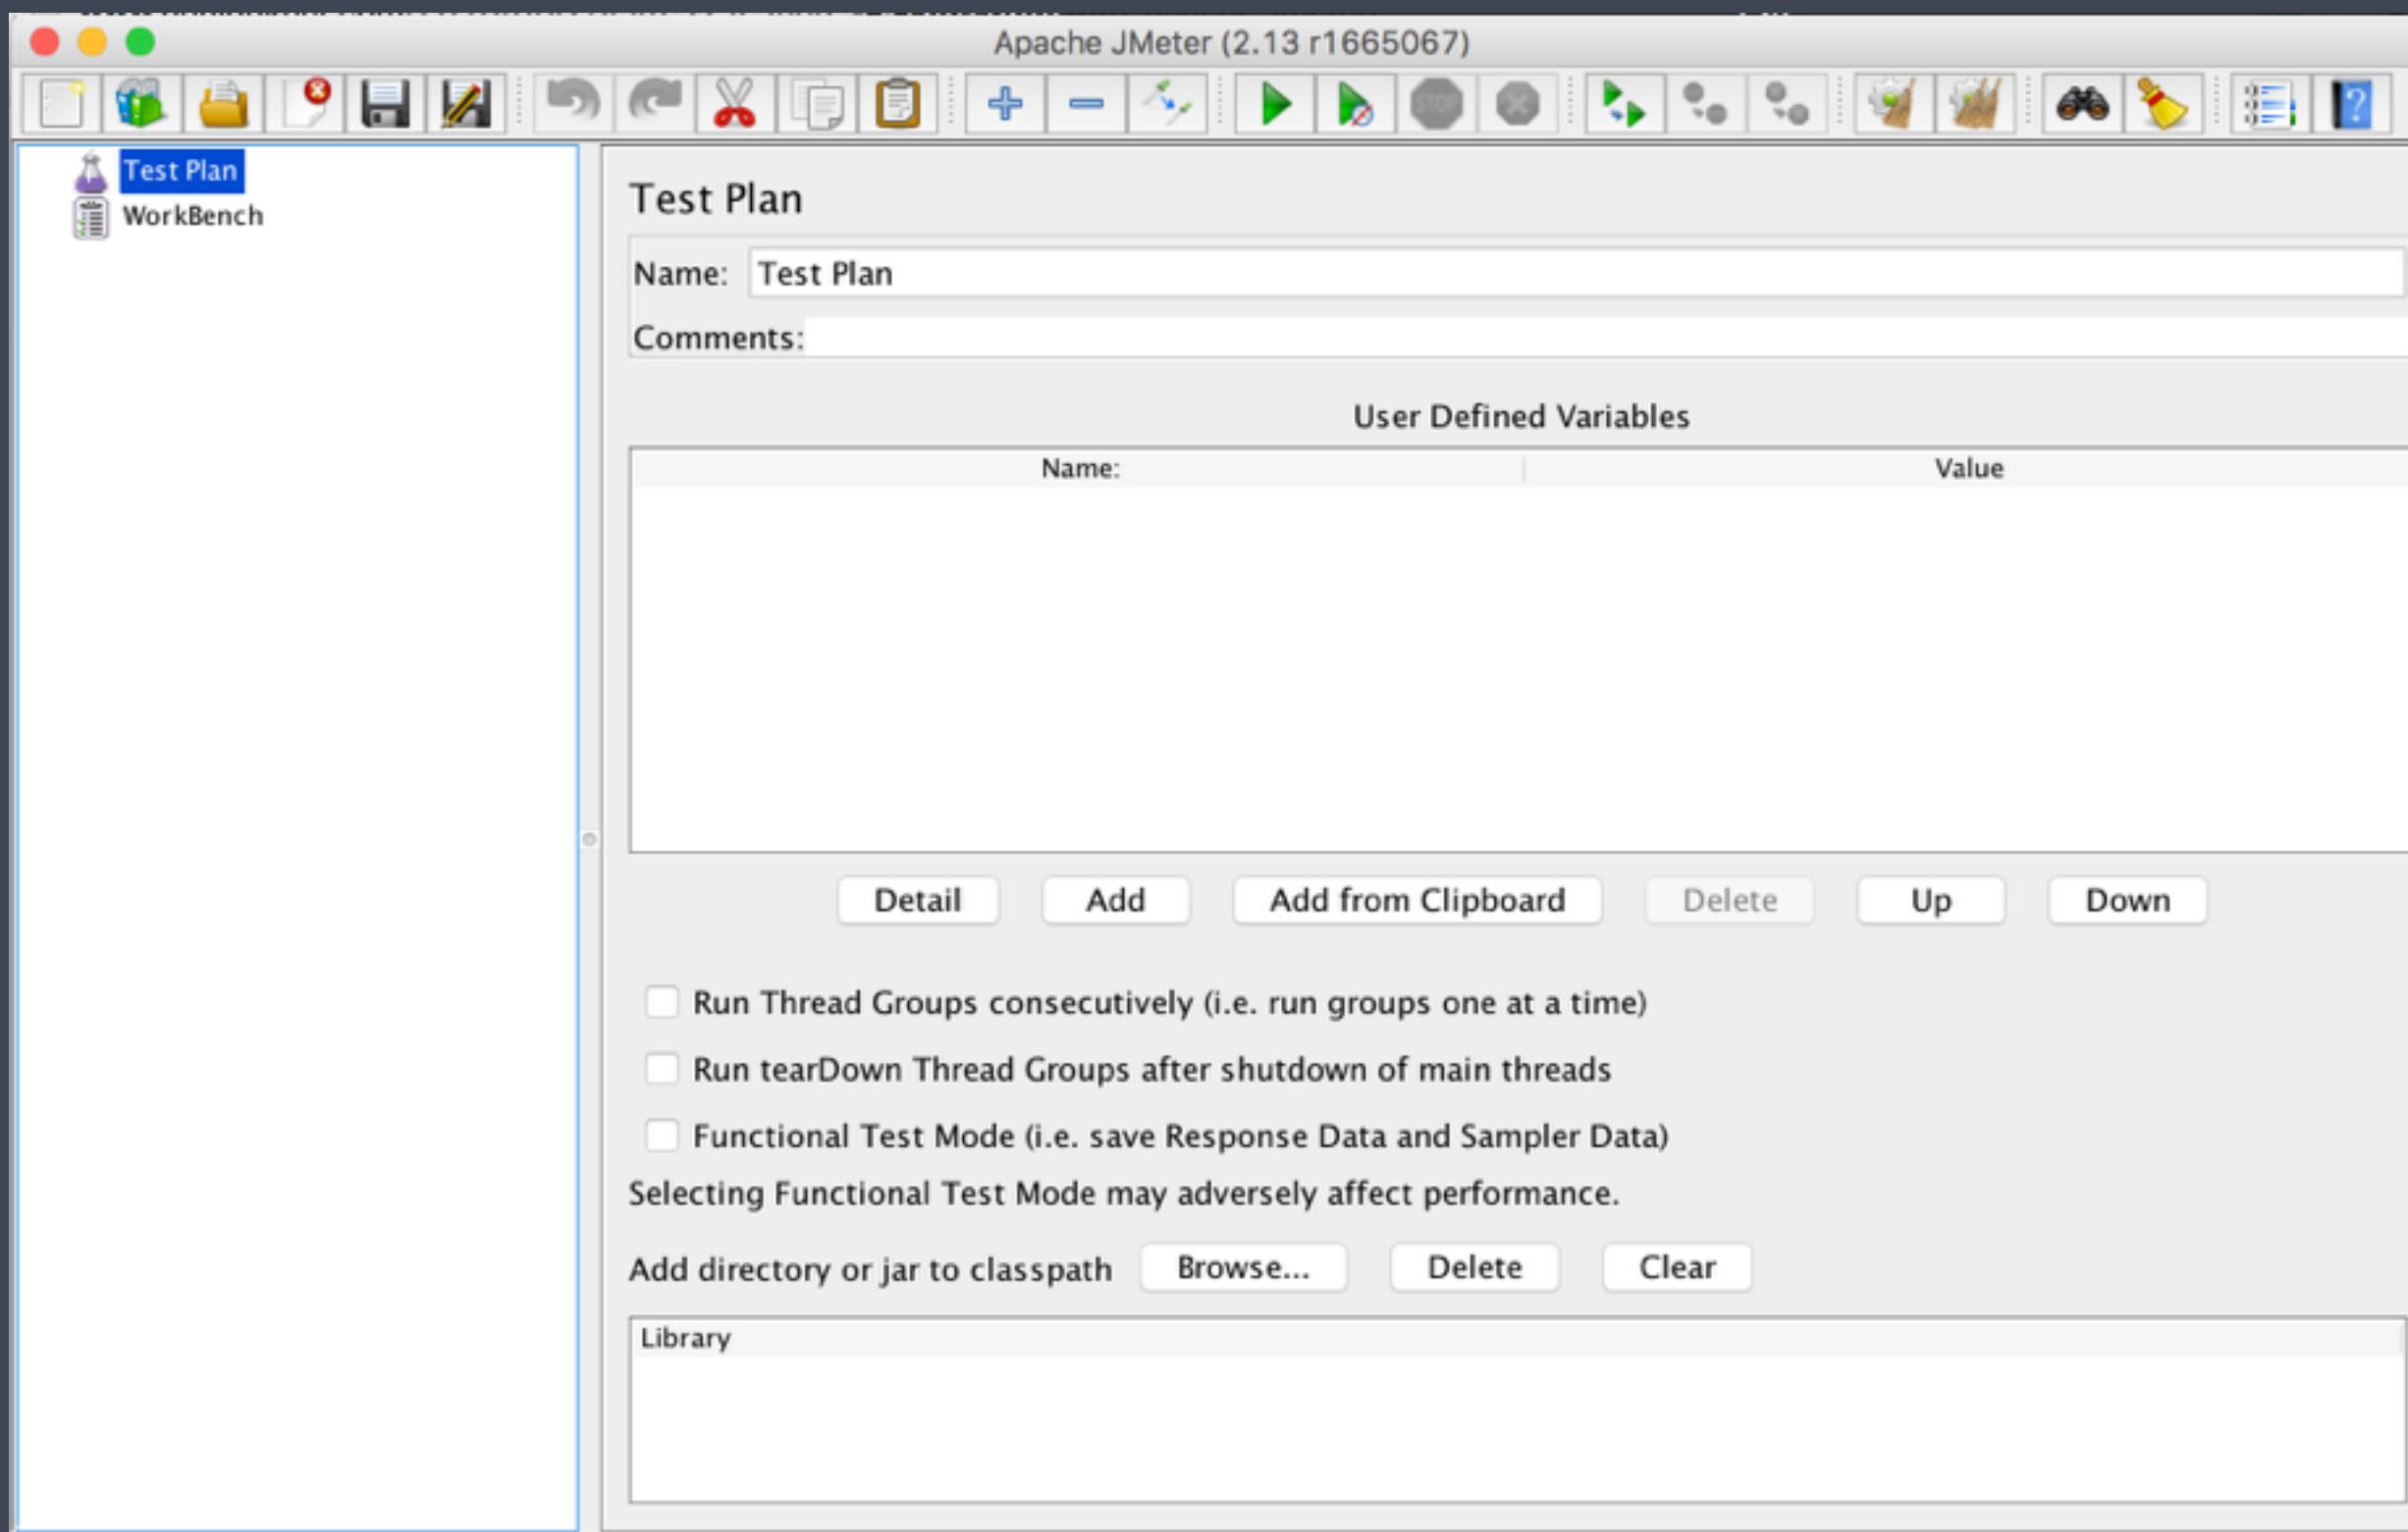
HOW JMETER WORK



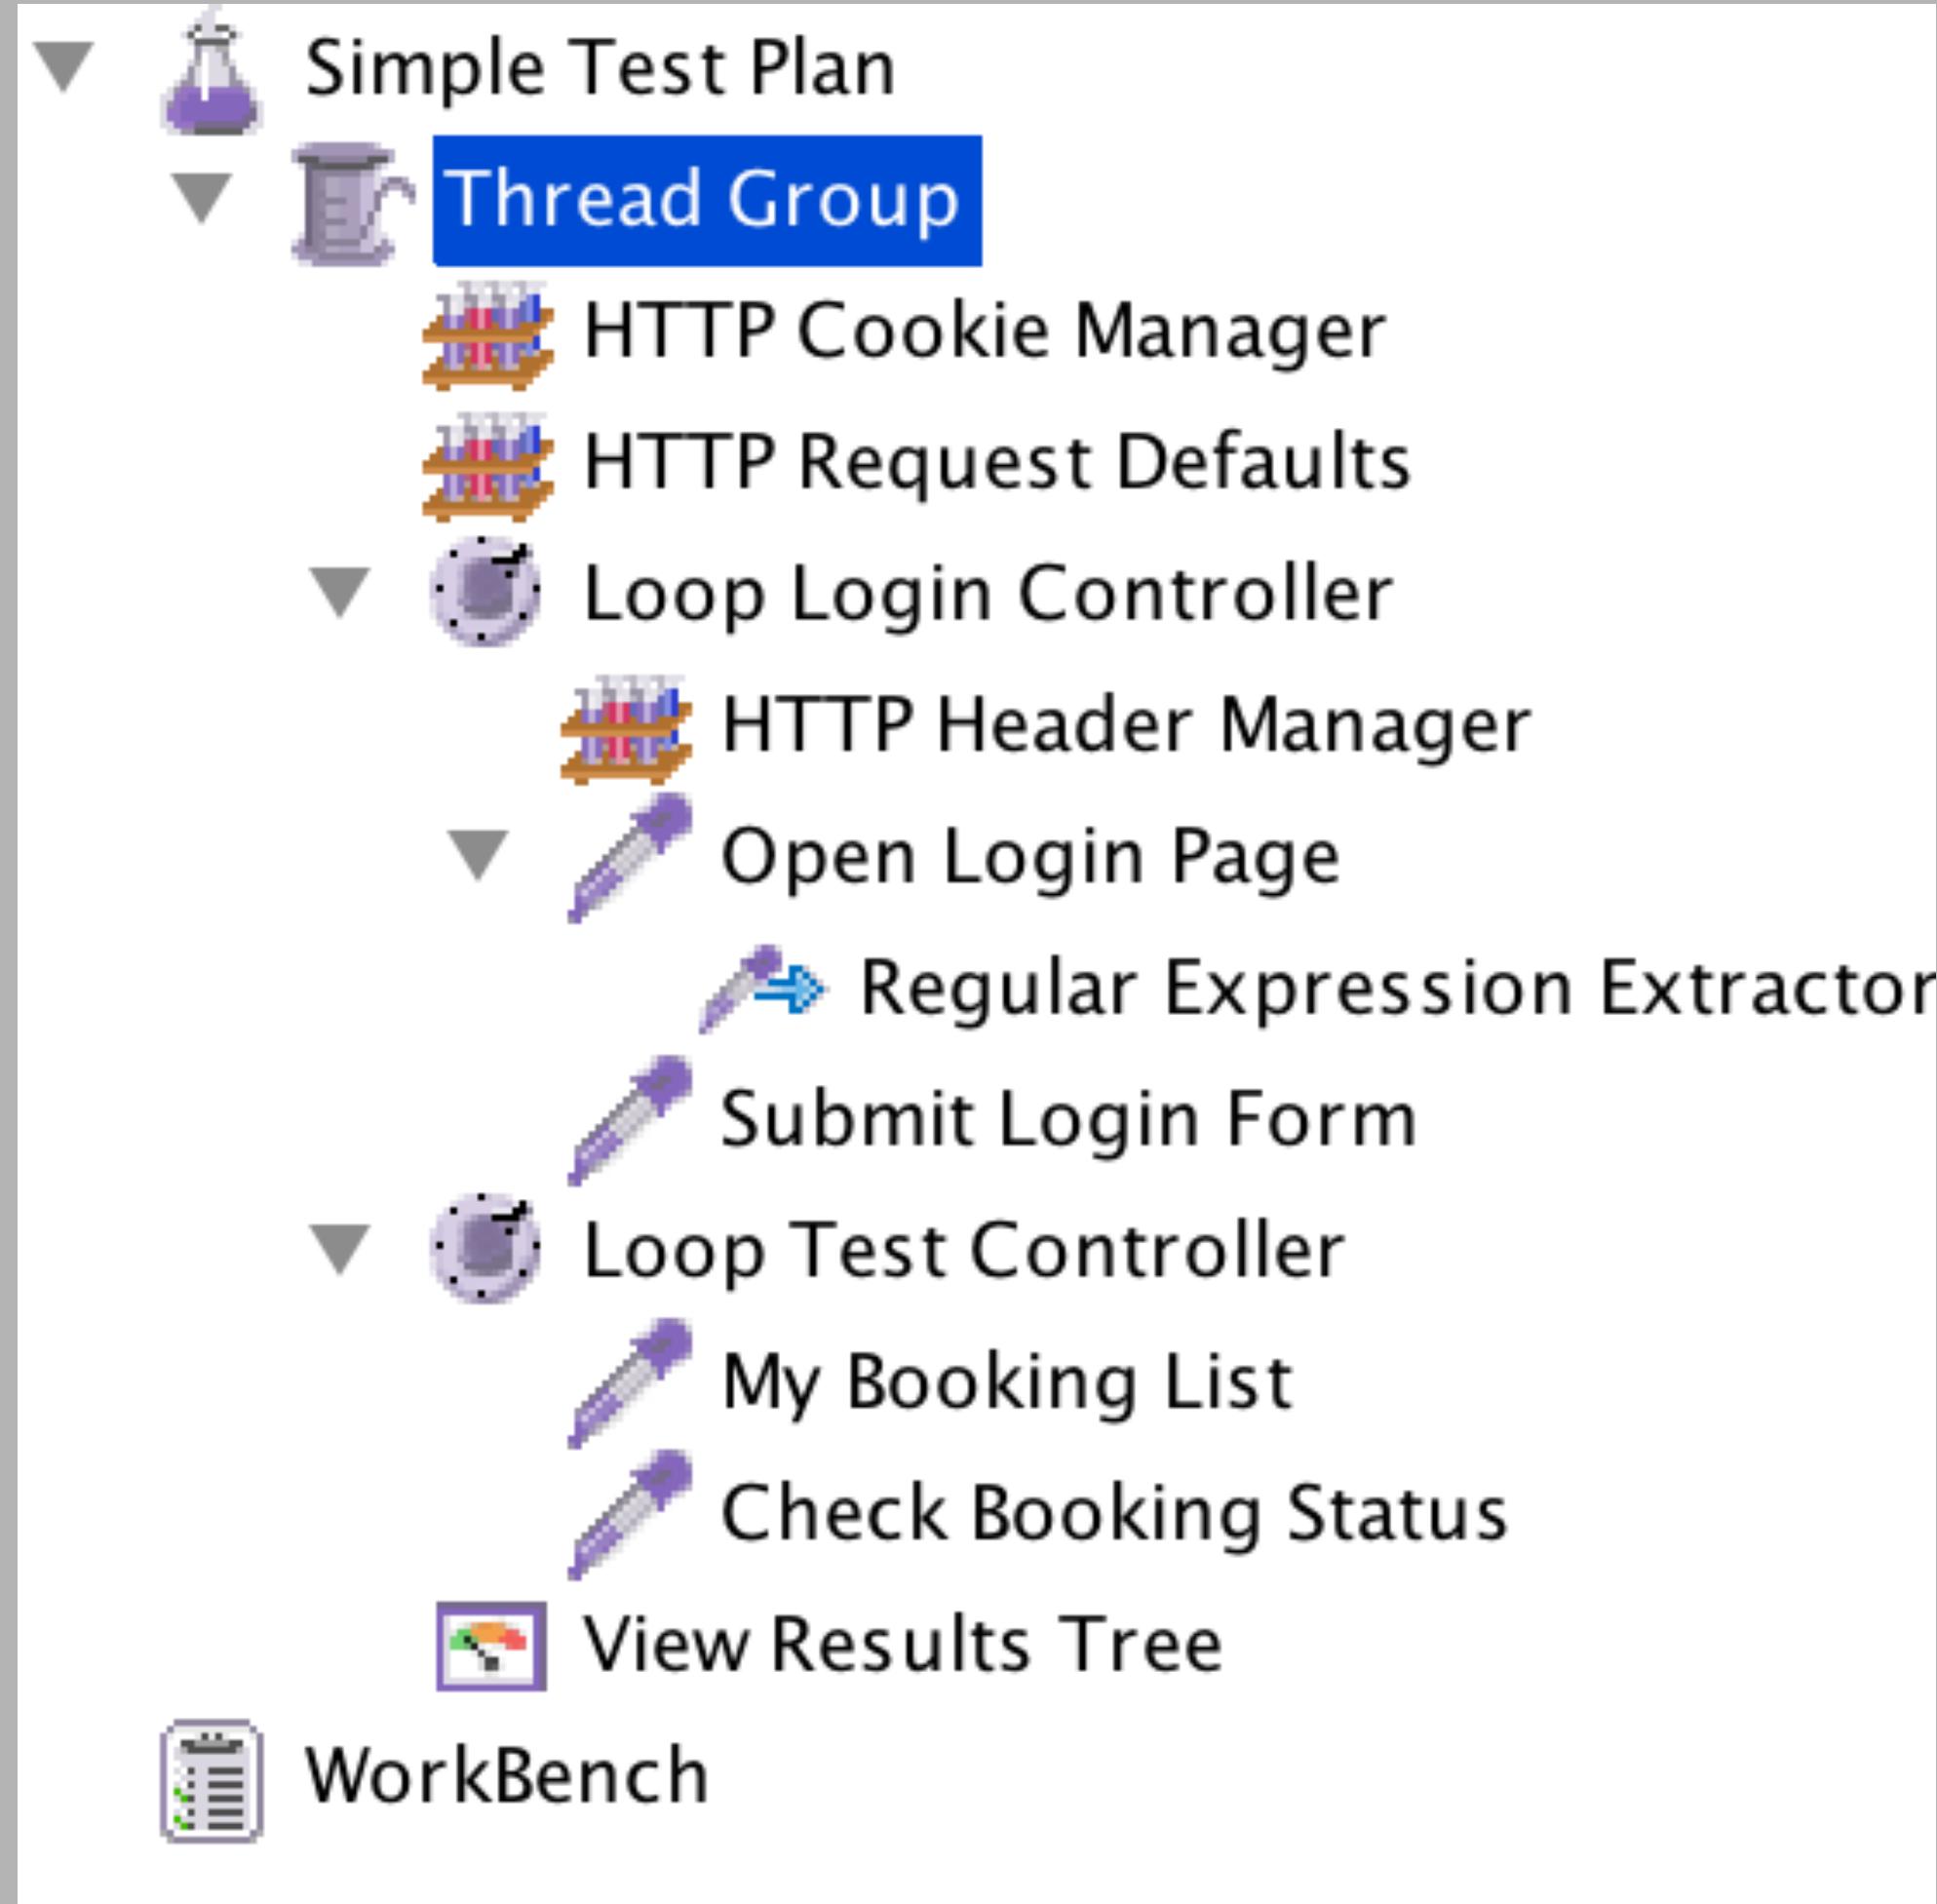
STARTING JMETER

\$JMETER_HOME/bin/jmeter

\$JMETER_HOME/bin/jmeter.bat



TEST PLAN BUILDING BLOCKS



Thread Group

Setup number of threads

Set up ramp up period

No. of times test executes

Controllers

Sampler (Send Request to Server)

Logical Controller (Customize logic to send request)

Listener

Graph Result

View Results Tree and many more.

Timers

Delay next request for certain amount of time

Assertions

Allow you to assert facts about responses received from HTTP requests

Configuration Elements

Allow you to configure settings

Preprocessor

Execute prior to sampler request

Post Processor

Execute some action after sampler request



LAB 2-6

Basic Load Testing

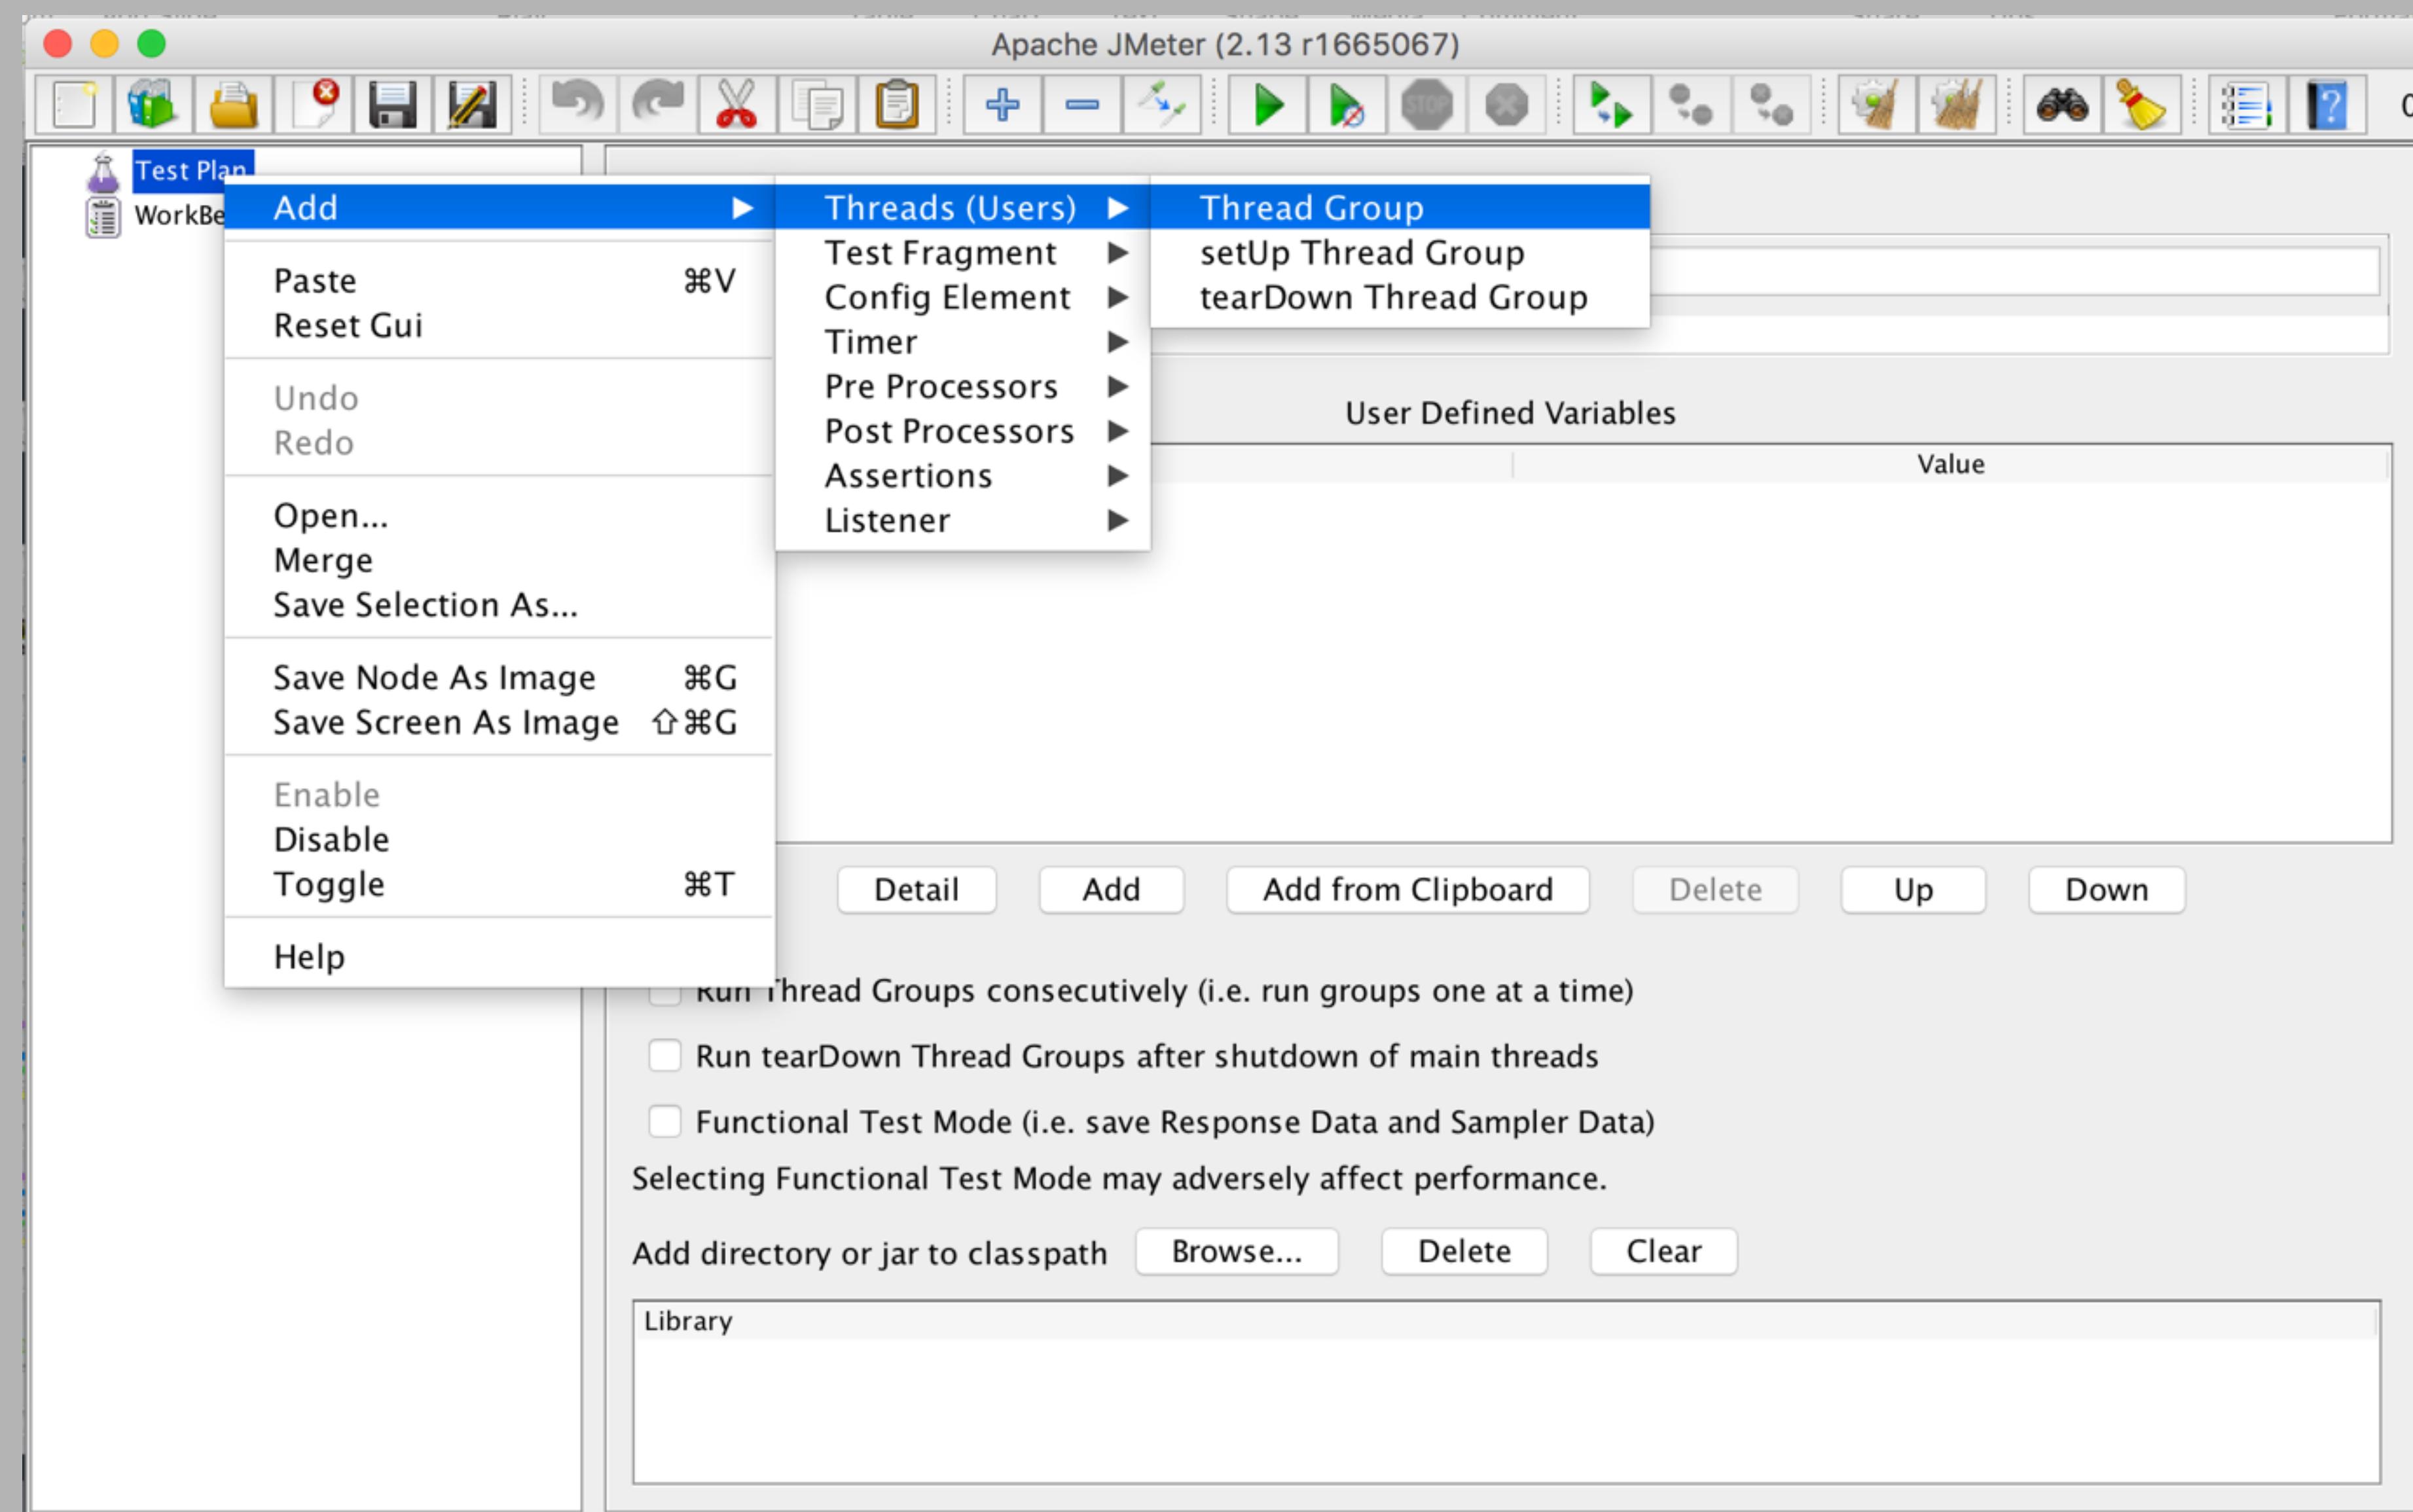
ASSIGNMENT

Given Scenarios:

Create a test plan having 5 Users which send 2 request to your API and repeat it twice.

$$= 5 \times 2 \times 2$$

CREATE THREAD GROUP



CREATE THREAD GROUP

Thread Properties

Number of Threads (users):

Ramp-Up Period (in seconds):

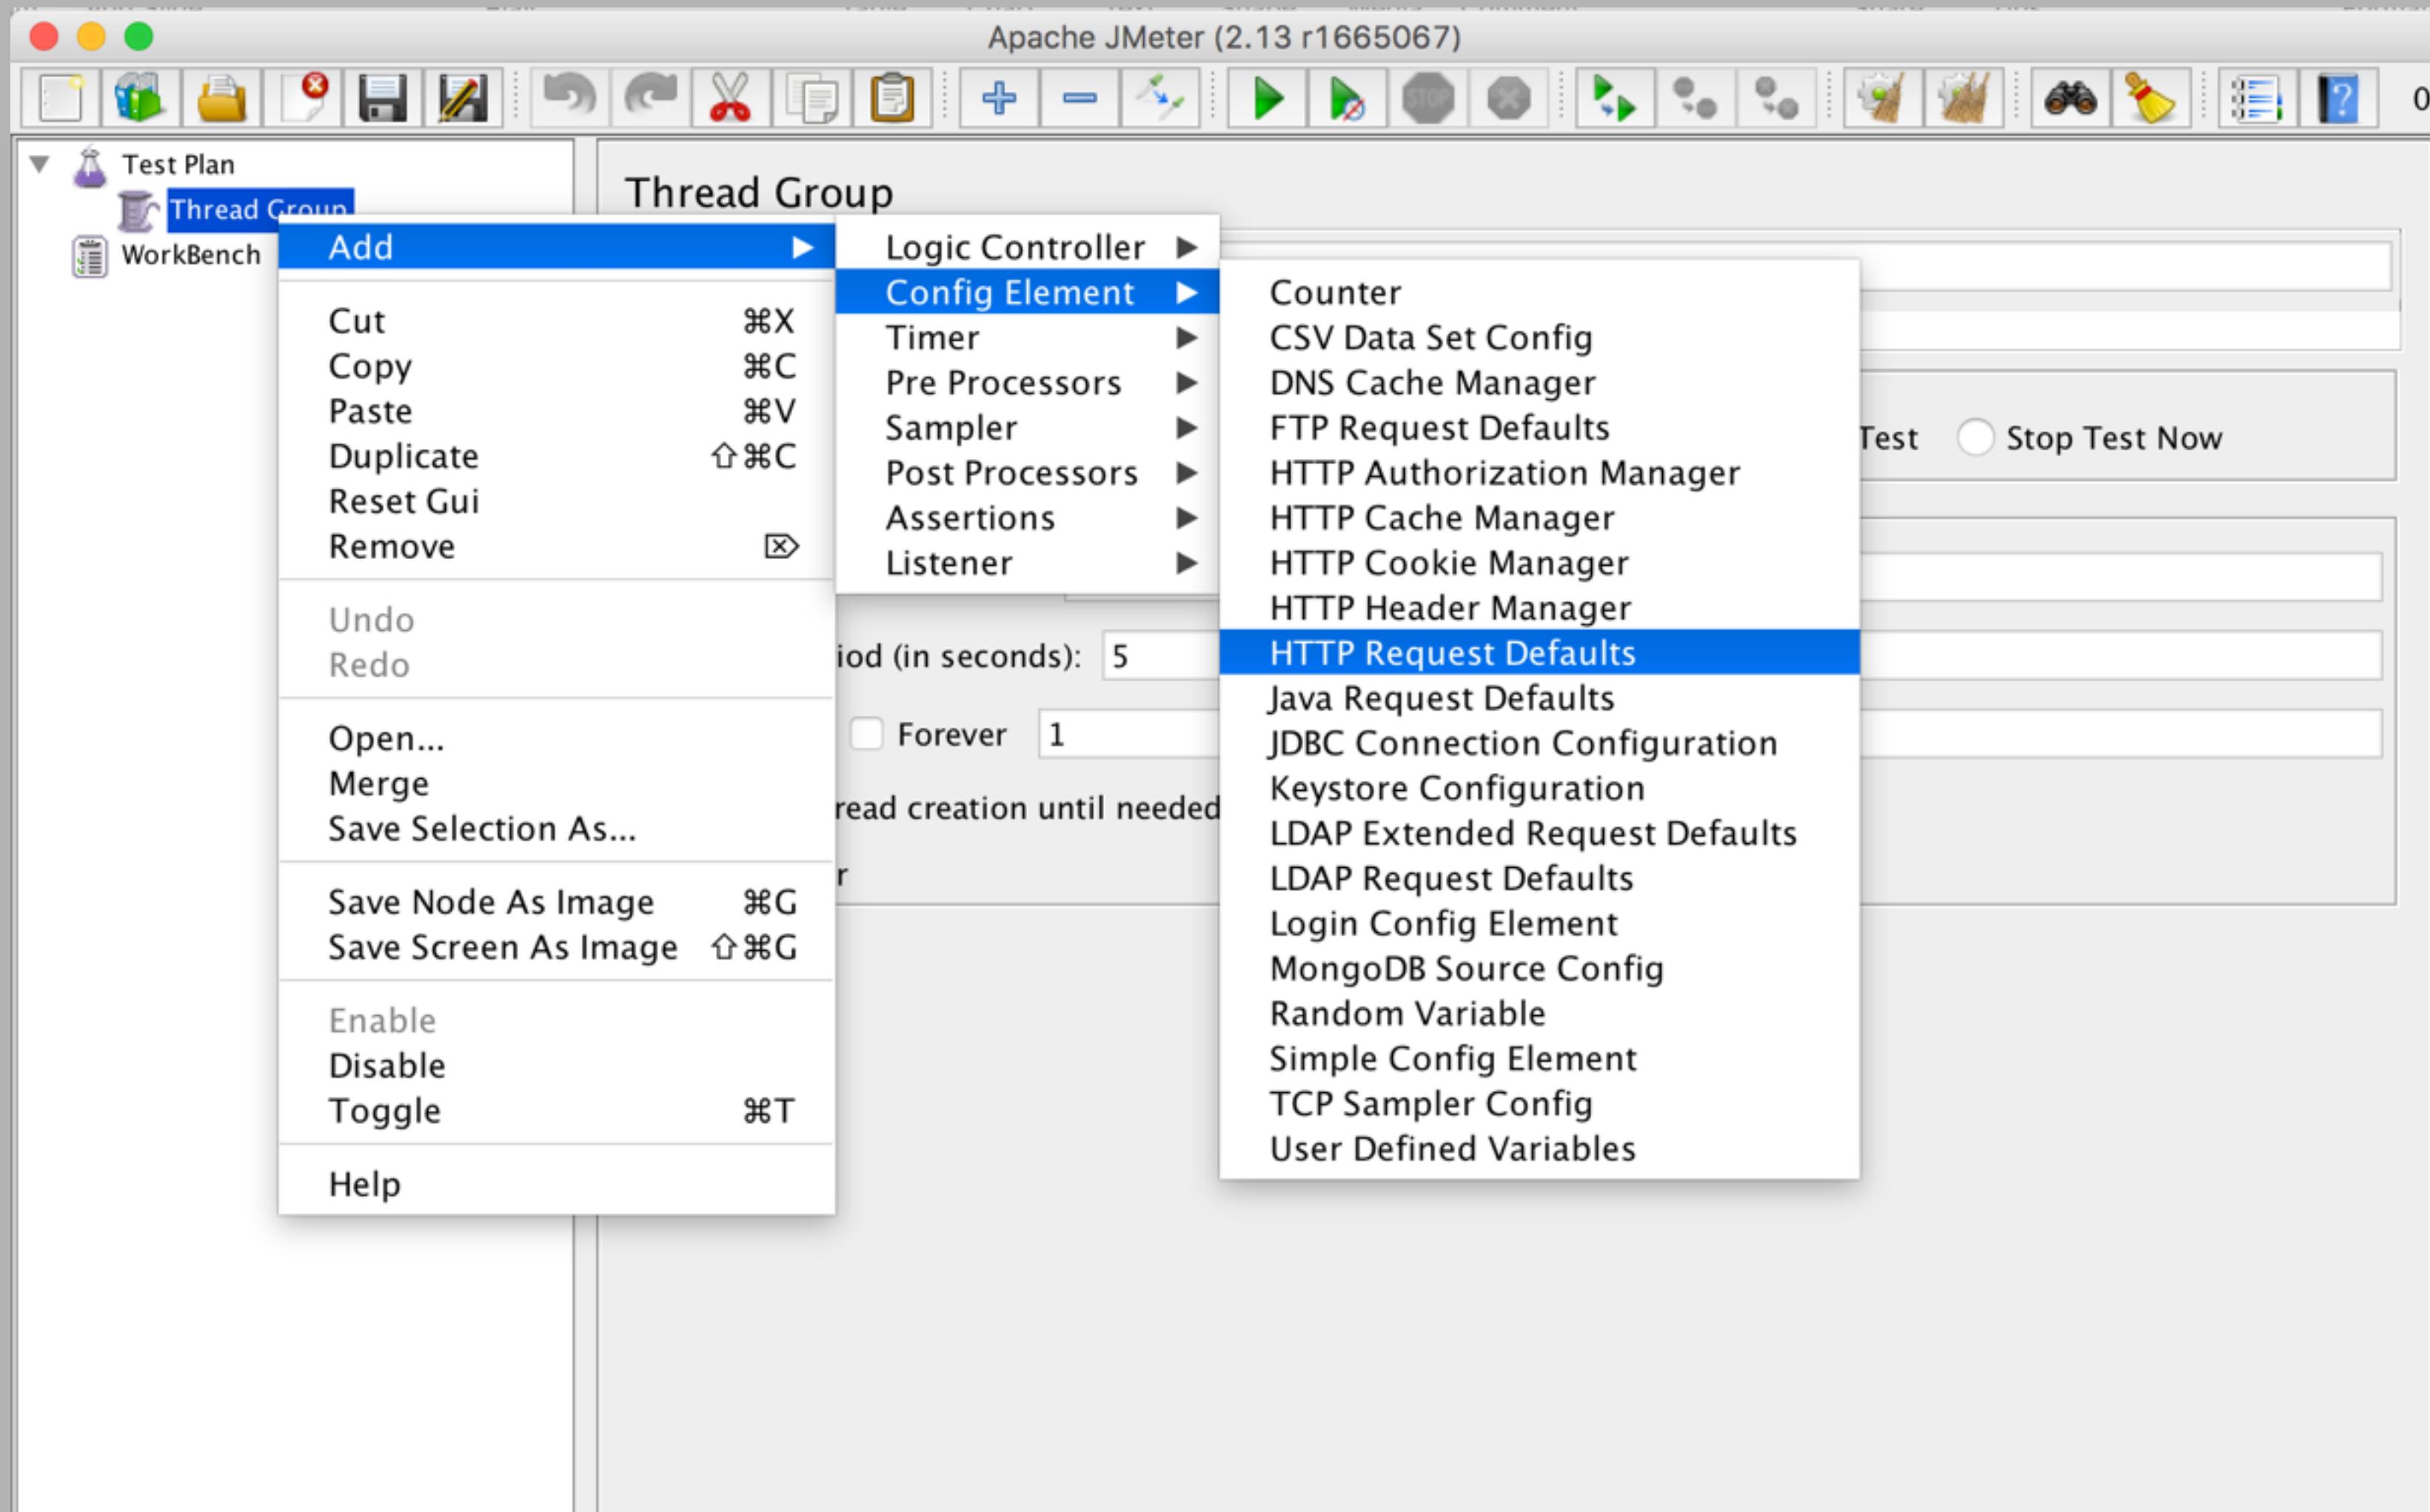
Loop Count: Forever

Delay Thread creation until needed

Scheduler

- Number of Threads (users): The number of users that JMeter will attempt to simulate.
- Ramp-Up Period (in seconds): The duration of time that JMeter will distribute the start of the threads over.
- Loop Count: The number of times to execute the test.

CREATE CONFIG ELEMENT



- HTTP Request Defaults
- HTTP Header Manager
- HTTP Cookie Manager
- HTTP Cache Manager
- HTTP Authorisation Manager

CREATE HTTP REQUEST DEFAULT

Apache JMeter (2.13 r1665067)

Test Plan

Thread Group

HTTP Request Defaults

WorkBench

HTTP Request Defaults

Name: **HTTP Request Defaults**

Comments:

Web Server

Server Name or IP/Port Number: **www.google.com**

Timeouts (milliseconds)

Connect: Response:

HTTP Request

Implementation: **HttpClient4**

Protocol [http]: **80**

Content encoding:

Path:

Parameters

Send Parameters With the Request:

Name:	Value	Encode?	Include Equals?

Detail Add Add from Clipboard Delete Up Down

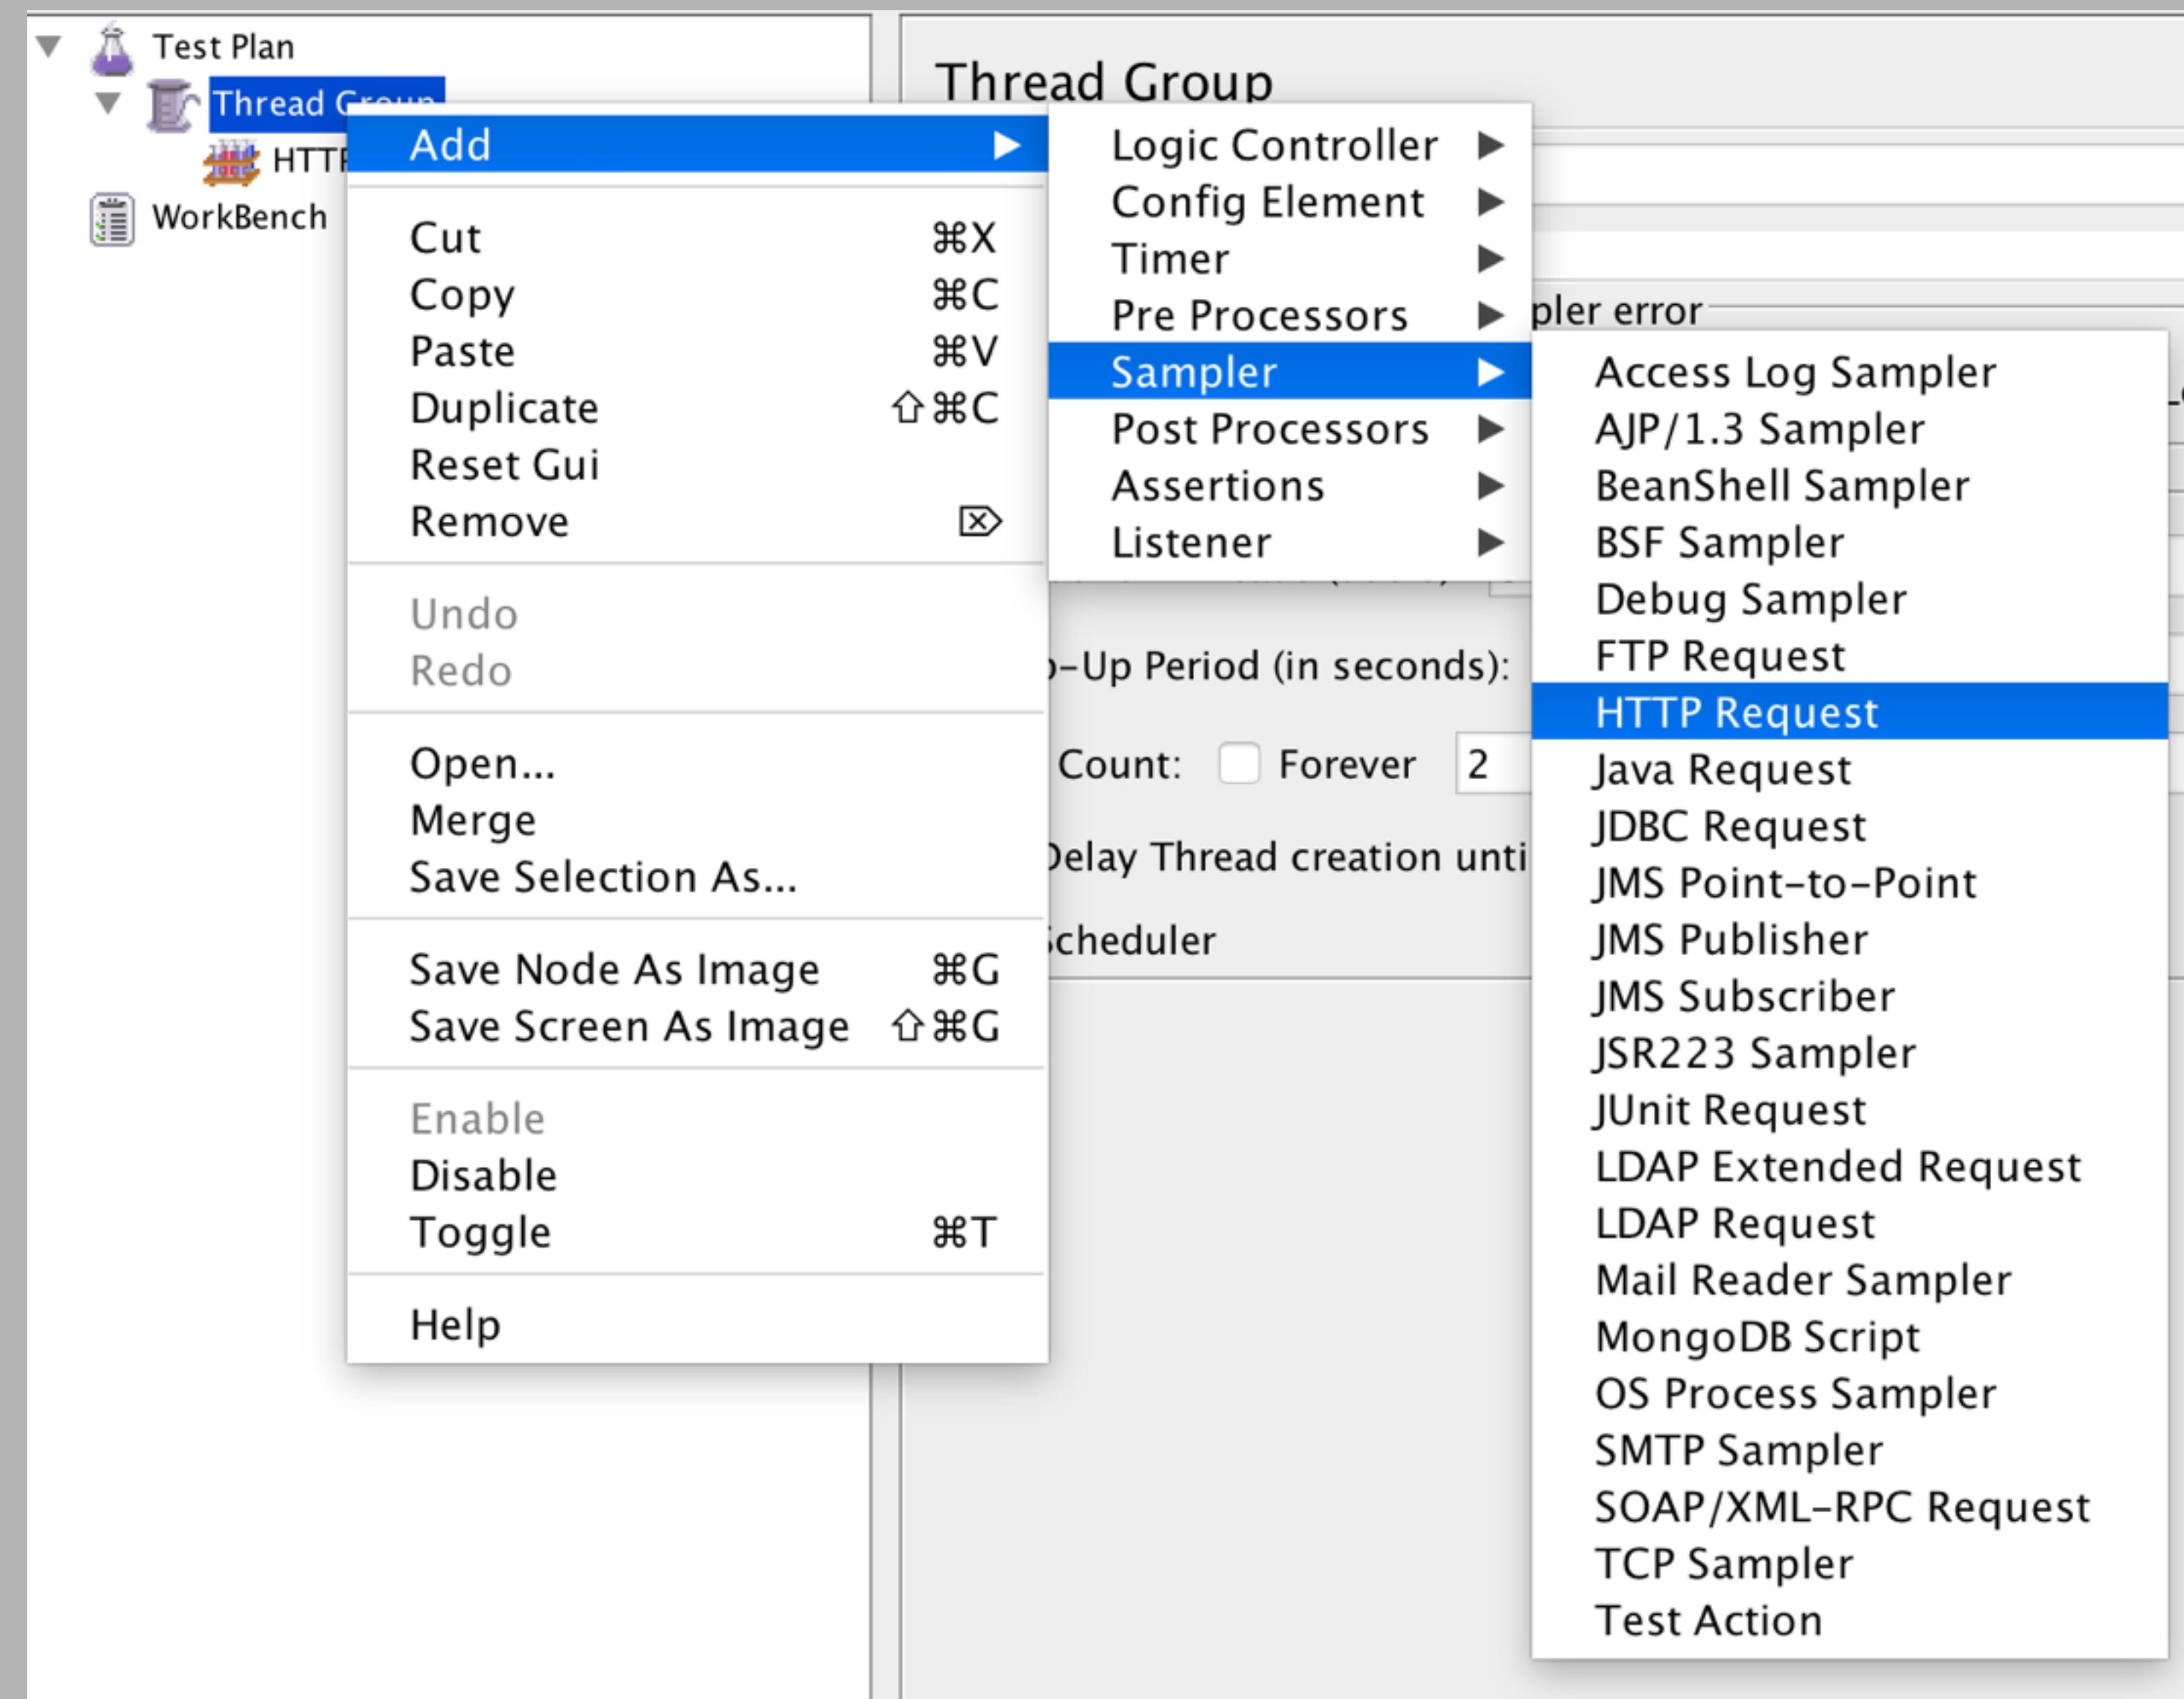
Proxy Server

Server Name or IP: Port Number: Username Password

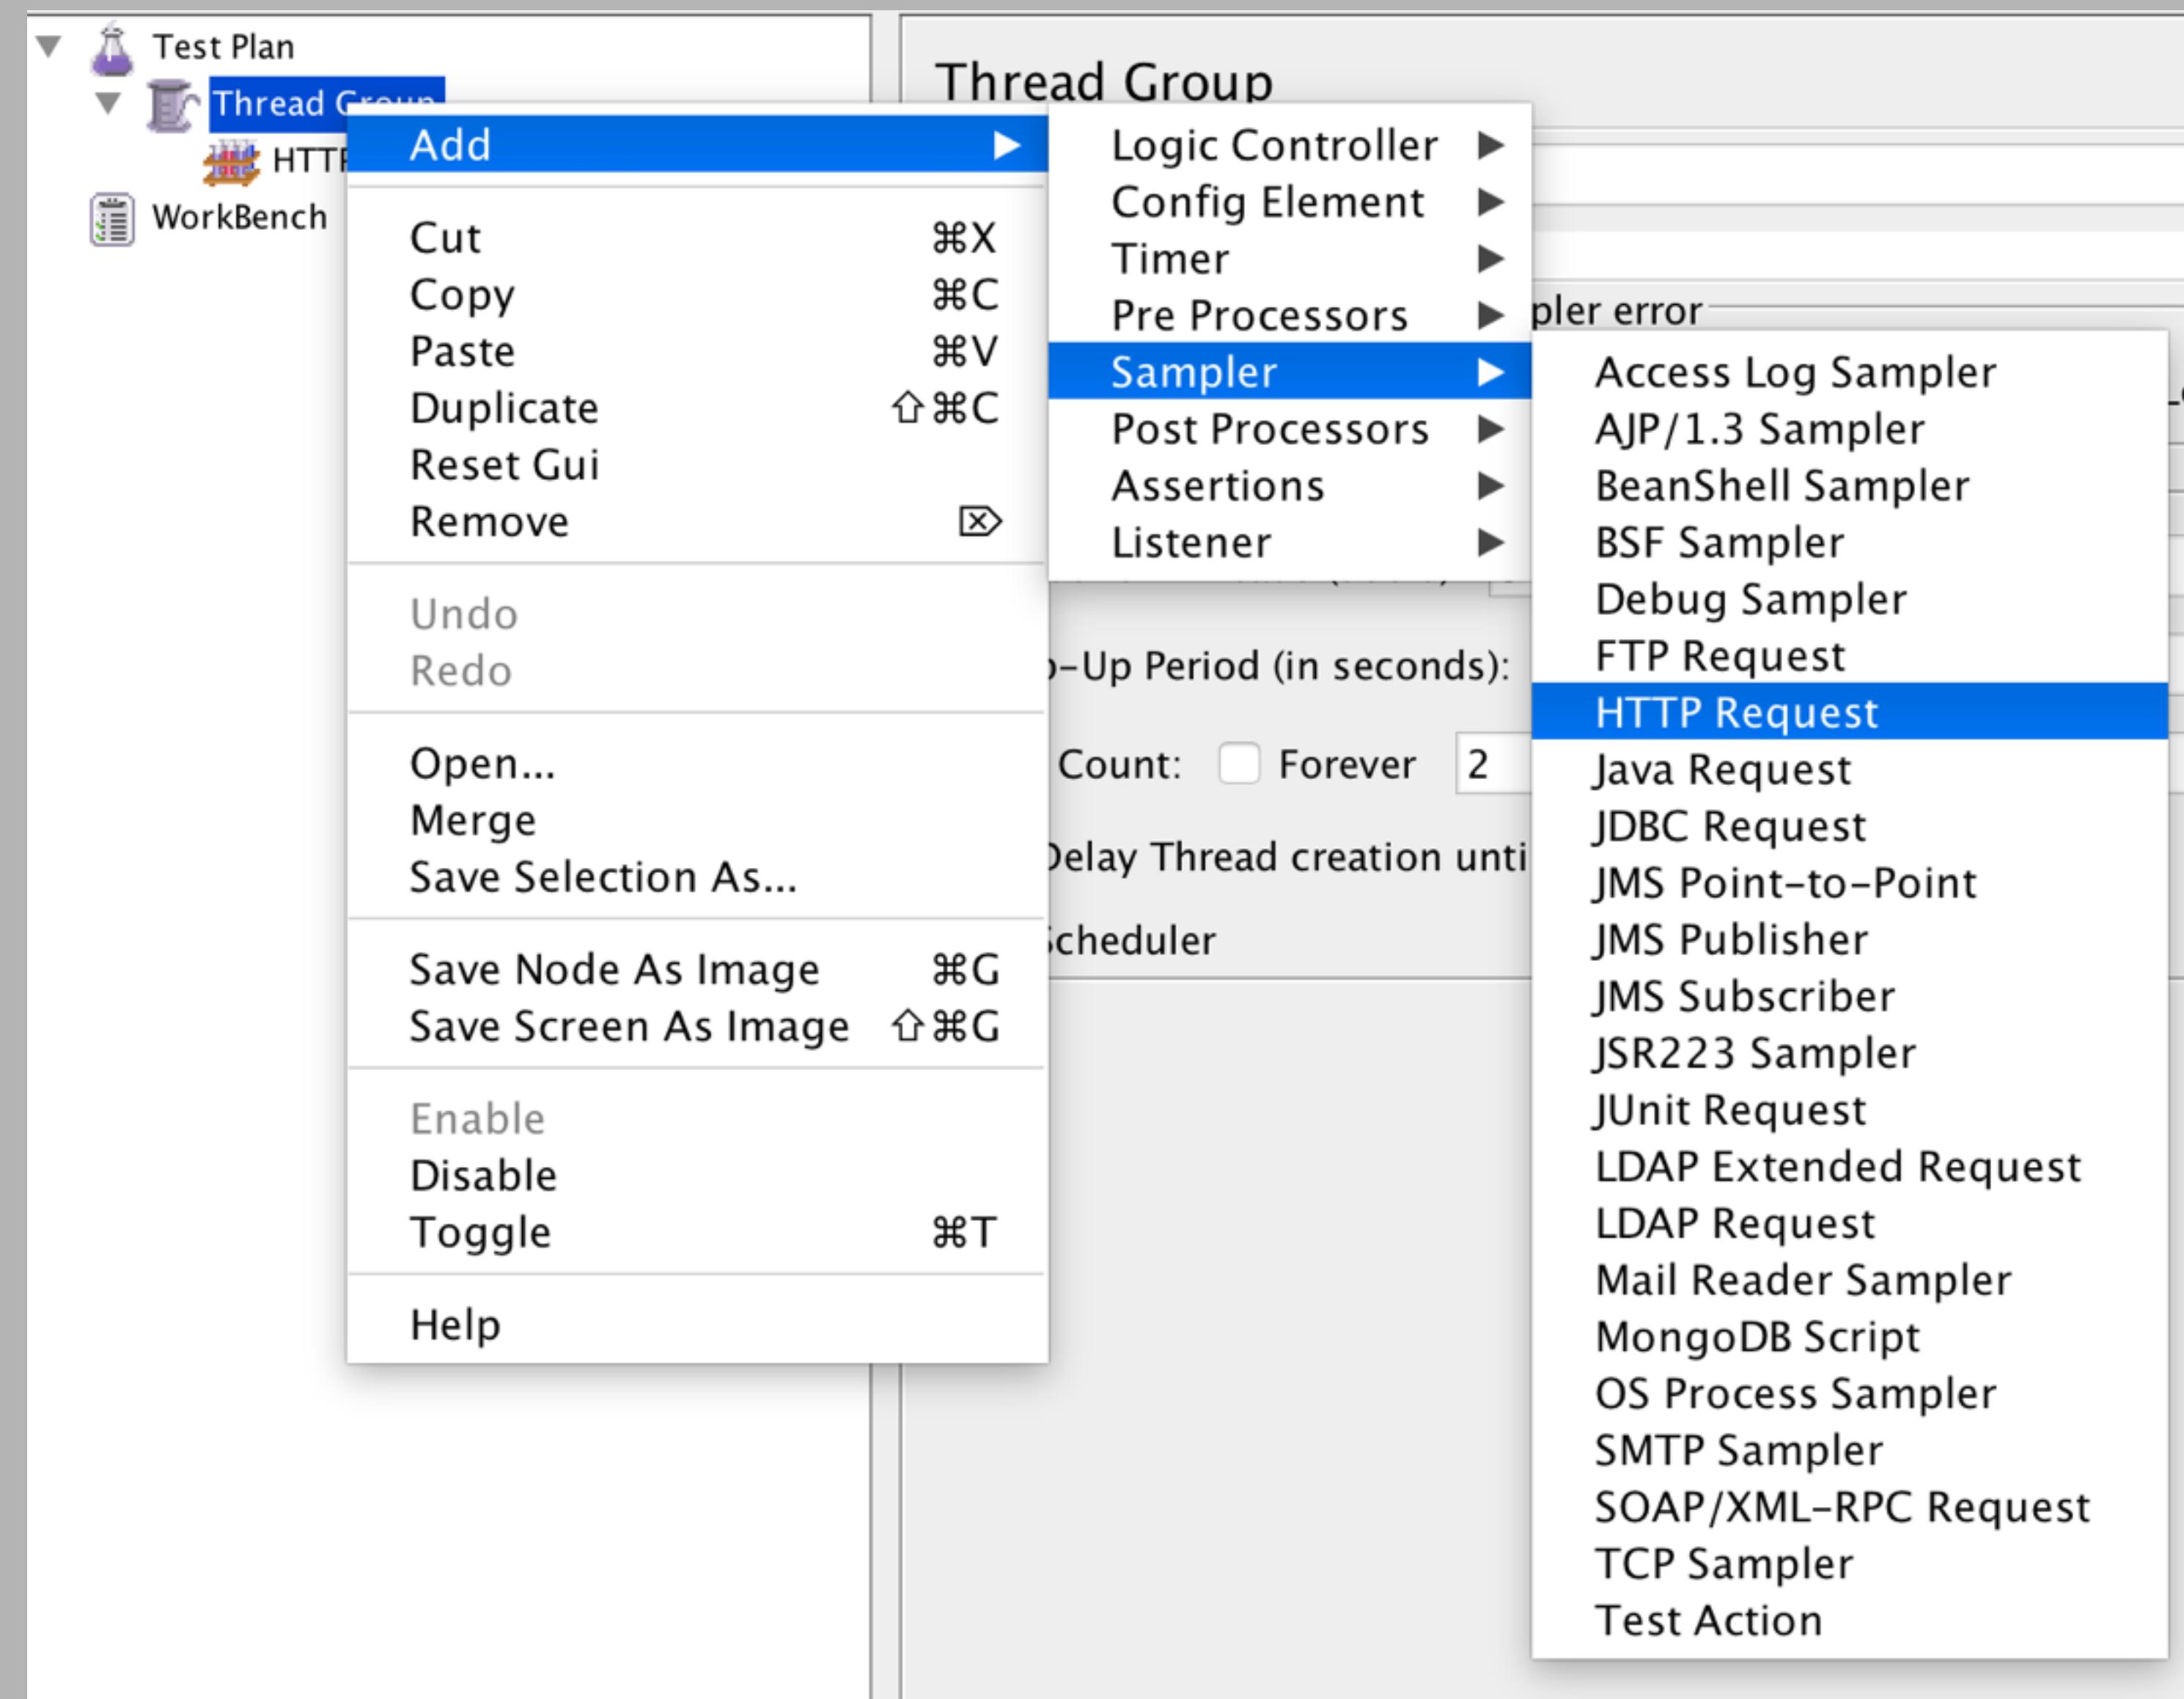
Embedded Resources from HTML Files

Retrieve All Embedded Resources Use concurrent pool. Size: **4** URLs must match:

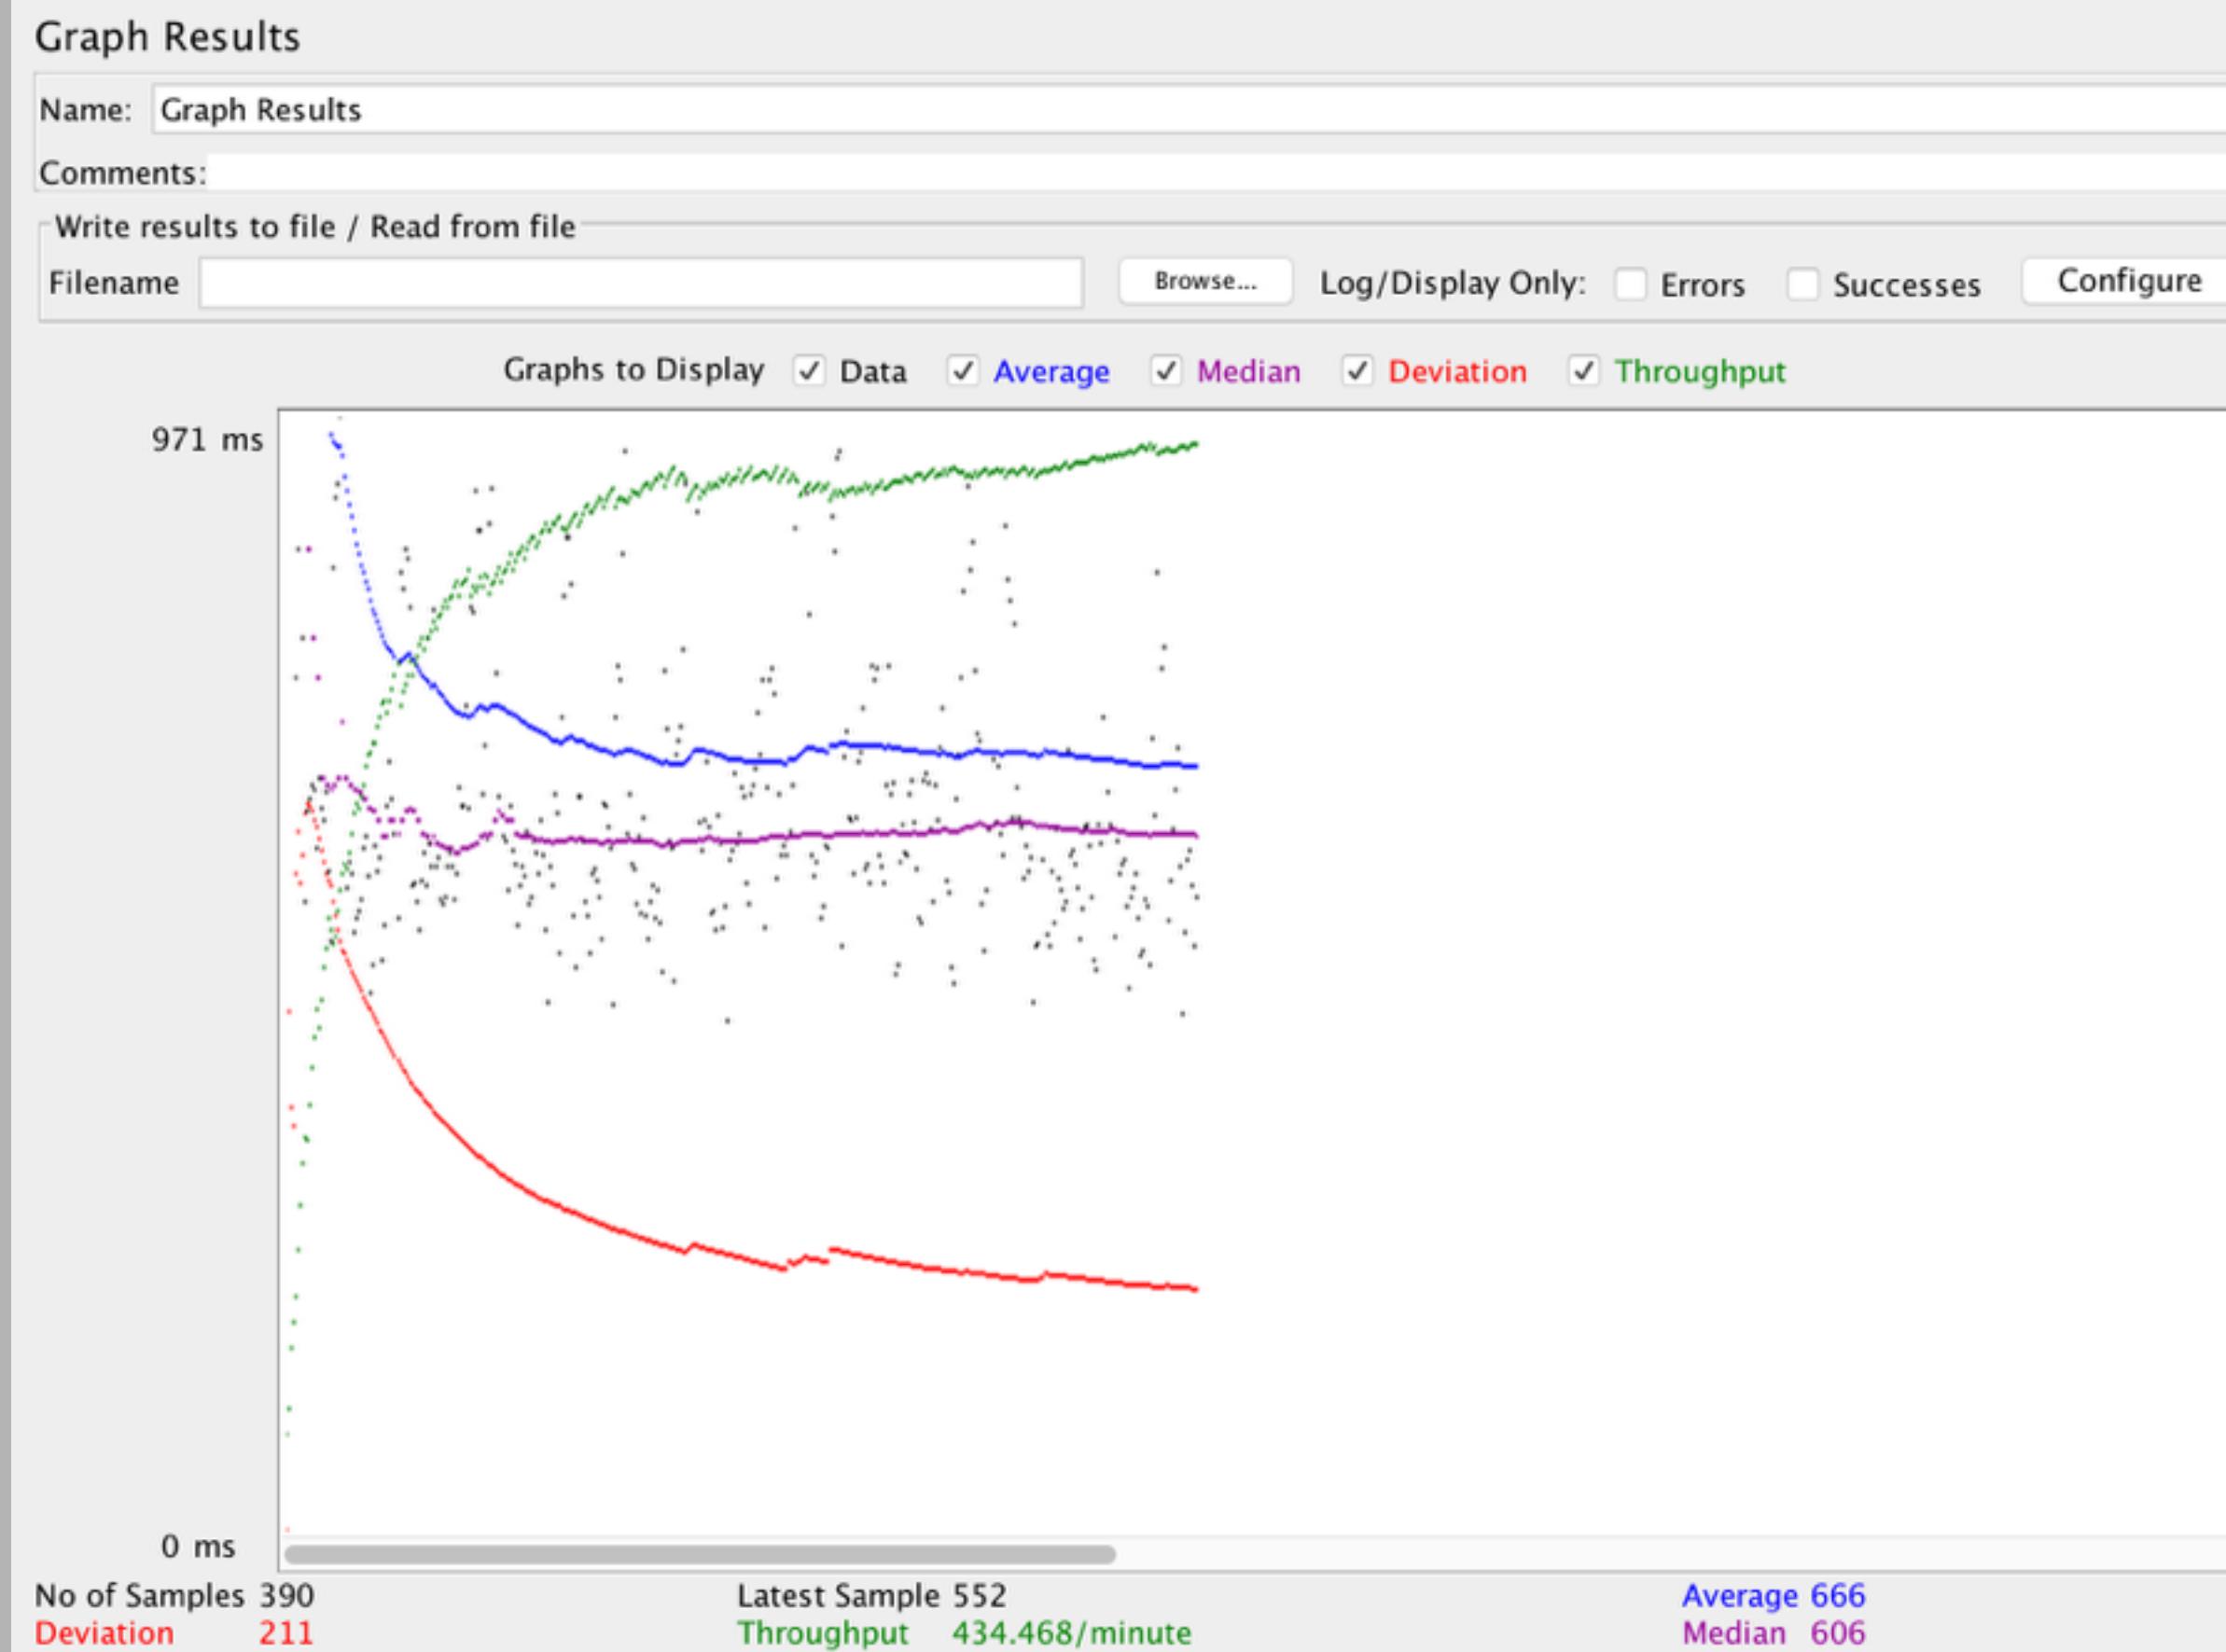
CREATE HTTP REQUEST



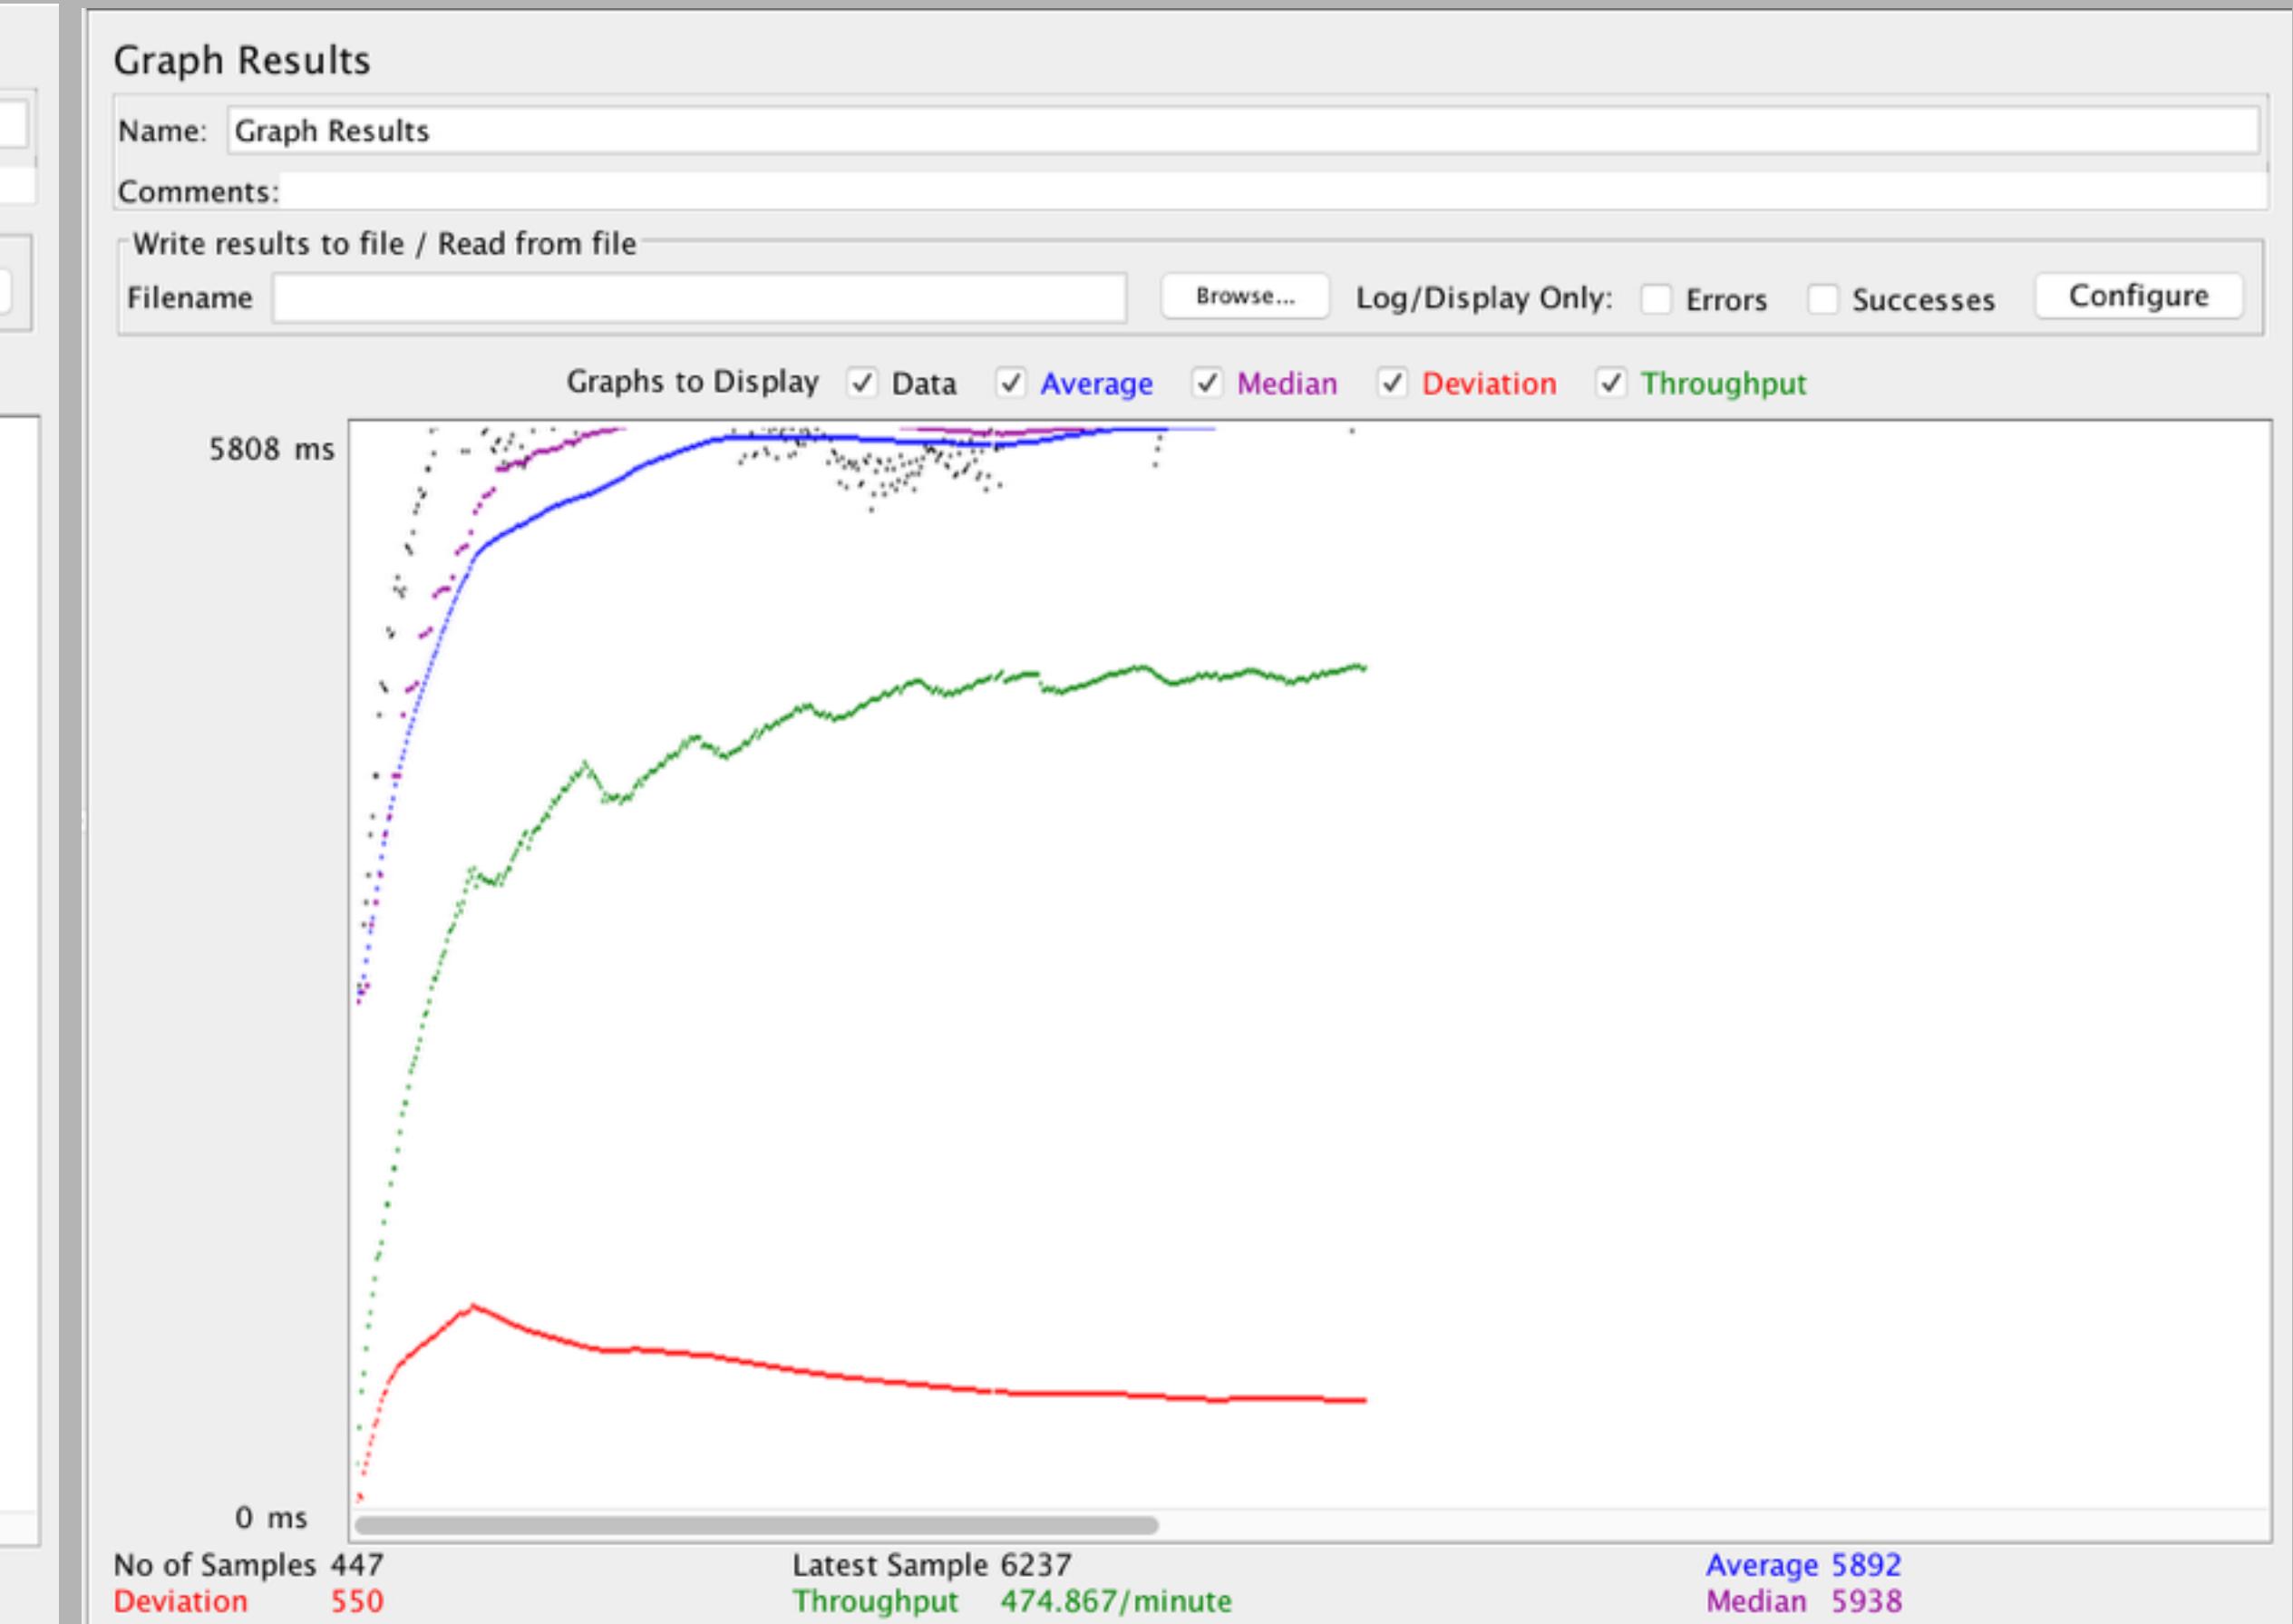
CREATE HTTP REQUEST



GRAPH RESULT



5 Users



50 Users

SUMMARY REPORT

Summary Report

Name:

Comments:

Write results to file / Read from file

Filename

[Browse...](#)

Log/Display Only:

Errors

Successes

[Configure](#)

Label	# Samples	Average	Min	Max	Std. Dev.	Error %	Throughput	KB/sec	Avg. Bytes
API1	490	620	67	1742	168.28	1.02%	7.7/sec	140.78	18660.7
TOTAL	490	620	67	1742	168.28	1.02%	7.7/sec	140.78	18660.7

5 Users

Summary Report

Name:

Comments:

Write results to file / Read from file

Filename

[Browse...](#)

Log/Display Only:

Errors

Successes

[Configure](#)

Label	# Samples	Average	Min	Max	Std. Dev.	Error %	Throughput	KB/sec	Avg. Bytes
API1	518	5273	3	7168	1070.90	9.46%	9.3/sec	157.35	17400.4
TOTAL	518	5273	3	7168	1070.90	9.46%	9.3/sec	157.35	17400.4

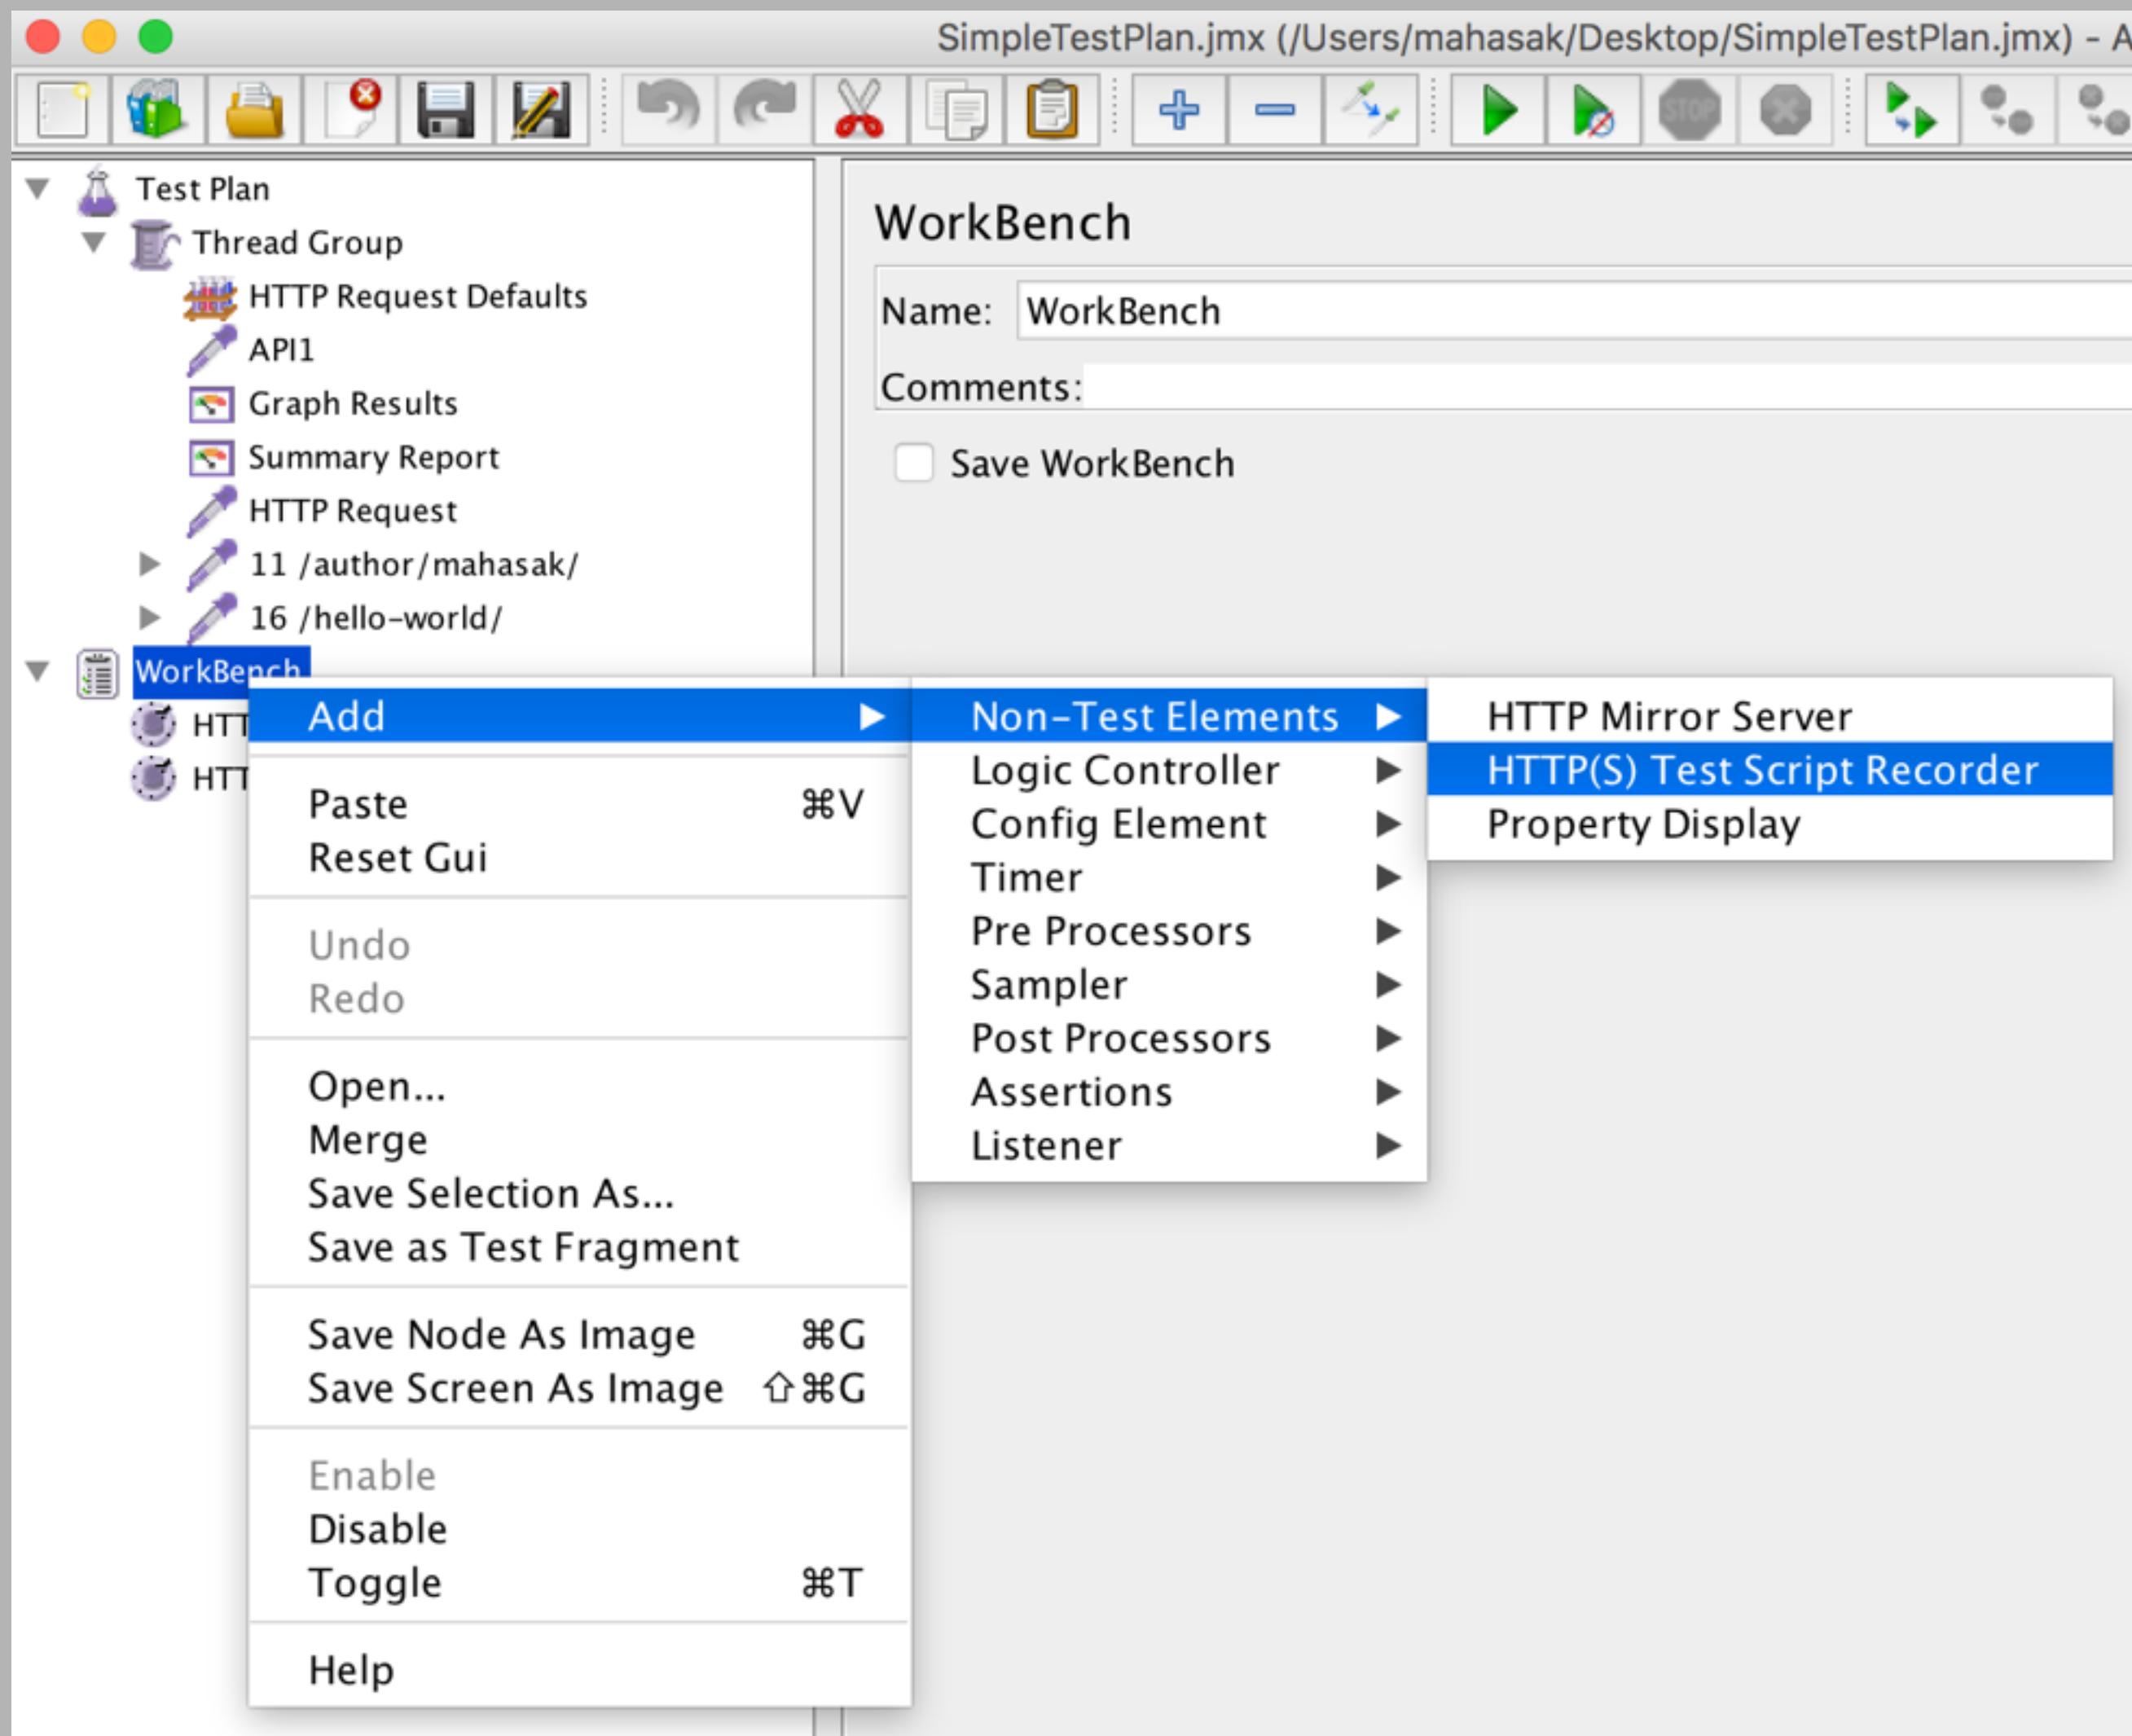
50 Users



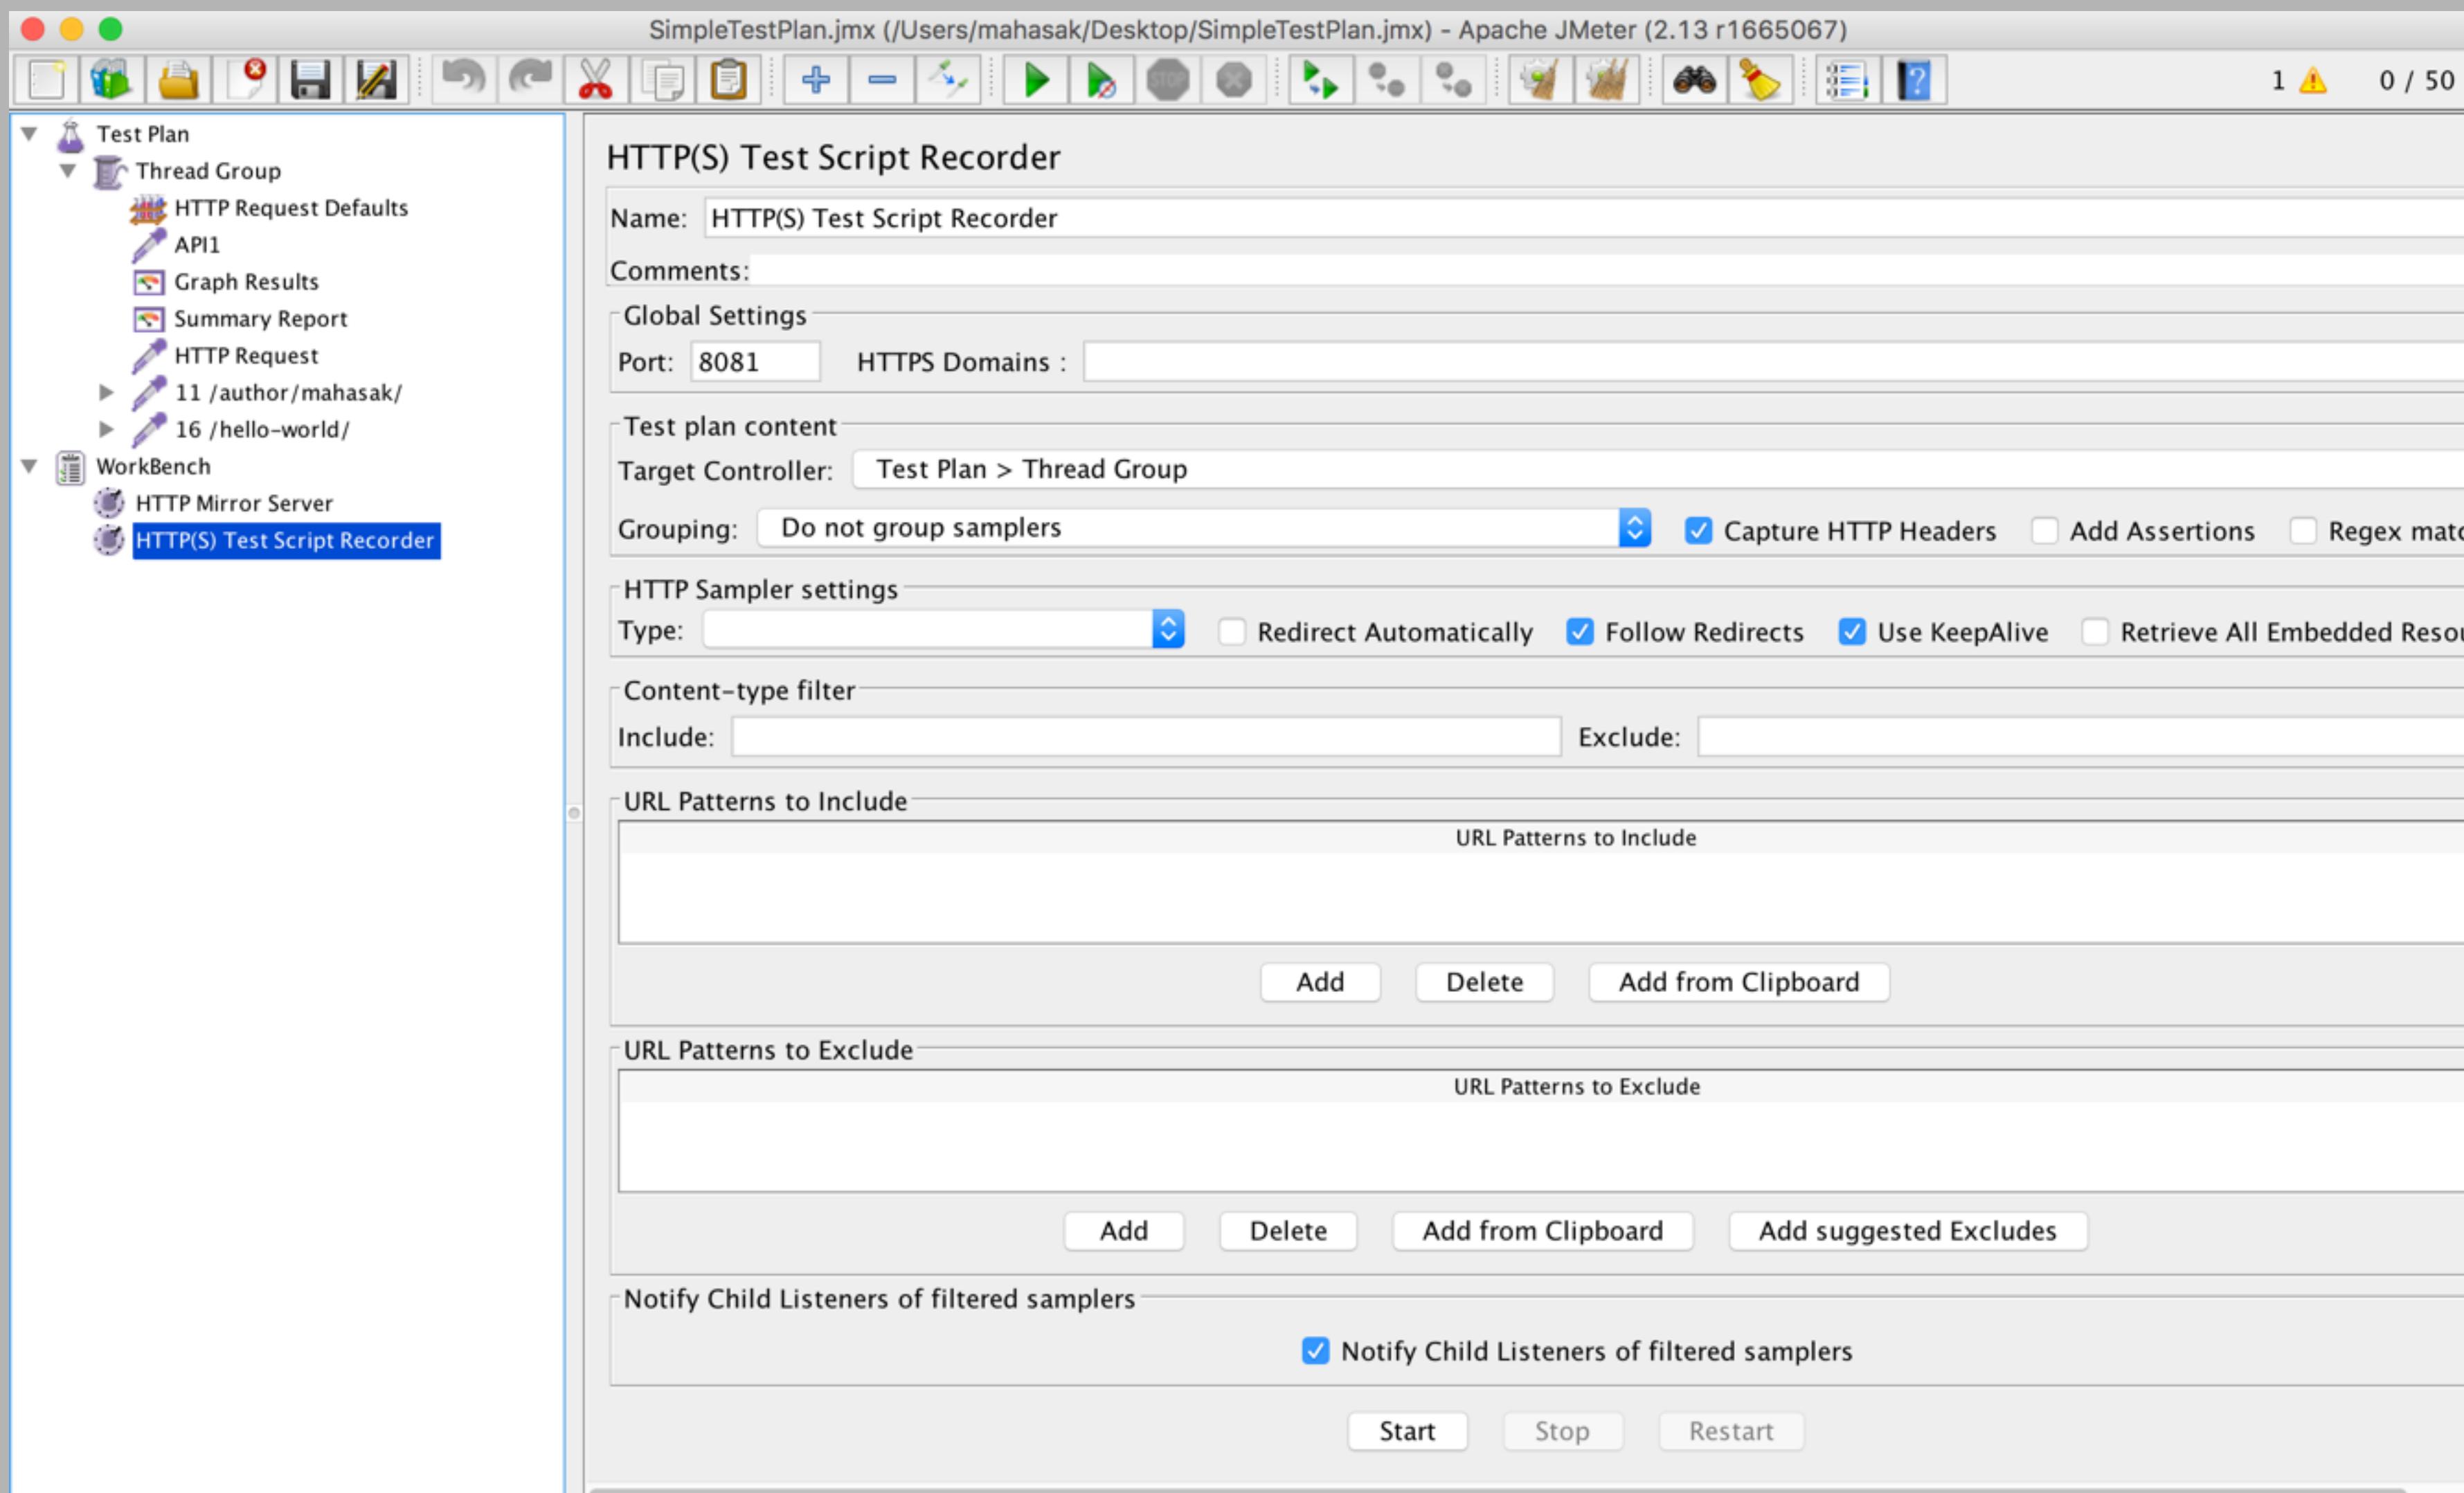
LAB 2-7

Recording Test Plan

CREATE TEST SCRIPT RECORDER



CONFIGURE TEST RECORDER





LAB 2-8

Design Workload Model

DESIGN WORKLOAD MODEL

Workload Model

User Scenarios	Percentage of workload distribution	Active users (500)
Dashboard	40	200
Feature A	10	50
Feature B	10	50
Feature C	20	100
Feature D	20	100

DESIGN WORKLOAD MODEL

The screenshot displays four separate JMeter Thread Group configurations within a single Test Plan structure. Each Thread Group is named "Workload model 1".

- Thread Group 1:** Contains a "Workload model 1" controller with five Transaction Controller elements (Dashboard, Feature A, Feature B, Feature C, Feature D) and one Transaction Controller element (Feature E). It has 500 users, a ramp-up period of 60 seconds, and a loop count of 2.
- Thread Group 2:** Contains a "Workload model 1" controller with five Transaction Controller elements (Dashboard, Feature A, Feature B, Feature C, Feature D) and one Transaction Controller element (Feature E). It has 500 users, a ramp-up period of 60 seconds, and a loop count of 2.
- Thread Group 3:** Contains a "Workload Dashboard" controller with five "Workload Feature" controllers (A-E), each containing a Transaction Controller and a Feature element (A-E). It has 500 users, a ramp-up period of 60 seconds, and a loop count of 2.
- Thread Group 4:** Contains a "Workload Dashboard" controller with five "Workload Feature" controllers (A-E), each containing a Transaction Controller and a Feature element (A-E). It has 200 users, a ramp-up period of 60 seconds, and a loop count of 2.

Common settings across all Thread Groups include:

- Action to be taken after a Sampler error: Continue
- Comments: Action to be taken after a Sampler error
- Thread Properties:
 - Number of Threads (users): 500 (for groups 1, 2)
 - Number of Threads (users): 200 (for group 4)
 - Ramp-Up Period (in seconds): 60
 - Loop Count:
 - Forever 2 (for groups 1, 2)
 - Forever 2 (for group 4)
 - Forever 2 (for group 3, which is incorrect)
 - Delay Thread creation until needed
 - Scheduler

<http://jmeter-plugins.org/>

RECOMMENDED PLUGIN TO EXPLORE

jp@gc—Stepping Thread Group

jp@gc—Response Time vs Threads

jp@gc—Active Threads Over Time

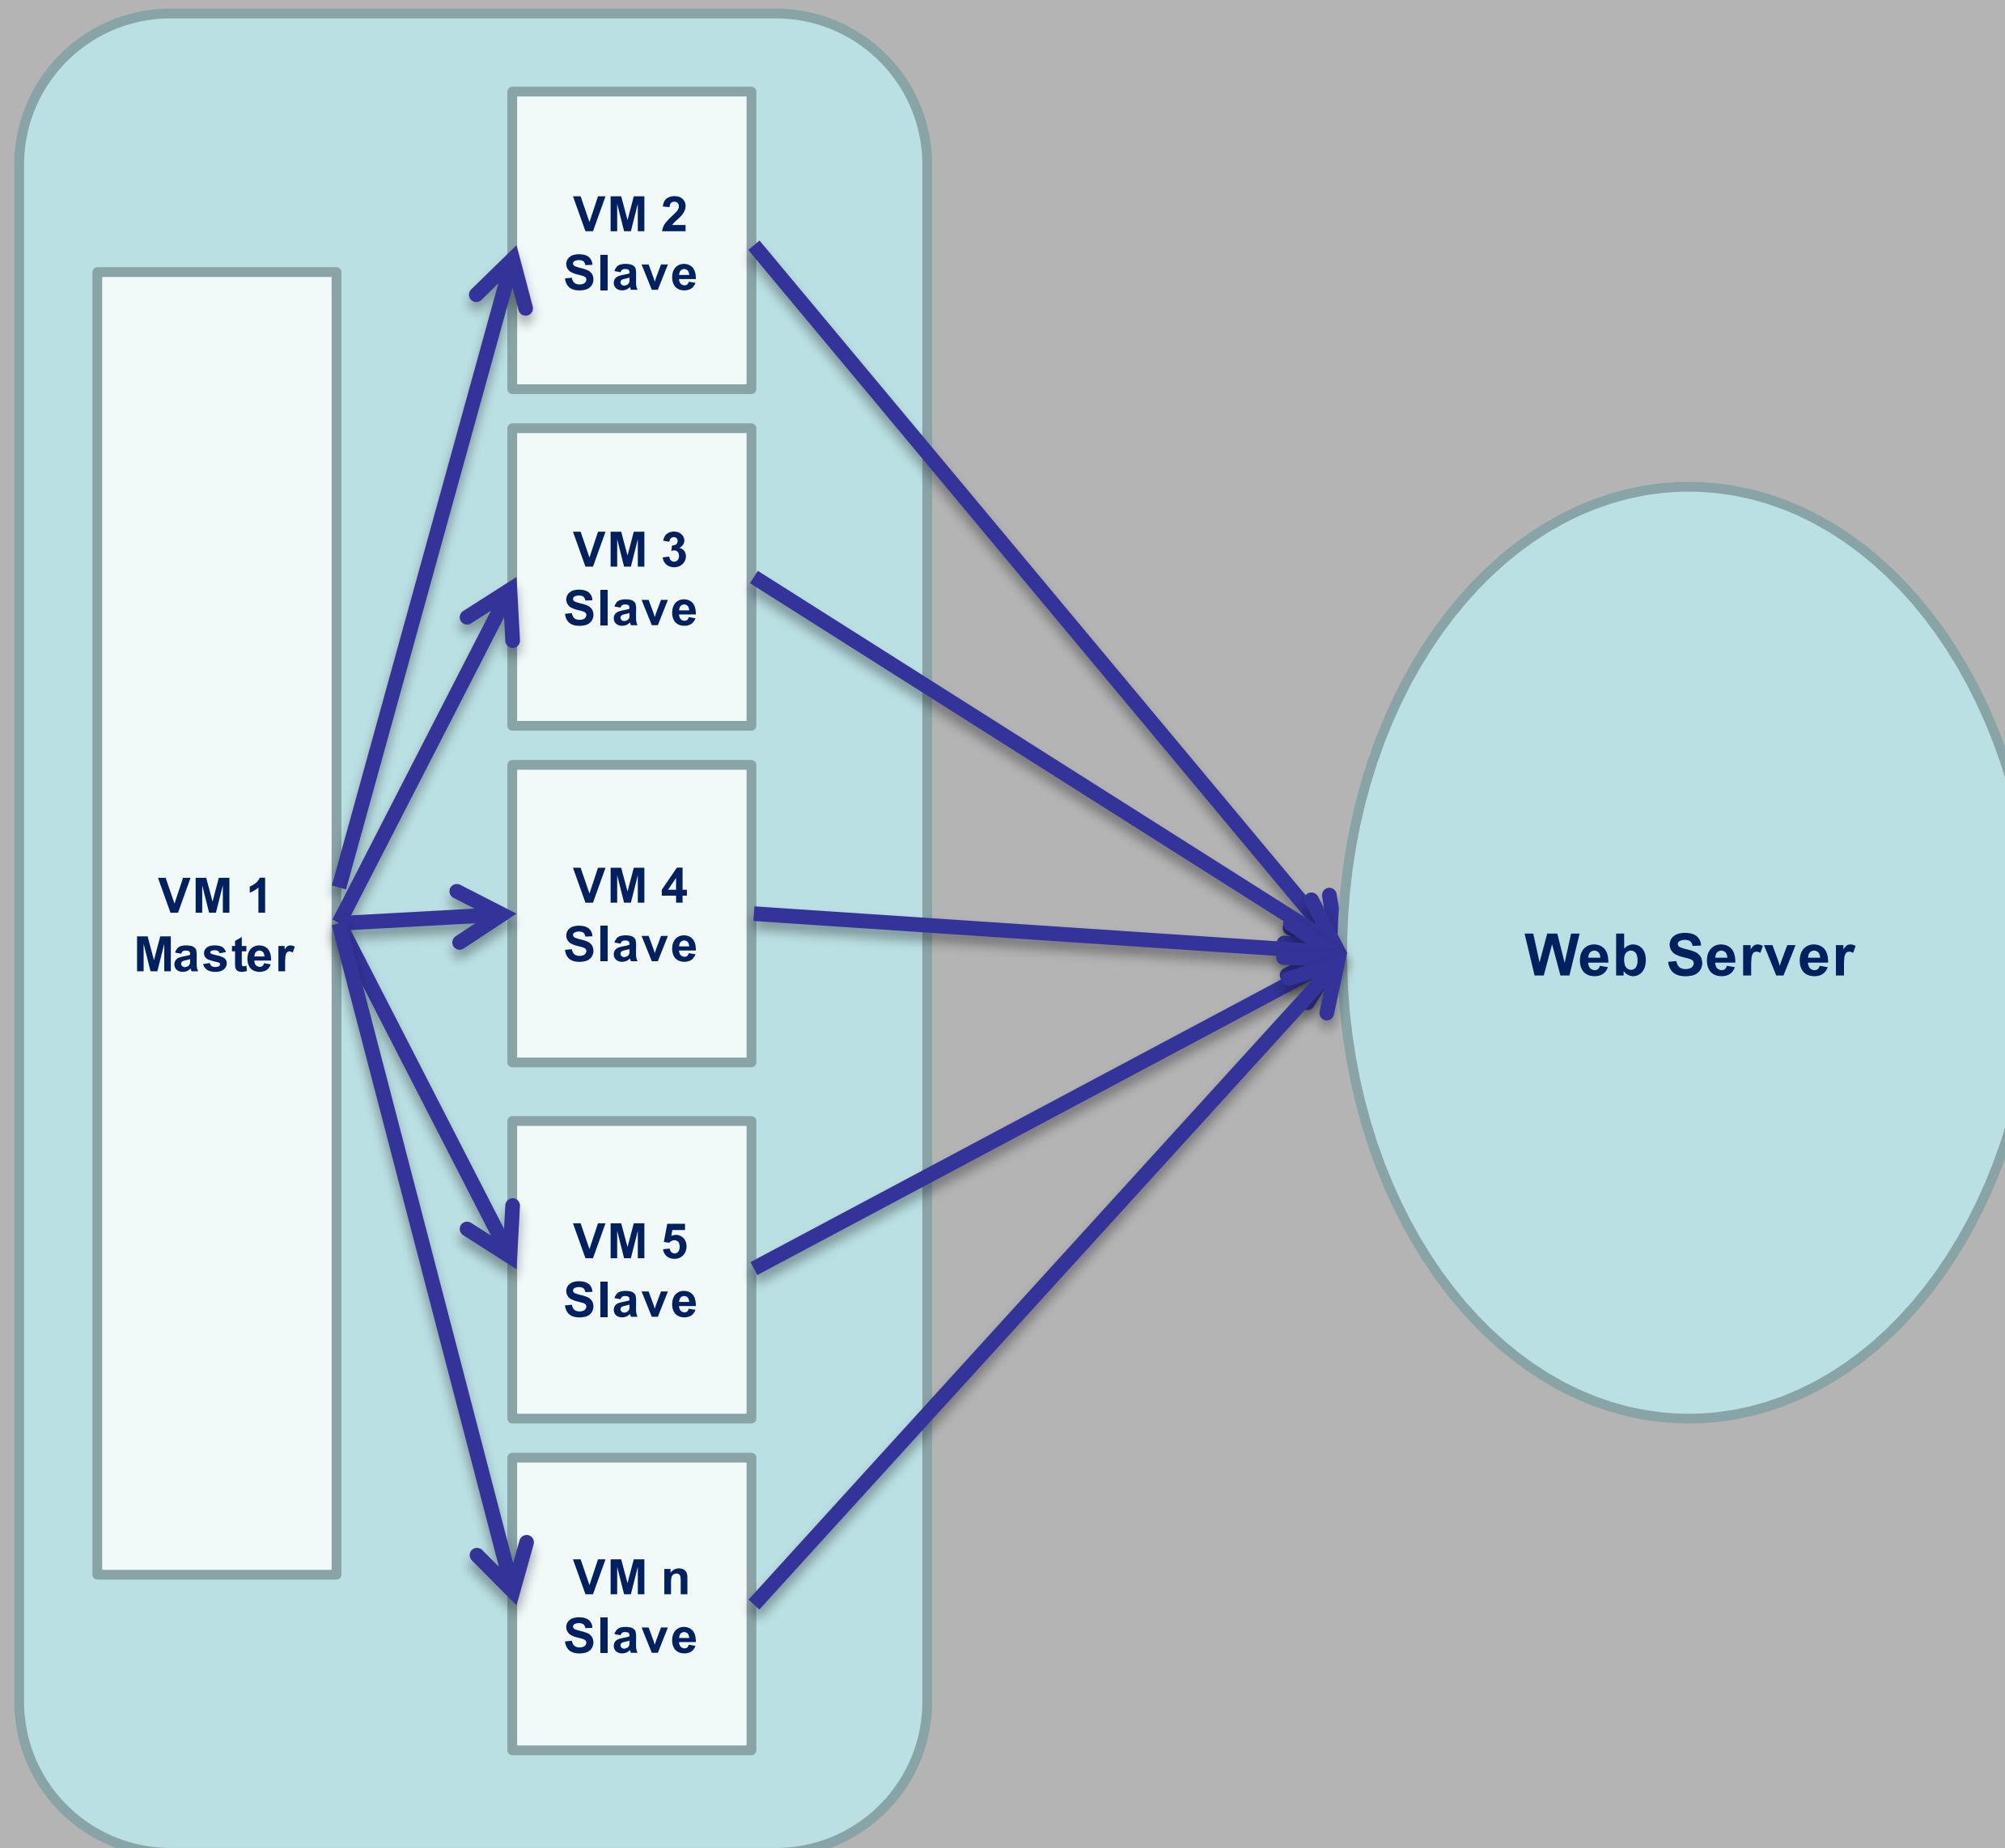
jp@gc—Response Codes per Second

jp@gc—Response Times Distribution

jp@gc - Transactions per Second

jp@gc - Response Codes per Second

CREATE HTTP REQUEST DEFAULT



Master: The system running Jmeter GUI which control the test.

Slave: The System running Jmeter-server which takes commands from the GUI and send the requests to the target system.

Target: The Web Server planned for the load test.



PERFORMANCE MONITORING

<http://newrelic.com/application-monitoring>

<https://keymetrics.io/>

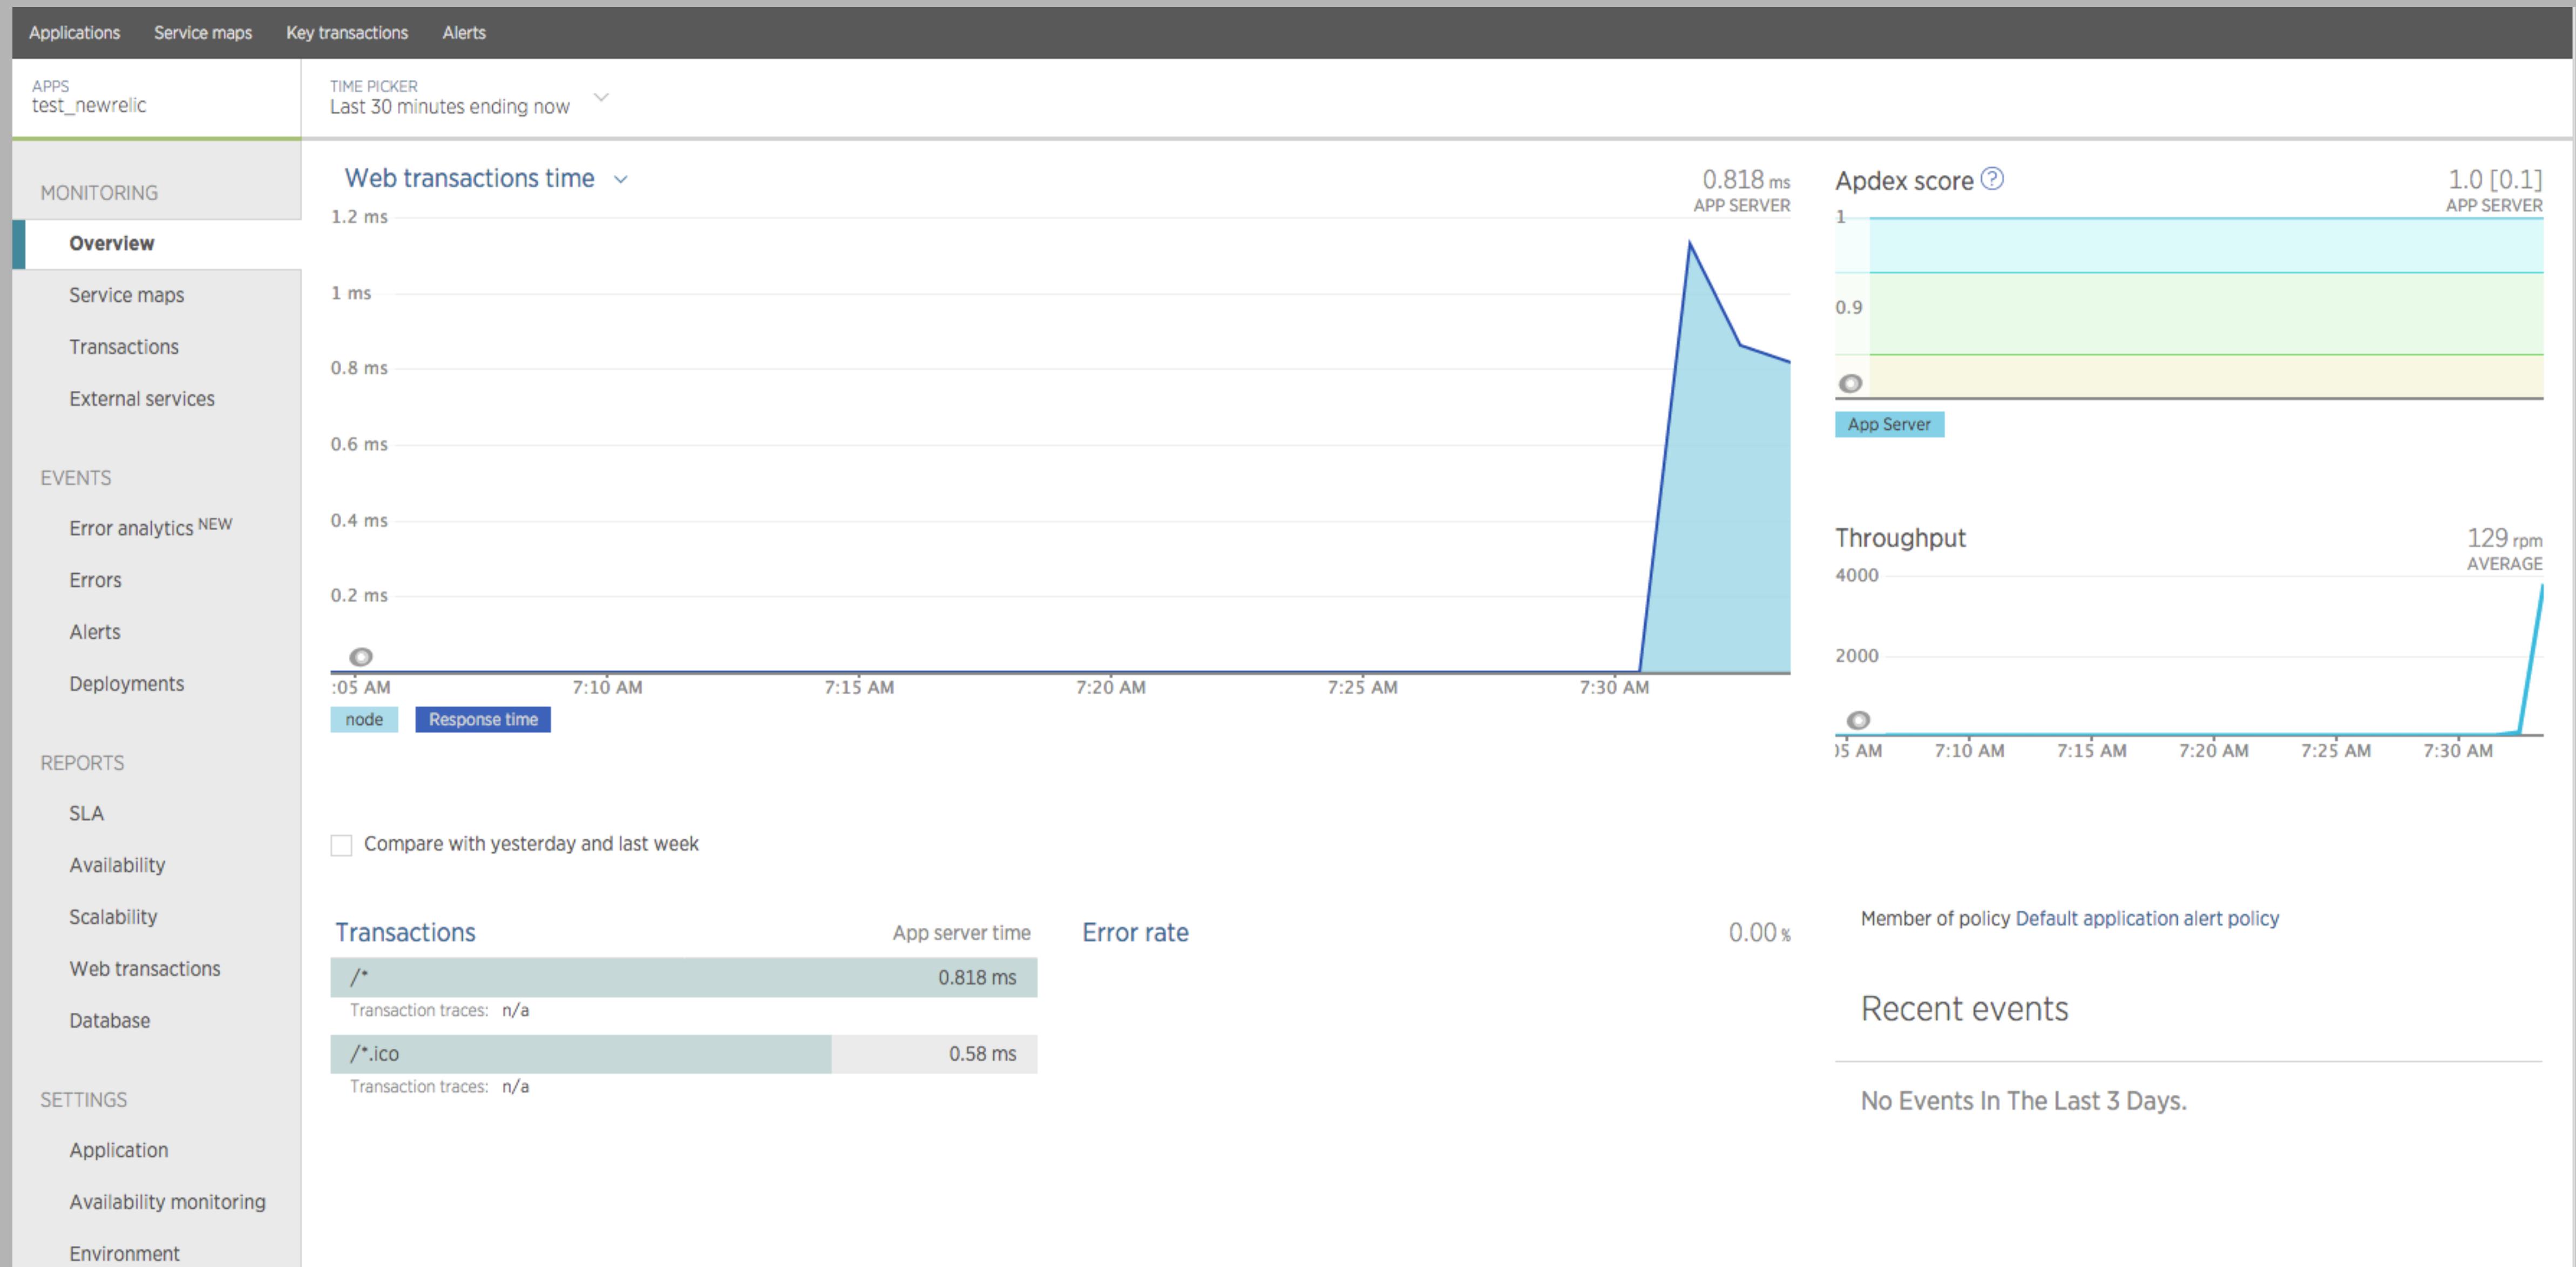
NEWRELIC

```
$ npm install newrelic
```

Copy newrelic.js from node_modules/newrelic into the root directory
Set a value for app_name, license_key

```
require('newrelic');
```

NEWRELIC





LAB 2-9

Newrelic

KEYMETRICS

```
$ npm install newrelic
```

Copy newrelic.js from node_modules/newrelic into the root directory
Set a value for app_name, license_key

```
require('newrelic');
```

KEYMETRICS

The screenshot shows the Keymetrics pricing page with three main plans: Starter, Premium, and Enterprise.

Starter (Free):
OUR OPEN SOURCE TOOL, PM2
Auto clustering
Log management
Os application update
Deployment system
Behavior configuration
Memory management
Live monitoring only
Live remote actions
No Keymetrics notifications
REGISTER

Premium Starting at \$29 per host/month
All features from the free plan
+ EMAIL NOTIFICATION
+ 14-DAYS DATA RETENTION
+ CUSTOM METRICS HISTORIC
+ EXCEPTION REPORTING
+ REALTIME LOGS
+ ACL USER MANAGEMENT
+ ROUTES MONITORING
+ CPU/MEM PROFILING
REGISTER

Enterprise Custom
Dedicated instance my-enterprise.keymetrics.io
Unlimited users
Unlimited buckets
Unlimited retention
Up to 10 servers monitored
PM2 professional support
CONTACT US

TOUR DOC PRICING BLOG LOGIN SIGNUP

Empower your Node.js application, now



LAB 2-9

Keymetrics

KEYMETRICS

Buckets

0 BUCKETS

You haven't created any Bucket.
A bucket allows you to manage server/apps of the same context and to share it with your team

+ New bucket



Keymetrics Bucket

test_keymetrics

Comment (about servers you will manage)

Create Bucket

What's a bucket?

A **BUCKET** is like a container on which multiple servers and multiple apps can be attached to. It's very convenient when you need to monitor a bunch of apps that belongs to the same companies or environment for example.

Each **BUCKET** has a public and private key. This will allow PM2 to interact with a determined Bucket and most importantly, to communicate in a fully secured manner (with AES256).

KEYMETRICS

```
$ pm2 link [private] [public]
```

The screenshot shows the Keymetrics dashboard interface. At the top, there's a header with the project name "test_keymetrics", a "Condensed Dashboard" button, and search icons. To the right are "Public key" and "Secret key" fields, with a "Show" button. A message encourages upgrading the bucket for notifications and persistence.

The main area displays "Server #1 - bahamutzero" with its IP and hostname. It includes a table with columns for Status, PM2 version, Node version, CPU load average, and Hostname. The status is "Online". The PM2 version is "1.0.0", Node version is "v5.0.0", and CPU load average is "1.9". The Hostname is "bahamutzero".

Below the table, there's a detailed view of the process "test". It shows the process is online and has been online for a few seconds. It has 0% CPU usage and 34.7 MB of memory. There are 0 errors. The HTTP average is listed as "Versioning module not activated (no .git found in the project)".

At the bottom of the dashboard, there are sections for monitoring more servers, installing pm2-server-monit, and documentation links.

Status	PM2 version	Node version	CPU load average	Hostname
Online	1.0.0	v5.0.0	1.9	bahamutzero

Process details for "test":

test online online for a few seconds More Alerts Logs	CPU 0 %	MEM 34.7 MB	Errors 0	HTTP avg. configure	Versioning Versioning module not activated (no .git found in the project) Processes: 1 Restart: 0 Mode: fork Interpreter: node	Process Metadata Deployments Restart
--	------------	----------------	-------------	------------------------	--	--

Ports 80 (tcp out) and 43554 (tcp in/out) must be opened

Monitor more servers and applications with:

```
$ pm2 link gipoi2qluqo6kw5 8eqatwqnn4cb1sv [machine name]
```

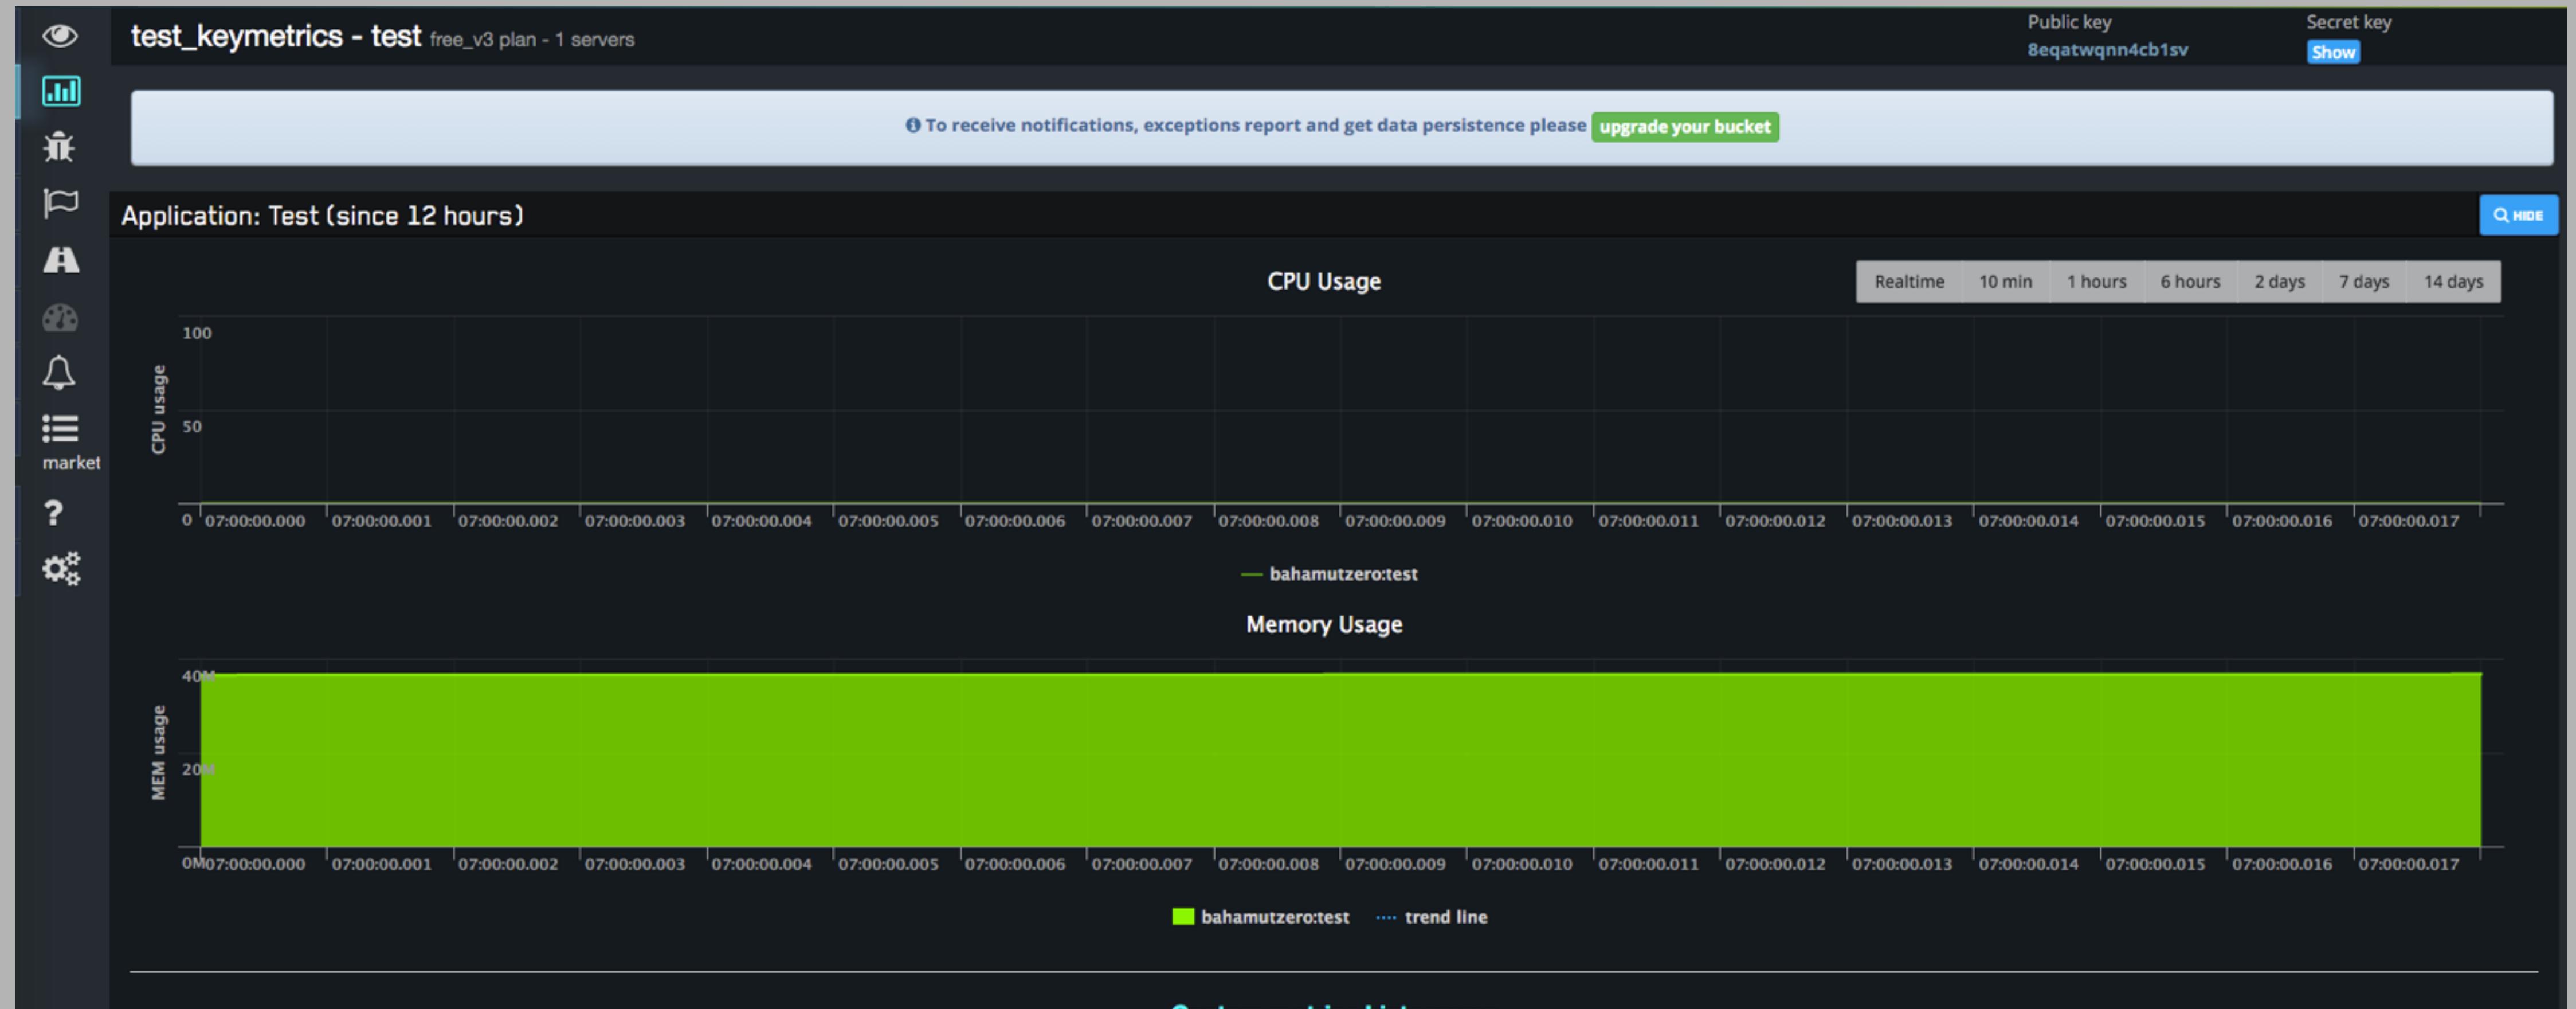
Monitor your server with:

```
$ pm2 install pm2-server-monit
```

Documentation

<http://docs.keymetrics.io/>

KEYMETRICS

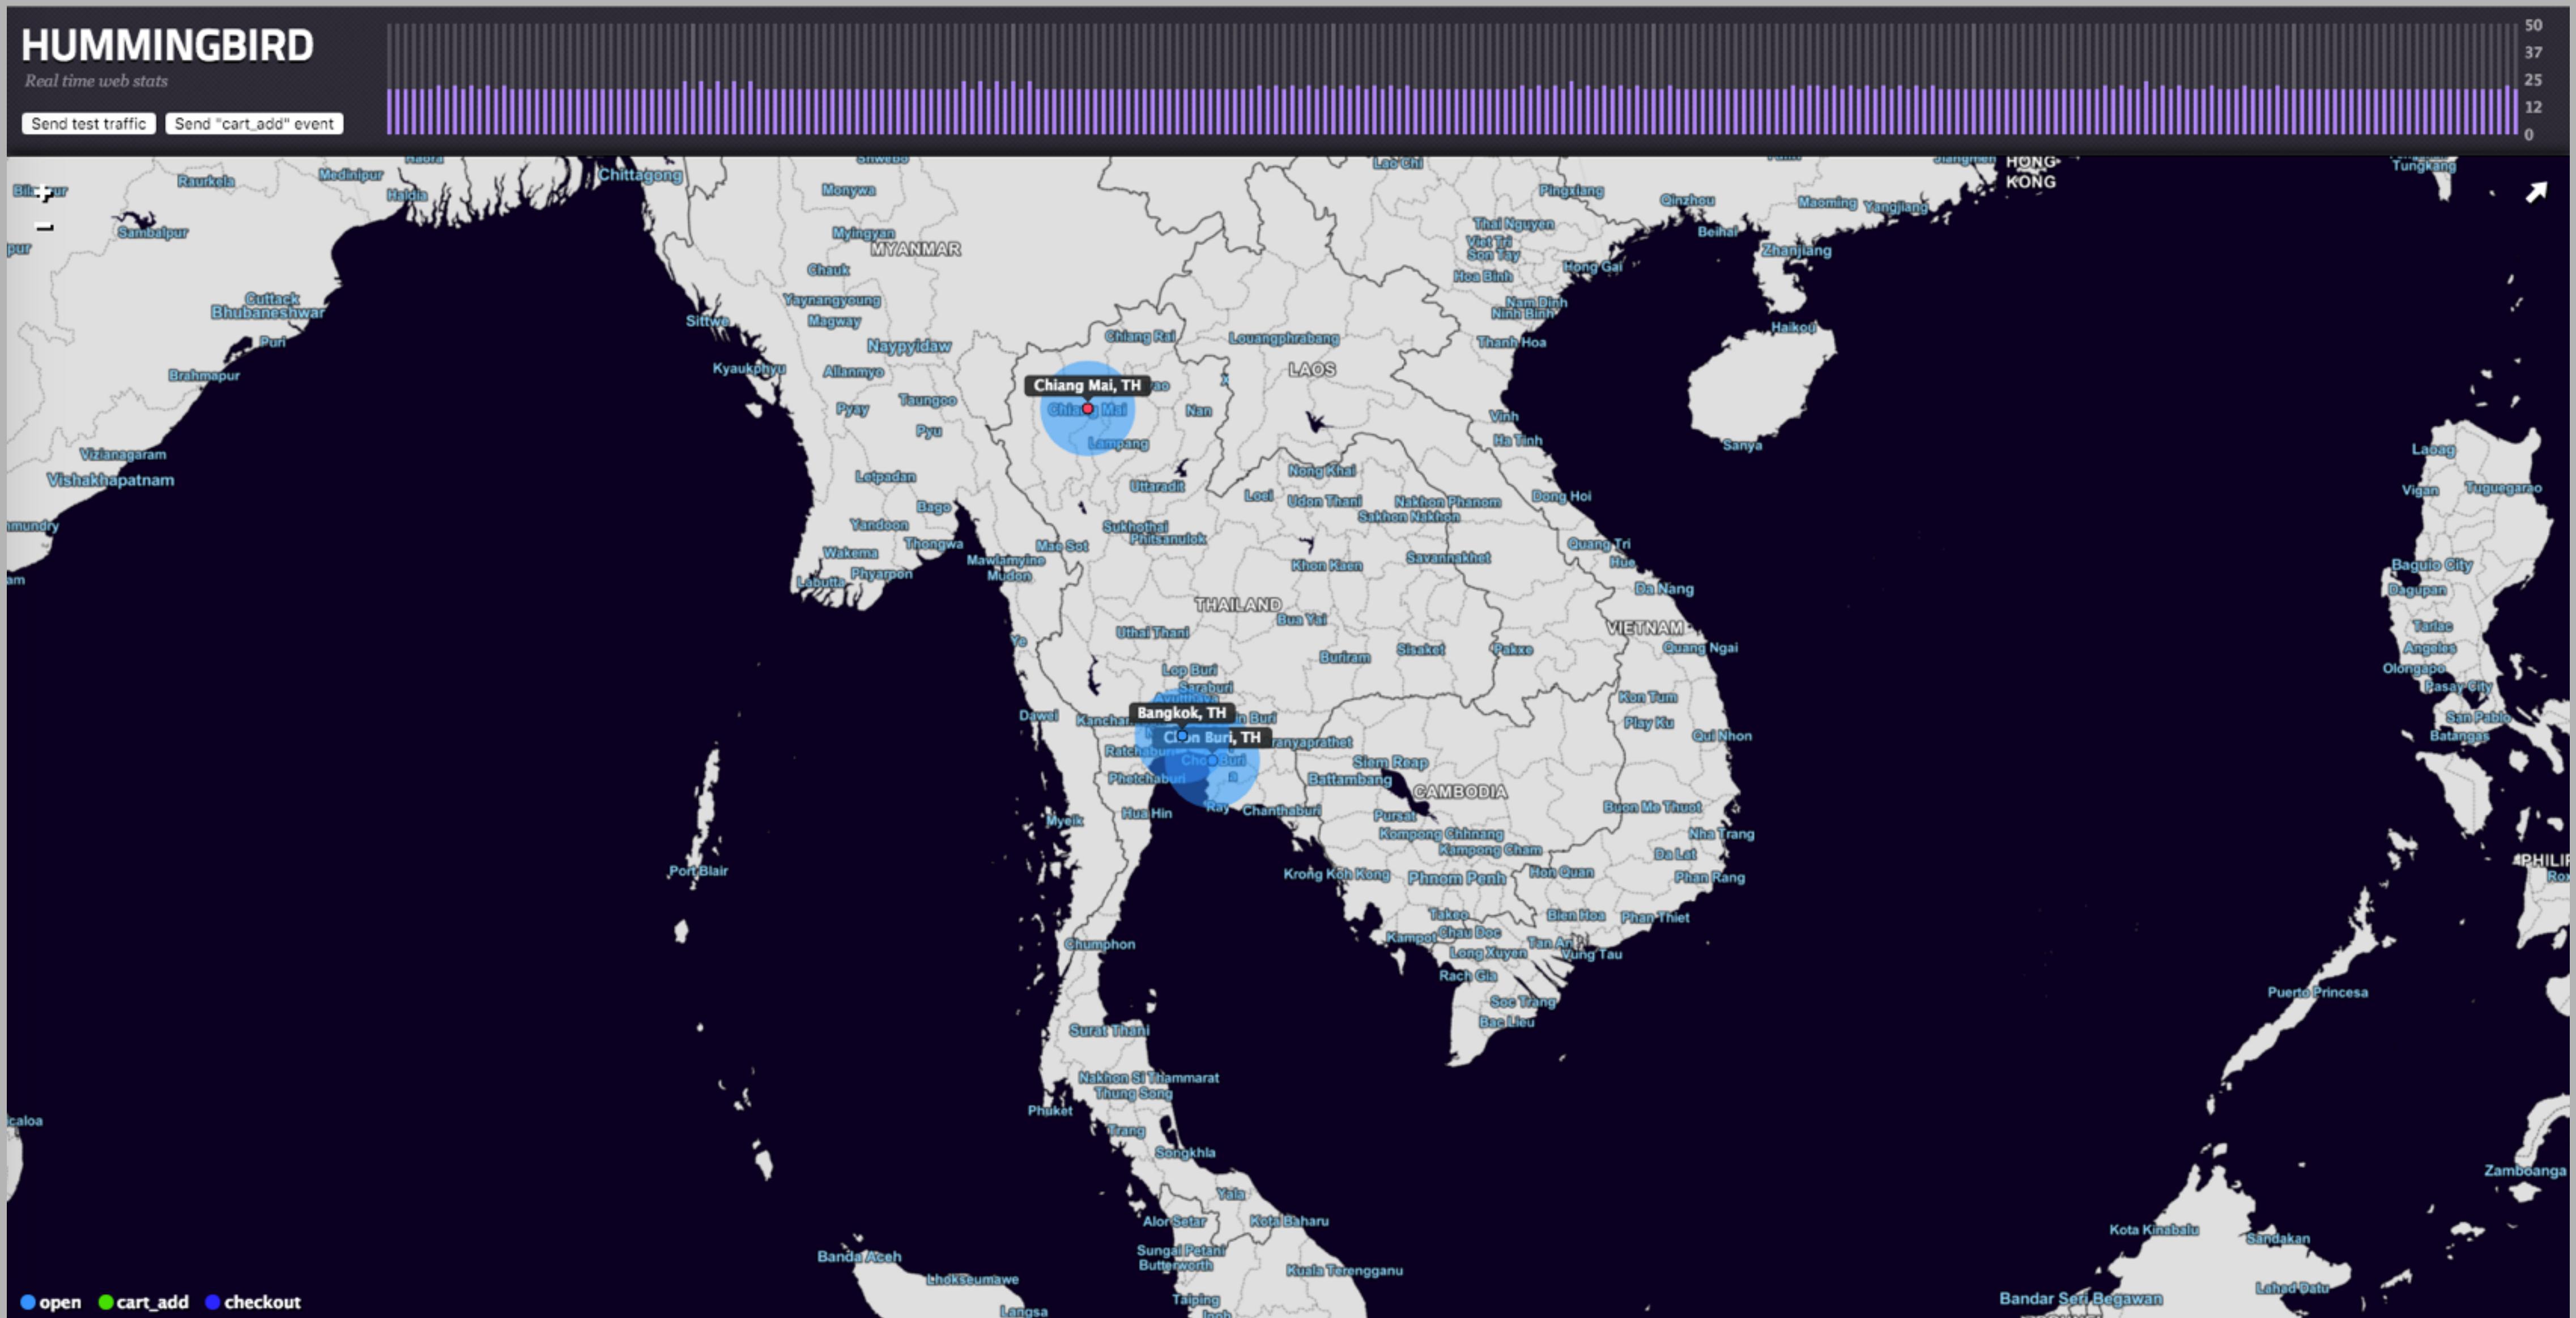




REAL-TIME APPLICATION ANALYTICS SAMPLE CASE

<https://github.com/mnutt/hummingbird>

HUMMINGBIRD



HUMMINGBIRD CLIENT

```
var Buffer = require('buffer').Buffer;
var dgram = require('dgram');

var sock = dgram.createSocket("udp4");

var data = {
  "ip": "127.0.0.1",
  "timestamp": "Sat Oct 23 2010 21:39:35 GMT+0700 (EDT)",
  "url_key": 123,
  "product_id": 456
};

var buf = new Buffer(JSON.stringify(data));

console.log('buf:' + buf);

setInterval(function() {
  sock.send(buf, 0, buf.length, 8000, "0.0.0.0");
}, 100);
```



SECURITY DISCUSSION

CONFIGURATION : SECURITY HTTP HEADER

- Strict-Transport-Security enforces secure (HTTP over SSL/TLS) connections to the server
- X-Frame-Options provides clickjacking protection
- X-XSS-Protection enables the Cross-site scripting (XSS) filter built into most recent web browsers
- X-Content-Type-Options prevents browsers from MIME-sniffing a response away from the declared content-type
- Content-Security-Policy prevents a wide range of attacks, including Cross-site scripting and other cross-site injections

```
# nginx.conf

add_header X-Frame-Options SAMEORIGIN;
add_header X-Content-Type-Options nosniff;
add_header X-XSS-Protection "1; mode=block";
add_header Content-Security-Policy "default-src 'self'"
```

AUTHENTICATION : BRUTE FORCE

```
$ npm install ratelimiter
```

```
#app.js
var email = req.body.email;
var limit = new Limiter({ id: email, db: db });

limit.get(function(err, limit) {

});
```

RATE LIMIT

SESSION MANAGEMENT

Cookie Flags

secure - this attribute tells the browser to only send the cookie if the request is being sent over HTTPS.

HttpOnly - this attribute is used to help prevent attacks such as cross-site scripting, since it does not allow the cookie to be accessed via JavaScript.

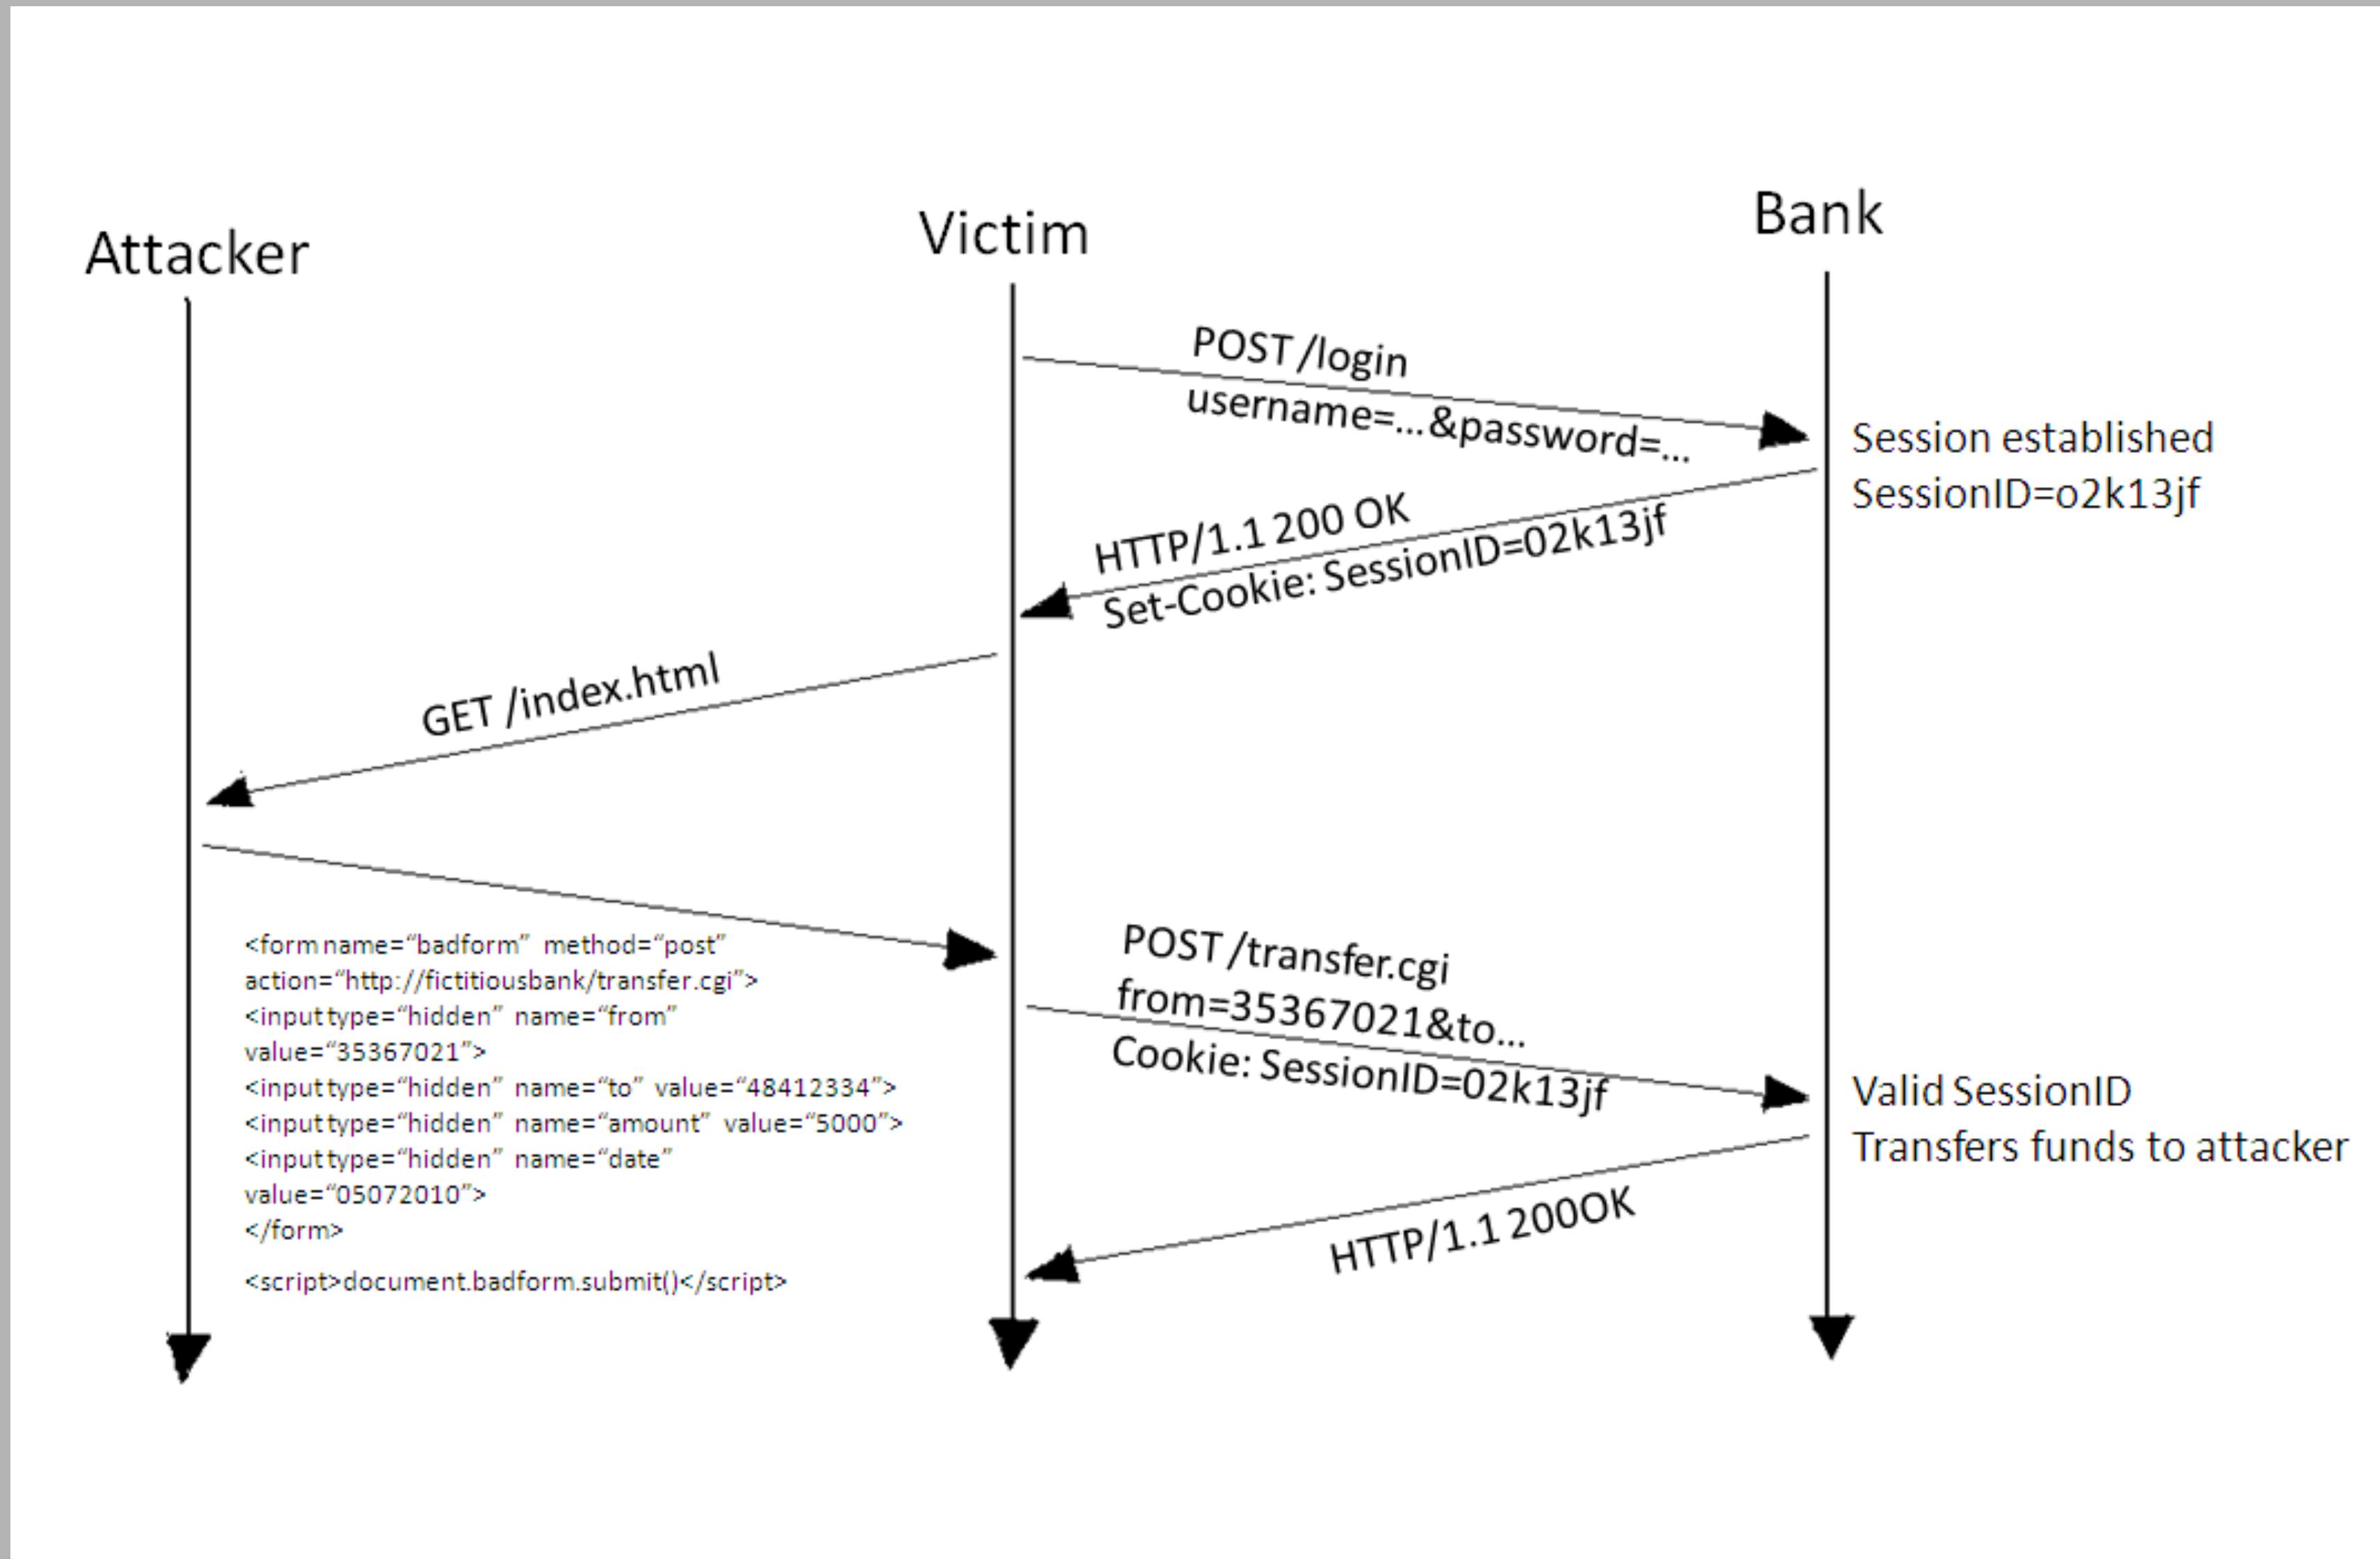
Cookie Scope

domain - this attribute is used to compare against the domain of the server in which the URL is being requested. If the domain matches or if it is a sub-domain, then the path attribute will be checked next.

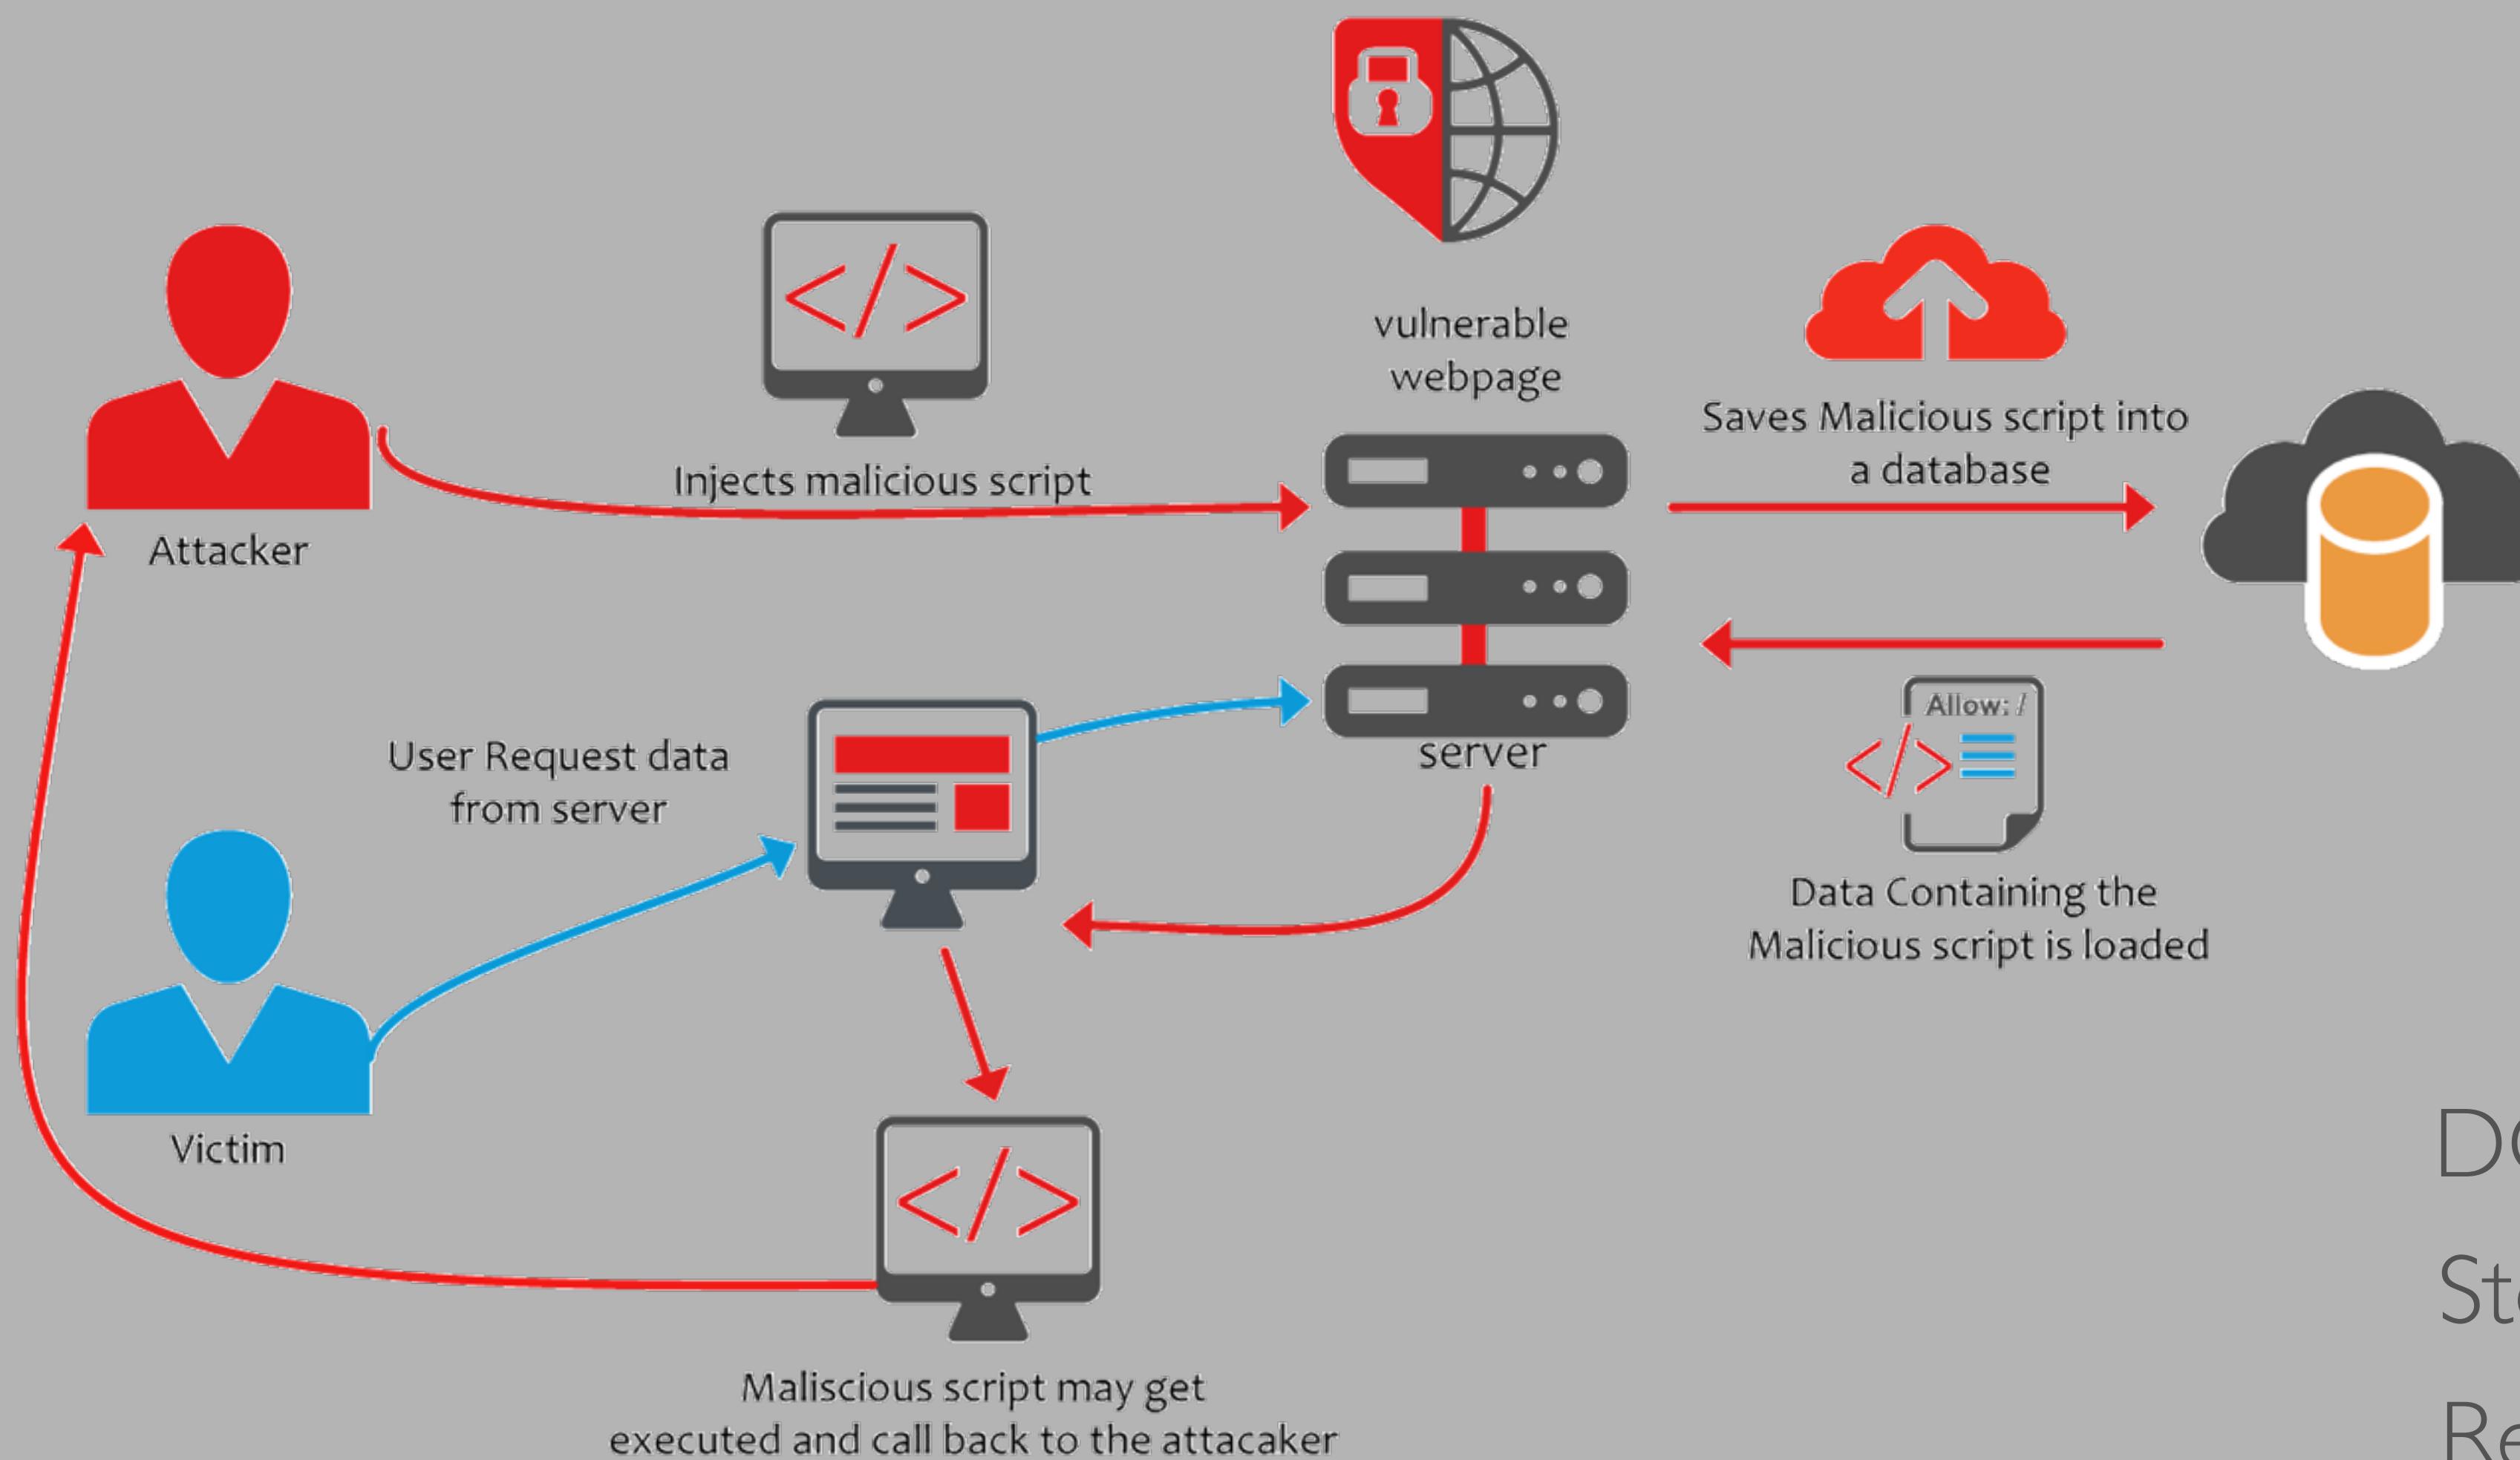
path - in addition to the domain, the URL path that the cookie is valid for can be specified. If the domain and path match, then the cookie will be sent in the request.

expires - this attribute is used to set persistent cookies, since the cookie does not expire until the set date is exceeded

CSRF: CROSS SITE REQUEST FORGERY



XSS : CROSS SITE SCRIPTING

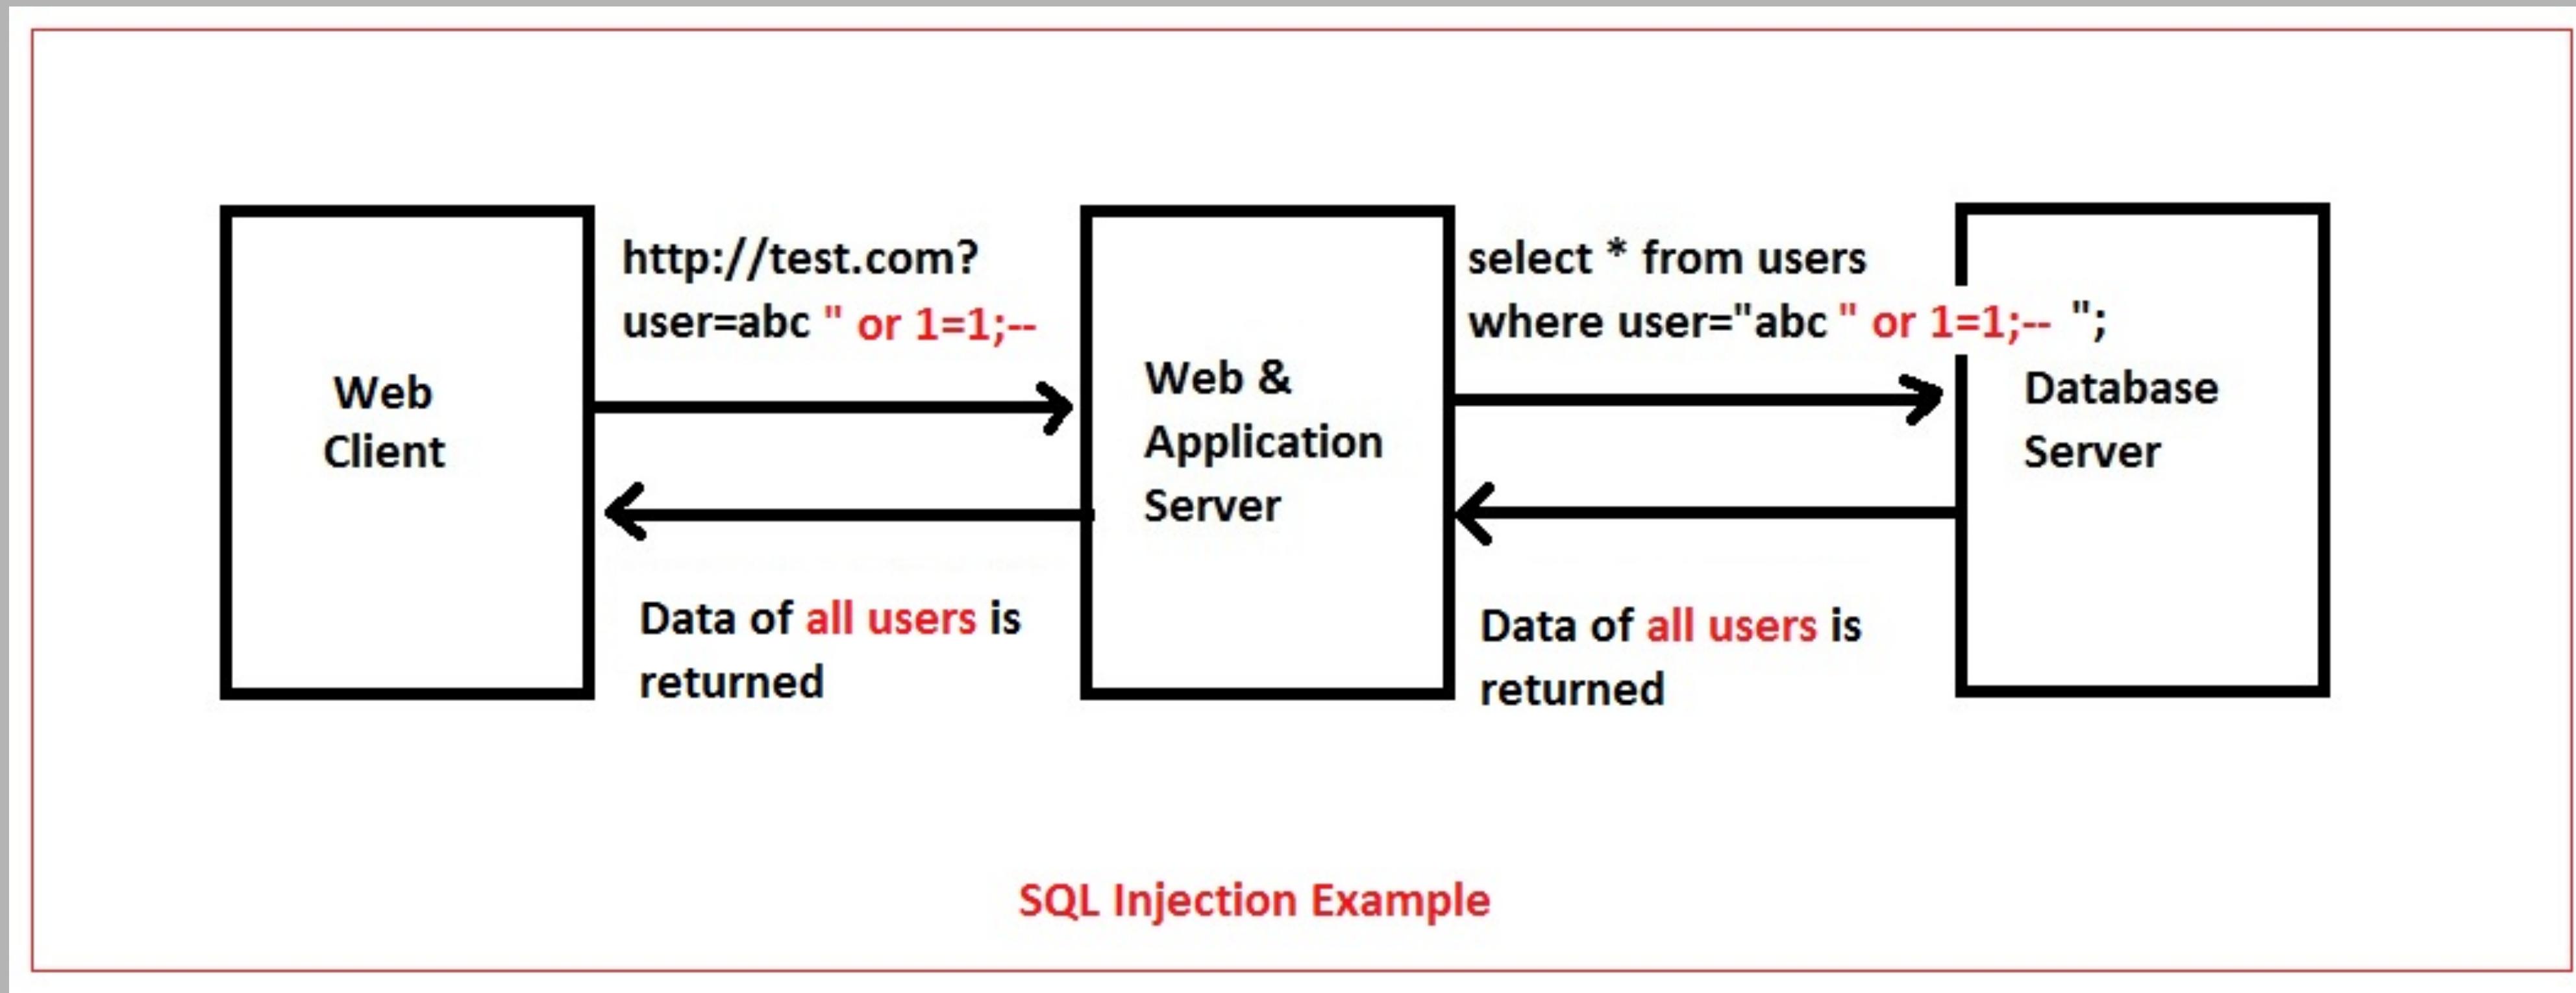


DOM Based XSS

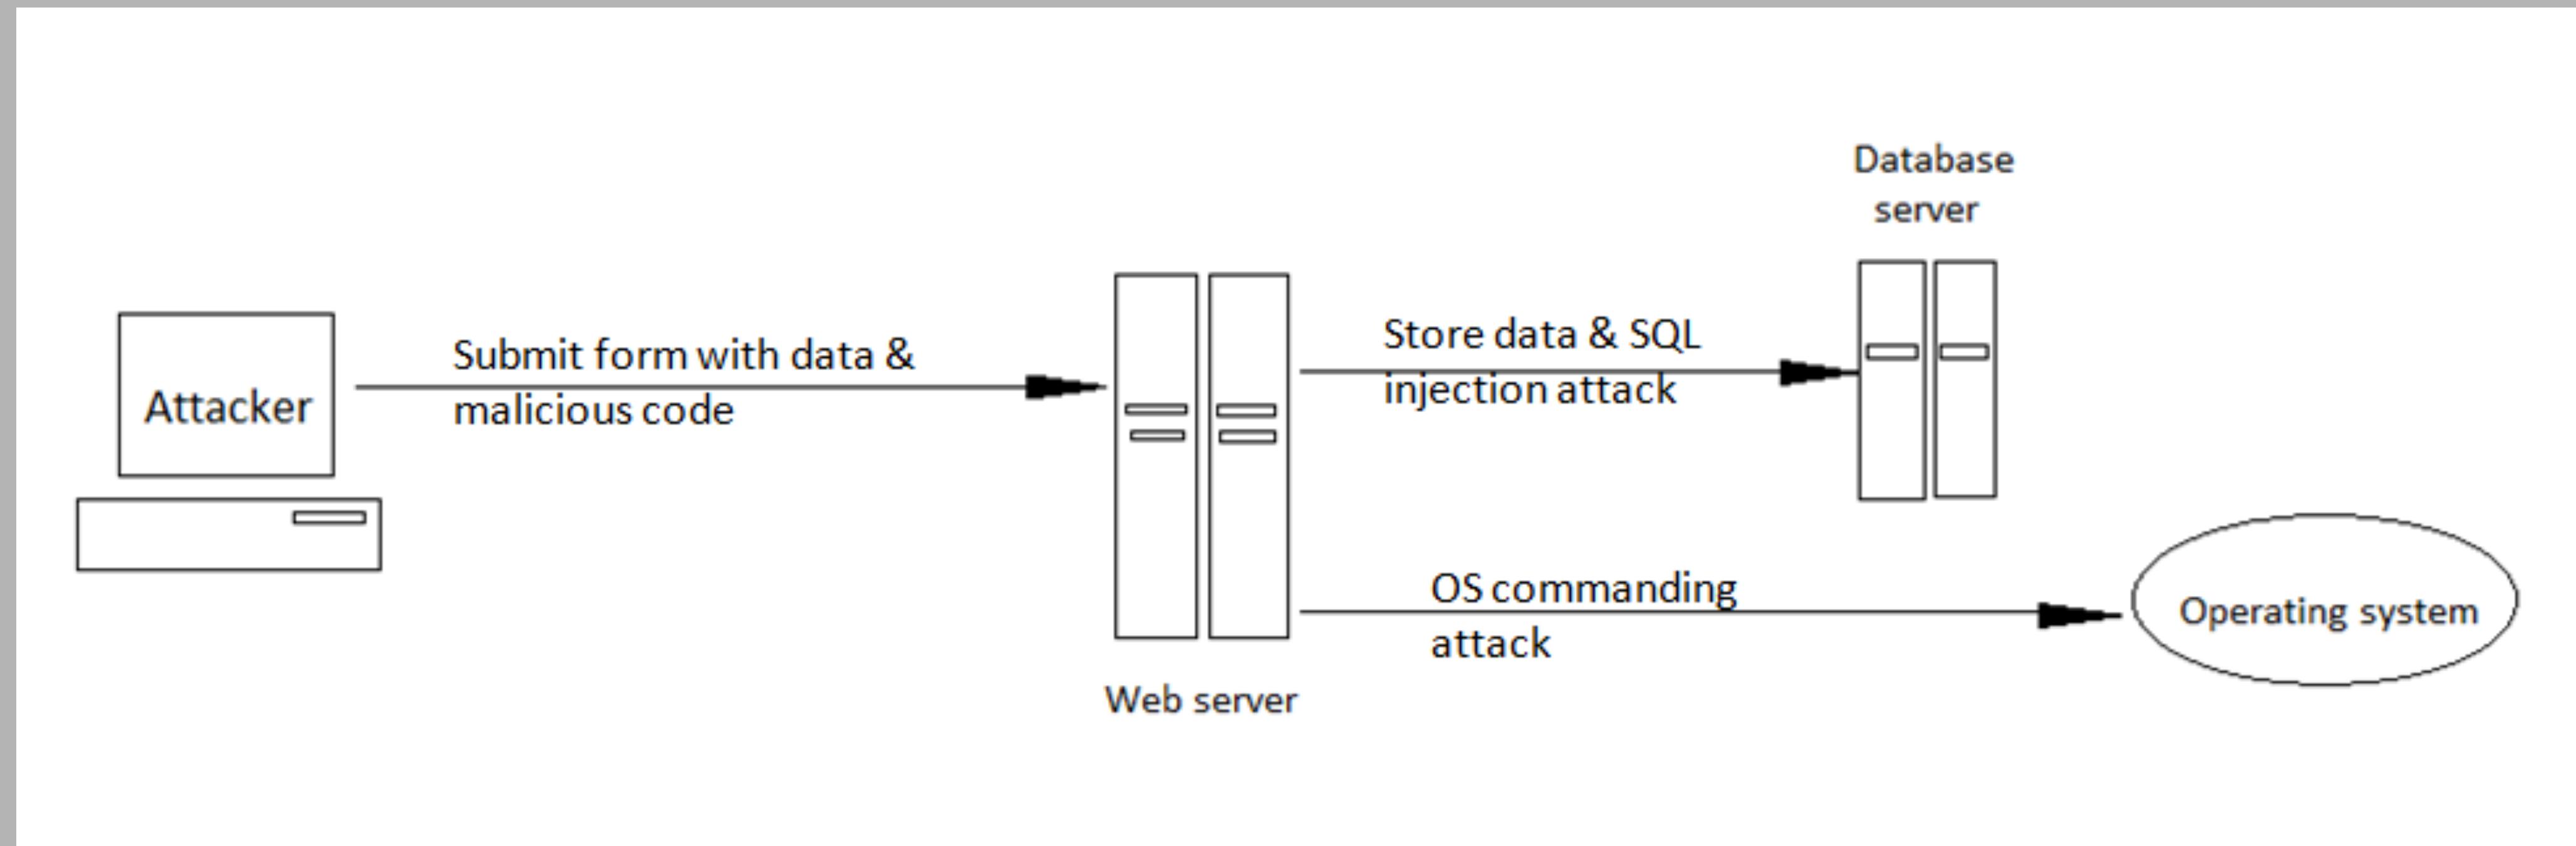
Stored XSS

Reflected XSS

SQL INJECTION



COMMAND INJECTION



Original URL: <https://example.com/downloads?file=user1.txt>

Attacking URL: <https://example.com/downloads?file=%3Bcat%20/etc/passwd>



Q & A